

Asymptotic expansions of inverse functions on a real branch

Power-logarithmic reversion with real exponents, from two Mathematica Stack Exchange questions to a unified theory and a Wolfram Language package

Vladimir Reshetnikov*

September 2026

Abstract

The inverse of $f(x) = x + x^2(1 + \log x)$ on the positive real axis has no Taylor expansion at 0: its expansion is $y - y^2(1 + \log y) + y^3(2 \log^2 y + 5 \log y + 3) - \dots$, with a polynomial in $\log y$ at every power. The inverse of $x + x^{\sqrt{2}}$ has irrational exponents, $y - y^{\sqrt{2}} + \sqrt{2} y^{2\sqrt{2}-1} - \dots$. The built-in functions `InverseSeries`, `Series`, `Asymptotic` and `InverseFunction` of the Wolfram Language either fail on such inputs or return the germ at a complex root. The same algebra of finite *power-log expansions* also governs forward asymptotics with irrational exponents, such as $(1 + x + x^{\sqrt{2}})^{\sqrt{2}}$ at infinity, where `Asymptotic` returns its input unchanged.

We develop an ordered calculus of power-log blocks with exact truncated arithmetic, marker Lagrange-Bürmann inversion, and an Euler-operator formula for every block of $g(y)^r$. The positive real branch is selected explicitly. An implicit-function argument in finitely many small parameters proves convergence, a Rouché argument supplies a conditional numerical majorant, and exponent cutoffs have remainders governed by the first omitted block. Forward errors and residuals are transported with the derivative scaling appropriate to the observable, including powers of inverses at infinity. A finite-smoothness theorem also covers perturbations outside the analytic power-log class. A logarithmic coordinate extends the method to Lambert cores, including the real inverse of $x \log x$ near 0 and of $x e^x$ at infinity, with a convergent expansion and explicit remainder scale. Further implemented algebras cover finite logarithmic hierarchies, exact-core perturbations, flat exponential sectors and Fourier coefficients. Special-function adapters distinguish Poincaré forward models from exact local threshold equations. The package provides composable results, retained-state refinement, exact rational residual certificates and independent generated validation. We distinguish implemented computations, hypotheses carried by the input data, and broader scales still outside the admitted classes, and compare the nine reports underlying the original implementation.

Contents

1	Introduction	5
1.1	Two questions	5
1.2	What is common to all of this	6
1.3	Summary of results	6
1.4	Conventions	7

*Written with Claude (Anthropic). The theory synthesises nine independently written reports, listed in Section 14, and adds the extensions needed for the unified package.

2	The power–log scale and its exact arithmetic	8
2.1	Blocks and their order	8
2.2	Exponent semigroups	8
2.3	Jets and their arithmetic	9
2.4	Blocks with nonconstant leading coefficient	10
3	Forward expansions	10
3.1	Convergent power–log expansions	10
3.2	The precision-tracking evaluator	12
3.3	The second question	14
4	The model and the real branch	14
4.1	The normal form	14
4.2	Existence and uniqueness of the local inverse	14
4.3	The complex germ returned by <code>InverseFunction</code>	15
5	Lagrange–Bürmann inversion with a marker	15
5.1	The marker theorem	16
5.2	Consequences	17
6	The coefficient theorem	17
6.1	Statement	17
6.2	First orders and cross terms	18
6.3	Degrees and leading logarithms	19
6.4	Constant coefficients: Fuss–Catalan numbers	19
6.5	Formal existence, uniqueness, and the two iterative engines	20
7	Convergence	21
7.1	Convergence in finitely many small parameters	21
7.2	An explicit conditional majorant	22
8	Remainders after a cutoff	23
8.1	Exponent cutoff: the first omitted block	23
8.2	Depth truncation	24
8.3	Numerical illustration	25
9	Transport of remainders and finite smoothness	25
9.1	From the forward function to the inverse	25
9.2	From a residual to the error of the inverse	26
9.3	Finite smoothness suffices	27
10	Finite endpoints, infinity, and powers of the inverse	28
10.1	Local variables	29
10.2	The theorem at an arbitrary endpoint	29

10.3 Powers of the inverse and the original variable	29
10.4 Examples	30
11 Worked examples	30
11.1 The two examples of the first question, once more	31
11.2 Nonunit leading power	31
11.3 Several gaps	31
11.4 Automatically expanded forward functions	32
11.5 Symbolic data	32
11.6 Powers and logarithms of the inverse	32
12 Lambert coordinates and boundaries of the theory	32
12.1 A leading logarithm: Lambert W and the log log scale	32
12.2 A convergent universal Lambert expansion	33
12.3 Exponential cores at infinity	34
12.4 Perturbations of a Lambert core	34
12.5 Zero gaps: corrections of relative size $1/\log$	36
12.6 Flat terms	36
12.7 Oscillatory coefficients	36
12.8 Symbolic exponents and nonuniformity	36
12.9 Nested powers of the logarithm	37
13 The package	37
13.1 Data structures and exact comparisons	38
13.2 The engines	38
13.3 Checks	39
13.4 Cost	39
14 The nine reports	40
14.1 Agreement	40
14.2 Individual contributions absorbed here	40
14.3 Errors found	40
14.4 Native execution	41
15 Inversion through explicit coordinate transformations	41
15.1 Target transformations	41
15.2 Exponential phases	41
15.3 Source reconstruction and its error	43
15.4 Pure logarithmic functions	43
15.5 Finite sums of exponentials	44
15.6 Result contracts and verification	44
16 Incremental and grouped inverse computation	45

16.1	A persistent exact-weight frontier	45
16.2	Newton iteration with growing precision	45
16.3	Combine resonances before applying Euler operators	46
16.4	Validation obligations	46
16.5	Recorded native measurements	47
17	Calculus of expansions with transported precision	47
18	Perturbations of a retained exact inverse	49
18.1	The coefficient identity	49
18.2	A class with an asymptotic tail bound	50
18.3	The Lambert core and its first correction	51
18.4	A divergent power plus a logarithm	52
18.5	A declared forward remainder	52
18.6	The implemented operation and its evidence boundary	53
19	Executable residual certificates	54
19.1	The certificate theorem with interval data	54
19.2	Rational arithmetic and elementary enclosures	54
19.3	Interface, phase coordinates, and adaptive accuracy	55
19.4	Validation scope	57
20	Implemented source charts and reconstruction checks	57
21	Finite logarithmic hierarchies and reciprocal-logarithmic order	59
21.1	The finite coefficient hierarchy	59
21.2	Positive source-power gaps with logarithmic coefficients	60
21.3	Reciprocal-logarithmic units	61
21.4	A leading logarithmic monomial	62
21.5	Implementation contracts and checks	63
22	Finite commensurate flat sectors	64
22.1	Computing complete sectors	64
22.2	A bound for the complete omitted tail	65
23	Finite Fourier coefficients in a logarithmic variable	66
23.1	Coefficient arithmetic and reality	66
23.2	Inverse coefficients and independent composition	66
23.3	Envelopes, convergence, and omitted terms	67
23.4	Interface and acceptance examples	67
24	Special functions, asymptotic forward models, and real thresholds	68
24.1	The complementary-error-function tail	68
24.2	A Stirling model with a retained Lambert inverse	69

24.3	The two real Lambert branches at their common threshold	70
24.4	An exact coalescing quadratic branch	71
24.5	API, numerical evidence, and switching policy	71
25	Refinement by retaining complete coefficient states	72
26	Independent generated validation and reproducible failures	74
27	Retaining an exact growing exponential core	76
27.1	The admitted core and source charts	76
27.2	Complete coefficients without expanding the core inverse	77
27.3	A uniform complete-tail theorem	77
27.4	An unknown input error remains a separate first sector	79
27.5	Implemented behavior and verification	79
28	Separate inner precision and calculus of flat sectors	80
28.1	Inner truncation and multiplication	80
28.2	Qualified differentiation	81
29	Integrated refinement and numerical evidence	82
29.1	Three meanings of refinement	82
29.2	Checking observables against the original equation	82
30	Composition and differentiation in a reciprocal logarithm	83
30.1	A common positive logarithmic coordinate	83
30.2	Composition of two monomial carriers	84
30.3	Differentiating the analytic remainder	85
30.4	Provenance, refinement, and evidence	86

1 Introduction

1.1 Two questions

The first question [1] asks for an asymptotic expansion, as $y \rightarrow 0^+$, of the inverse of

$$f_1(x) = x + x^2(1 + \log x), \quad x > 0, \quad (1.1)$$

on the branch that maps $(0, \infty)$ to itself. The natural attempts fail in instructive ways. `AsymptoticSolve[x + x^2 (1 + Log[x]) == y, {x}, y -> 0]` returns unevaluated; `InverseSeries` returns unevaluated because the forward series has a coefficient containing $\log x$; and

```
Series[InverseFunction[Function[x, x + x^2 (1 + Log[x])]][y], {y, 0, 2}]
```

returns the perfectly correct Taylor germ

$$c + \frac{y}{c-1} - \frac{3 + 2W_{-1}(-e)}{2(c-1)^3} y^2 + \mathcal{O}(y^3), \quad c = e^{W_{-1}(-e)-1} \approx -0.1178 - 0.5332i, \quad (1.2)$$

of the inverse at a *different* preimage of 0: the nonzero complex root c of $1 + x(1 + \log x) = 0$, where f_1 is analytic and $f_1'(c) = c - 1 \neq 0$ (Section 4). The real branch has no Taylor expansion at all. Its expansion is

$$g_1(y) = y - y^2(1 + L) + y^3(2L^2 + 5L + 3) - y^4\left(5L^3 + \frac{41}{2}L^2 + 27L + \frac{23}{2}\right) + y^5\left(14L^4 + \frac{241}{3}L^3 + \frac{335}{2}L^2 + 151L + \frac{299}{6}\right) + \mathcal{O}(y^6(1 + |L|)^5), \quad (1.3)$$

with $L = \log y$: every power of y carries a polynomial in $\log y$, of degree growing with the power, so the scale of the expansion is neither a Taylor nor a Puiseux scale. The question also asks, as a second example, for

$$g_2(y) = y - y^{\sqrt{2}} + \sqrt{2}y^{2\sqrt{2}-1} - \frac{6 - \sqrt{2}}{2}y^{3\sqrt{2}-2} + \left(\frac{17\sqrt{2}}{3} - 4\right)y^{4\sqrt{2}-3} + \mathcal{O}(y^{5\sqrt{2}-4}), \quad (1.4)$$

the inverse of $f_2(x) = x + x^{\sqrt{2}}$: no logarithms this time, but exponents in the semigroup $1 + (\sqrt{2} - 1)\mathbb{N}$, which no rational grid contains.

The second question [2] asks for the forward asymptotics of $(1 + x + x^{\sqrt{2}})^{\sqrt{2}}$ as $x \rightarrow \infty$. The built-in `Asymptotic` correctly returns its leading term x^2 , but adding `SeriesTermGoal -> 7` returns the input unchanged. The expansion is

$$(1 + x + x^{\sqrt{2}})^{\sqrt{2}} = x^2 + \sqrt{2}x^{3-\sqrt{2}} + \left(1 - \frac{1}{\sqrt{2}}\right)x^{4-2\sqrt{2}} + \left(\frac{2\sqrt{2}}{3} - 1\right)x^{5-3\sqrt{2}} + \sqrt{2}x^{2-\sqrt{2}} + \frac{13-9\sqrt{2}}{12}x^{6-4\sqrt{2}} + (2 - \sqrt{2})x^{3-2\sqrt{2}} + \mathcal{O}(x^{7-5\sqrt{2}}), \quad (1.5)$$

and its inverse, as $y \rightarrow \infty$, begins $\sqrt{y} - \frac{1}{\sqrt{2}}y^{(2-\sqrt{2})/2} + \frac{3-\sqrt{2}}{4}y^{(3-2\sqrt{2})/2} + \mathcal{O}(y^{(4-3\sqrt{2})/2})$.

1.2 What is common to all of this

All these expansions are finite or infinite sums of *power-log blocks* $w^\beta P(\log w)$, where $w \rightarrow 0^+$ is a local variable ($w = y$, $w = 1/x$, $\dots$), β is a real exponent, and P is a polynomial. Their supports lie in translates of semigroups generated by finitely many positive gaps. These semigroups are locally finite, even when the groups generated by the same gaps are dense in $\mathbb{R}$, so every truncation is a finite computation. The logarithmic polynomials are handled by the Euler identity $\frac{d}{dw}[w^\beta P(\log w)] = w^{\beta-1}(\beta + \frac{d}{dL})P(L)$; and remainders must keep their logarithmic factor, because $y^3(2\log^2 y + 5\log y + 3)$ is not $\mathcal{O}(y^3)$.

This article gives the complete theory of this algebra as far as it is needed to compute and to justify such expansions, forward and inverse, with proofs, and describes the package `AsymptoticInverse` that implements it. It grew out of nine independently written reports on the first question (Section 14 and [12]), which agree on the core of the theory and each contribute something the others lack; here their contents are unified, the proofs are completed, and the theory is extended to finite endpoints from either side, to infinity, to powers and logarithms of the inverse, and to the forward expansion problem of the second question.

1.3 Summary of results

- **Scale** (Section 2). Power-log blocks are totally ordered by size, the finitely generated exponent semigroups are locally finite, and the truncated arithmetic (products, real powers, logarithms, exponentials, composition with polynomials in the logarithm) is exact and terminates. The single primitive is the composition recurrence $Q_{k+1} = ((a - k)Q_k + Q'_k)/(k + 1)$ for $(1 + t)^a P(L + \log(1 + t))$.

- **Forward expansions** (Section 3). Expressions built from constants, the variable, arithmetic, real powers, logarithms, bounded exponentials and analytic functions of expressions with a limit have convergent power–log expansions; a precision-tracking evaluator produces them with a rigorous remainder class, which answers the second question.
- **Branch** (Section 4). For $f(x) = ax^p(1 + \sum_i x^{\delta_i} B_i(\log x))$ with $a, p \neq 0$, $\delta_i > 0$, the real branch is the one with $x > 0$ and $x/(y/a)^{1/p} \rightarrow 1$; it is unique, and the complex germ (1.2) is explained.
- **Lagrange–Bürmann with a marker** (Section 5). A residue proof of $\Phi(u) = \Phi(q) + \sum_N \frac{(-\eta)^N}{N!} \partial_q^{N-1} [\Phi'(q)h(q)^N]$ at a positive point, valid analytically and formally, with the transport of observables and the general-core form.
- **Coefficients** (Section 6). For every multi-index $\mathbf{k}$ of weight σ , the block of $g(y)^r$ at $z^{r+\sigma}$, $z = (y/a)^{1/p}$, is $\frac{(-1)^{|\mathbf{k}|} r}{p^{|\mathbf{k}|} \mathbf{k}!} \prod_{j=1}^{|\mathbf{k}|-1} (\mathcal{D} + r + \sigma + pj) \prod_i B_i(L)^{k_i}$, $\mathcal{D} = d/dL$; equal weights are added. Degrees, leading logarithms (Catalan numbers), closed forms for constant coefficients (Fuss–Catalan), $\log g$, and the formal uniqueness through the triangular equation, Picard iteration and Newton iteration.
- **Convergence** (Section 7). The expansion converges absolutely for small positive y after substituting finitely many small parameters $z^{\delta_i}(\log z)^j$, with coefficient bounds $|C_{\mathbf{k}}(L)| \leq MK^{|\mathbf{k}|}(1 + |L|)^{D(\mathbf{k})}$; an explicit conditional tail majorant follows from Rouché’s theorem.
- **Remainders** (Section 8). Truncation below any exponent cutoff leaves $z^{1+\sigma_*} C_{\sigma_*}(L) + o(z^{1+\sigma_*})$ where σ_* is the first omitted weight and C_{σ_*} the complete coefficient there, computed from the one-step boundary of the retained multi-index set; depth truncation has its own bound.
- **Transport** (Section 9). Forward remainders and residuals are transported to the inverse with the correct scaling z^{1-p} ; a residual bracket gives existence and a two-sided error bound; and a theorem with a complete proof shows that finite smoothness of the perturbation, with derivative bounds in the power–log scale, suffices for the marker formula with an explicit remainder, which covers perturbations such as $x^2 \sin(\log x)$.
- **Endpoints** (Section 10). Finite points from either side and $\pm\infty$ reduce to the same engine through a local variable and a signed leading power $p \neq 0$; the limit y_0 may be finite or infinite; powers $(x - x_0)^r$ and x^r of the inverse are obtained directly. Worked examples include the Wright omega function and the inverse of (1.5).
- **Lambert coordinates and boundaries** (Section 12). Leading logarithmic and exponential cores reduce to a convergent Lambert expansion in a reciprocal-logarithm coordinate, with a polynomial in its logarithm at every order. This covers $x \log x$ near 0 and xe^x at infinity. Zero gaps, flat perturbations and oscillatory coefficients require further scale choices; symbolic exponents can use depth truncation when their branch assumptions are proved.
- **Package and reports** (Sections 13, 14). Data structures, exact comparison of algebraic exponents, the two engines, precision rules, checks; and the comparison of the nine reports, whose Wolfram packages were run natively for the first time.

1.4 Conventions

$\mathbb{N} = \{0, 1, 2, \dots\}$. For a local variable $w \rightarrow 0^+$ we write $L = \log w$ and $M(w) = 1 + |\log w|$. $\mathcal{O}$ and o refer to $w \rightarrow 0^+$ with all other data fixed; no uniformity in parameters is asserted unless stated. $\mathcal{D} = d/dL$ acts on polynomials in L . For a multi-index $\mathbf{k} = (k_1, \dots, k_m) \in \mathbb{N}^m$ we write

$|\mathbf{k}| = \sum k_i$, $\mathbf{k}! = \prod k_i!$, and $\mathbf{k} \cdot \boldsymbol{\delta} = \sum k_i \delta_i$. Real powers of positive quantities are $w^\beta = e^{\beta \log w}$; complex powers are never used on the real branch. Cutoffs are *exclusive*: a cutoff q retains the blocks of exponent $< q$. Wolfram Language names are set in typewriter font.

2 The power–log scale and its exact arithmetic

2.1 Blocks and their order

Throughout, w denotes a positive real variable tending to 0^+ , $L = \log w \rightarrow -\infty$ and $M(w) = 1 + |\log w|$.

Definition 2.1 (Blocks). A *power–log block* is an expression $w^\beta P(L)$ with $\beta \in \mathbb{R}$ and $P \in \mathbb{R}[L]$, $P \neq 0$. Its *exponent* is β and its *logarithmic degree* is $\deg P$.

Lemma 2.2 (Power beats logarithm). *For every $\varepsilon > 0$ and every real b , $w^\varepsilon M(w)^b \rightarrow 0$ as $w \rightarrow 0^+$. Consequently, if $\beta < \gamma$ and $P, Q \neq 0$, then $w^\gamma Q(L) = o(w^\beta P(L))$; and within a fixed exponent, a higher logarithmic degree dominates: $w^\beta L^j = o(w^\beta L^k)$ for $j < k$.*

Proof. Put $w = e^{-t}$, $t \rightarrow +\infty$: $w^\varepsilon M(w)^b = e^{-\varepsilon t} (1+t)^b \rightarrow 0$. For the second claim the ratio is $w^{\gamma-\beta} Q(L)/P(L)$, and $|Q(L)/P(L)| \leq CM(w)^{\deg Q}$ for w small because $P(L)$ has no zeros for L sufficiently negative and $|P(L)| \geq c|L|^{\deg P} \geq c$ there; the first claim finishes. The third claim is clear. $\square$

The growth classes of nonzero blocks are totally ordered by increasing exponent and, at equal exponents, by decreasing logarithmic degree; blocks of the same exponent and degree have the same growth class. A finite sum of blocks with distinct exponents is asymptotic to its block of least exponent (whole polynomial included), and vanishes identically only if every block does. This is why truncations are taken by *complete blocks*: cutting a block inside its polynomial would drop terms larger than the retained terms of the next exponent.

Definition 2.3 (Remainder classes). For $\beta \in \mathbb{R}$ and an integer $k \geq 0$ we write $R(w) = \mathcal{O}_{\beta,k}$, or in the package `PowerLogRemainder[w,  $\beta$ , k]`, for $R(w) = \mathcal{O}(w^\beta M(w)^k)$ as $w \rightarrow 0^+$.

By Lemma 2.2, $\mathcal{O}_{\beta,k} \subseteq \mathcal{O}_{\beta',k'}$ whenever $\beta > \beta'$, or $\beta = \beta'$ and $k \leq k'$; the class with the smaller exponent, and at equal exponents the larger degree, is the coarser one. A block $w^\beta P(L)$ with $\deg P = k$ belongs to $\mathcal{O}_{\beta,k}$ and to no $\mathcal{O}_{\beta,k'}$ with $k' < k$. Sums and products of remainder classes obey

$$\mathcal{O}_{\beta,k} + \mathcal{O}_{\beta',k'} = \mathcal{O}_{\min(\beta,\beta'),k''}, \quad \mathcal{O}_{\beta,k} \cdot \mathcal{O}_{\beta',k'} = \mathcal{O}_{\beta+\beta',k+k'}, \quad (2.1)$$

where k'' is the degree attached to the smaller exponent, and the larger of the two when the exponents coincide.

2.2 Exponent semigroups

Lemma 2.4 (Local finiteness). *Let $\delta_1, \dots, \delta_m > 0$ and $S = \{\mathbf{k} \cdot \boldsymbol{\delta} : \mathbf{k} \in \mathbb{N}^m\}$. For every real H the set $\{\mathbf{k} \in \mathbb{N}^m : \mathbf{k} \cdot \boldsymbol{\delta} \leq H\}$ is finite, hence $S \cap (-\infty, H]$ is finite, and every $\sigma \in S$ has only finitely many representations $\sigma = \mathbf{k} \cdot \boldsymbol{\delta}$.*

Proof. If $\mathbf{k} \cdot \boldsymbol{\delta} \leq H$ then $k_i \leq H/\delta_i$ for each i . $\square$

When the gaps are incommensurable the *group* generated by them is dense in $\mathbb{R}$ (for instance $\mathbb{Z} + \sqrt{2}\mathbb{Z}$), but the semigroup of nonnegative combinations has only finitely many elements below any bound. This is the whole reason why expansions with irrational exponents can be

computed term by term: the support of every expansion in this article lies in a translate $\beta_0 + S$ of such a semigroup. Following [11], a subset of $\mathbb{R}$ that is finite below every bound is called *left-finite*; the semigroups S are left-finite, and so are their translates and finite unions.

2.3 Jets and their arithmetic

Definition 2.5 (Jets). Fix gaps as above and $\beta_0 \in \mathbb{R}$. A *jet* (of cutoff H) is a finite sum $U = \sum_{\sigma \in T} w^\sigma C_\sigma(L)$ with $T \subset \beta_0 + S$ finite, $C_\sigma \in \mathbb{R}[L]$, all $\sigma < H$; equal exponents are merged and zero polynomials are discarded. Its *valuation* $\text{val } U$ is its least exponent ($+\infty$ for the empty jet). For nonnegative exponents and $H > 0$, the jet algebra $\mathcal{P}_{<H}$ uses addition and multiplication followed by truncation below H .

In that nonnegative-exponent ring, blocks of exponent $\geq H$ form an ideal and $\mathcal{P}_{<H}$ is the corresponding quotient. This quotient description does not apply to arbitrary real exponents: multiplication by a negative power can move a discarded term below the cutoff. General shifted jets therefore require precision tracking; a factor of valuation $\gamma < 0$ lowers an unknown exponent cutoff by $|\gamma|$ (Section 3.2). In the package a jet is a sorted list of pairs $\{\sigma, C_\sigma\}$ with σ exact and C_σ a polynomial in a private symbol ℓ . Algebraic exponents have canonical **RootReduce** forms; more general exact exponents require provable equality and order comparisons (Section 13).

Differentiation

Lemma 2.6 (Euler identity). For $\beta \in \mathbb{R}$, $P \in \mathbb{R}[L]$ and $r \in \mathbb{N}$,

$$\frac{d}{dw}[w^\beta P(L)] = w^{\beta-1}(\beta + \mathcal{D})P(L), \quad \frac{d^r}{dw^r}[w^\beta P(L)] = w^{\beta-r} \prod_{j=0}^{r-1} (\beta - j + \mathcal{D})P(L). \quad (2.2)$$

The operators $\beta - j + \mathcal{D}$ commute, and $(\beta + \mathcal{D})$ preserves $\deg P$ when $\beta \neq 0$ and lowers it by one when $\beta = 0$.

Proof. $\frac{d}{dw}[w^\beta P(\log w)] = \beta w^{\beta-1}P + w^\beta P'(\log w)/w$. Iterate. $\square$

The identity says that differentiation acts on a block as the scalar β plus $\mathcal{D}$ on the polynomial; the exponent bookkeeping and the polynomial arithmetic separate completely. In the coefficient formulas below operators $\mathcal{D} + c$ appear with real c (not necessarily positive), and their action on the coefficient vector $(z_0, \dots, z_d)$ of $P = \sum z_s L^s$ is bidiagonal: $(c + \mathcal{D})P = \sum_s (cz_s + (s + 1)z_{s+1})L^s$.

Unit series

Let U be a jet with $\text{val } U > 0$. Then U^k has valuation $\geq k \text{val } U$, so in $\mathcal{P}_{<H}$ only finitely many powers of U are nonzero, and every power series $\Phi(t) = \sum_k \phi_k t^k$ can be substituted:

$$\Phi(U) = \sum_{k < H/\text{val } U} \phi_k U^k \in \mathcal{P}_{<H}. \quad (2.3)$$

In particular $(1 + U)^a = \sum_k \binom{a}{k} U^k$ for every real a (with $\binom{a}{k} = a(a-1)\cdots(a-k+1)/k!$), $\log(1 + U) = \sum_{k \geq 1} (-1)^{k+1} U^k/k$, $e^U = \sum_k U^k/k!$, and $\phi(c + U) = \sum_k \phi^{(k)}(c) U^k/k!$ for any ϕ analytic at c . These are identities between finite sums of blocks modulo exponents $\geq H$; their analytic content is discussed in Section 7.

Composition with the logarithm

The one operation that mixes the exponent bookkeeping with the polynomial coefficients is the substitution of a jet into the logarithm: if w is replaced by $w(1 + U)$ then L becomes $L + \log(1 + U)$, and a block $w^a P(L)$ becomes $w^a (1 + U)^a P(L + \log(1 + U))$. The following recurrence computes this without expanding $\log(1 + U)$ separately.

Lemma 2.7 (Composition recurrence). *Let $P \in \mathbb{R}[L]$ and $a \in \mathbb{R}$. For $|t| < 1$,*

$$(1 + t)^a P(L + \log(1 + t)) = \sum_{k \geq 0} Q_k(L) t^k, \quad Q_0 = P, \quad Q_{k+1} = \frac{(a - k)Q_k + Q'_k}{k + 1}, \quad (2.4)$$

with $Q_k \in \mathbb{R}[L]$, $\deg Q_k \leq \deg P$, and, when $a \in \mathbb{N}$ and P is constant, $Q_k = 0$ for $k > a$. Consequently, for a jet U of positive valuation, $(1 + U)^a P(L + \log(1 + U)) = \sum_k Q_k(L) U^k$ in $\mathcal{P}_{<H}$.

Proof. Put $F(t) = (1 + t)^a P(L + \log(1 + t))$; it is analytic in $|t| < 1$ with values in $\mathbb{R}[L]$, so $F = \sum Q_k t^k$ with $Q_k \in \mathbb{R}[L]$ and $Q_0 = P$. Differentiating, $(1 + t)\partial_t F = aF + \partial_L F$, because ∂_t acting on the two factors gives $a(1 + t)^{a-1}P + (1 + t)^{a-1}P'$. Comparing coefficients of t^k in $(1 + t)\sum (k + 1)Q_{k+1}t^k = \sum (aQ_k + Q'_k)t^k$ gives $(k + 1)Q_{k+1} + kQ_k = aQ_k + Q'_k$. The degree does not increase, and for $a \in \mathbb{N}$ and constant P the factor $(a - k)$ annihilates Q_{a+1} and hence all later Q_k (for nonconstant P the derivative keeps the sequence alive: $\log(1 + t)$ is not a polynomial). The jet statement follows because substitution of U for t is a ring homomorphism on finite sums. $\square$

Remark 2.8. The recurrence is the whole composition calculus needed here: the forward model $ax^p(1 + \sum x^{\delta_i} B_i(\log x))$ evaluated at $x = z(1 + U)$ is $az^p(1 + U)^p(1 + \sum_i z^{\delta_i}(1 + U)^{\delta_i} B_i(L + \log(1 + U)))$, each factor of which is an instance of (2.4) (with $P = 1$ for the first). Its derivative is obtained from the same lemma with a replaced by $a - 1$ and P by $aP + P'$, since $\frac{d}{dx}[x^a P(\log x)] = x^{a-1}(aP + P')(\log x)$.

2.4 Blocks with nonconstant leading coefficient

Everything above allows the polynomial in the leading block to be nonconstant. In general, a real non-integer power or a logarithm of such a block produces $P(L)^a$ or $\log P(L)$ outside $\mathbb{R}[L]$. Special algebraic simplifications may stay in the class, but generic expressions require larger scales (rational functions of L , or iterated logarithms, Section 12). Nonnegative integer powers of any jet stay in the class.

3 Forward expansions

Before inverting anything we need the forward function in the normal form $ax^p(1 + \sum_i x^{\delta_i} B_i(\log x))$, and the second question [2] is a forward problem in its own right. This section shows that a large class of expressions has convergent power-log expansions, and describes the precision-tracking evaluator that computes them with a rigorous remainder class.

3.1 Convergent power-log expansions

Definition 3.1. A function F on $(0, w_0)$ has a *convergent power-log expansion* with support in $\beta_0 + S$ (S a finitely generated positive semigroup with gaps $\delta_1, \dots, \delta_m$) if there are polynomials

$C_\sigma \in \mathbb{R}[L]$, constants $M, K \geq 1$, and integers $d_0, d \geq 0$ such that

$$F(w) = \sum_{\sigma \in S} w^{\beta_0 + \sigma} C_\sigma(L) \quad \text{for } 0 < w < w_1, \quad \|C_\sigma\|_1 \leq MK^{n(\sigma)}, \quad \deg C_\sigma \leq d_0 + dn(\sigma), \quad (3.1)$$

where $n(\sigma)$ is the least $|\mathbf{k}|$ with $\mathbf{k} \cdot \boldsymbol{\delta} = \sigma$, $\|P\|_1$ is the sum of the absolute values of the coefficients of P , and the series converges absolutely. Thus $|C_\sigma(L)| \leq MK^{n(\sigma)} M(w)^{d_0 + dn(\sigma)}$. Zero coefficients satisfy the degree bound by convention. We write $F \in \mathcal{E}$. The fixed allowance d_0 is essential: it admits a polynomial in $\log w$ at the leading exponent, including $\log w$ itself.

The coefficient bound makes the tail after any cutoff controllable (Lemma 3.2 below), and it is stable under all the operations we need.

Lemma 3.2 (Tail of a convergent expansion). *Let $F \in \mathcal{E}$ as in (3.1), $\delta_* = \min_i \delta_i$, and $H > 0$. Then*

$$F(w) - \sum_{\sigma < H} w^{\beta_0 + \sigma} C_\sigma(L) = \mathcal{O}(w^{\beta_0 + H} M(w)^{d_0 + Nd}), \quad N = \lceil H/\delta_* \rceil. \quad (3.2)$$

More precisely, the omitted part equals $w^{\beta_0 + \sigma_} C_{\sigma_*}(L) + o(w^{\beta_0 + \sigma_*})$ where $\sigma_* = \min\{\sigma \in S : \sigma \geq H\}$.*

Proof. Split the omitted terms into those with $n(\sigma) \leq N - 1$ and the rest. The first group is finite (Lemma 2.4) and each term is $\mathcal{O}(w^{\beta_0 + H} M^{d_0 + (N-1)d})$. For the rest, the number of weights whose shortest representation has length n is at most m^n , and $\sigma \geq n\delta_*$, so their sum is bounded by

$$Mw^{\beta_0} M^{d_0} \sum_{n \geq N} (mKw^{\delta_*} M^d)^n = \mathcal{O}(w^{\beta_0 + N\delta_*} M^{d_0 + Nd}).$$

Since $N\delta_* \geq H$, this proves the bound. For the refined statement let σ_{**} be the next element of S after σ_* , which exists by local finiteness when S has a generator, and choose $\sigma_* < H' < \sigma_{**}$. Applying the bound at H' shows that everything beyond σ_* is $\mathcal{O}(w^{\beta_0 + H'} M^{d_0 + N'd}) = o(w^{\beta_0 + \sigma_*})$. The difference $\sigma_{**} - \sigma_*$ need not be at least δ_* . Finite support with no further terms has zero tail. $\square$

Theorem 3.3 (Closure). *The class $\mathcal{E}$ contains the constants and w and is closed under sums, products, and the following operations, with the supports and gaps indicated.*

1. *If $F \in \mathcal{E}$ has leading block cw^β with a constant $c > 0$, then $F^a \in \mathcal{E}$ for every real a and $\log F \in \mathcal{E}$; their supports are contained in $a\beta + S'$ and $\{0\} \cup S'$ respectively, with block $\log c + \beta L$ at exponent 0 in the logarithm. Here S' is a finitely generated positive semigroup containing the exponents of $F/(cw^\beta) - 1$; the proof establishes that such a semigroup exists.*
2. *If $F \in \mathcal{E}$ has no negative-exponent blocks and its exponent-0 block is $c + kL$ with $k \in \mathbb{R}$, then $e^F \in \mathcal{E}$ with leading block $e^c w^k$. The argument F may grow logarithmically.*
3. *If $F \in \mathcal{E}$ has $\beta_0 \geq 0$ and constant exponent-0 block c , and ϕ is analytic at c , then $\phi \circ F \in \mathcal{E}$.*

Proof. First note that a positive-valuation expansion U can be represented with support in a positive semigroup based at 0: if necessary, add its positive leading exponent as a generator and enlarge the constants in (3.1). Choose one shortest multi-index $\mathbf{k}$ for each resulting weight σ , put $D = d_0 + d$, and introduce

$$T_{ij} = w^{\delta_i} L^j, \quad 1 \leq i \leq m, \quad 0 \leq j \leq D.$$

For $n = |\mathbf{k}| \geq 1$, every $w^\sigma L^s$ with $s \leq d_0 + dn \leq Dn$ is a product of n such parameters: distribute the s logarithmic factors among the n generator factors, each receiving at most

D. Fix one such distribution for every monomial. The coefficient-norm bound makes the resulting series analytic in a sufficiently small polydisc, since its terms of degree n have total absolute coefficient at most $M(mK)^n$. Analytic composition with $(1+U)^a$, $\log(1+U)$, e^U , or $\phi(c+U)$ is therefore legitimate. Cauchy estimates give exponential coefficient bounds in the new parameters, and their logarithmic degree grows at most linearly in parameter degree.

Regrouping by weight preserves bounds of the form (3.1): a representation of weight σ has length at most σ/δ_* , whereas its shortest length is at least $\sigma/\delta_{\max}$; their ratio is bounded, and the number of parameter monomials of a given total degree is at most exponential. Enlarge M, K, d to absorb both factors.

For sums and products, use the union of the generator sets and, for a sum of differently shifted supports, the positive difference between their shifts as an additional generator. Coefficient norms are subadditive and submultiplicative, so the same bounds apply. If leading blocks cancel, rebasing at the next nonzero exponent still admits finitely many positive generators: the multi-indices above that exponent form an upward-closed subset of $\mathbb{N}^m$, with finitely many coordinatewise minimal elements. Subtracting the new leading exponent from the weights of those elements supplies finitely many nonnegative shifts, together with the original generators. This also justifies normalizing F by its actual leading monomial in (1). Separate the finite leading polynomial before applying the small-parameter construction; in (2) separate $c + kL$ to obtain the factor $e^c w^k$. $\square$

The theorem supplies sufficient closure conditions, and every truncation is a genuine asymptotic expansion by Lemma 3.2. They are not necessary conditions for every individual expression: for example $\sqrt{L^2} = -L$ for $L < 0$, and polynomials may be applied to an unbounded logarithm. In general, non-integer powers of a nonconstant leading polynomial, its logarithm, exponentials of negative-power arguments, and oscillatory functions of an unbounded logarithm require larger scales (Section 12).

Corollary 3.4 (Laurent and Puiseux composition). *If $U \in \mathcal{E}$ has a positive constant leading coefficient and positive valuation, and a selected real branch of $\phi(c+v)$ has a convergent Puiseux expansion $v^{n_0/d} A(v^{1/d})$ for $v \rightarrow 0^+$, with $d \geq 1$ an integer and A analytic at 0, then $\phi(c+U) \in \mathcal{E}$. The same holds from below using $\phi(c-v)$.*

Proof. Apply the closure theorem to $U^{1/d}$, the analytic function A , and $U^{n_0/d}$. Negative n_0 includes poles. The side of approach fixes the real branch of a ramified head. $\square$

Variable powers also reduce to the closure theorem through $F^G = \exp(G \log F)$ when $F > 0$ and the resulting exponential argument satisfies its hypotheses; this includes $w^w = \exp(w \log w)$. For a nonzero leading block $w^\beta P(L)$, the eventual sign is the sign of $(-1)^{\deg P} [L^{\deg P}]P$, so absolute values reduce to multiplication by that sign whenever it is provable.

3.2 The precision-tracking evaluator

The package computes forward expansions by structural recursion over the expression tree, carrying for every subexpression a triple (T, P, D) : a jet T and the assertion

$$\text{subexpression} = T + \mathcal{O}(w^P M(w)^D), \quad P \in \mathbb{R} \cup \{+\infty\}, \quad (3.3)$$

where $P = +\infty$ means that T is exact. A global working order K_w bounds the number of terms of every unit series that is expanded. The rules, all justified by (2.1) and Lemma 3.2, are:

- constants and w are exact; **Plus**: compute $T_1 + T_2$ and truncate below $P = \min(P_1, P_2)$. Combine the input remainder degrees at P with the degree of every explicit block discarded at exponent P ;

- **Times**: if R_i denotes the input remainder, the omitted terms include R_1T_2 , R_2T_1 and R_1R_2 . Their exponent/degree pairs follow from (2.1); use the coarsest pair and truncate T_1T_2 below that exponent. Also include the degrees of explicit product blocks discarded exactly at this boundary. A nonempty jet has valuation strictly below its remainder exponent; an empty jet contributes its remainder's exponent and degree when estimating its size;
- **Power** with an exact real exponent a : nonnegative integers by binary exponentiation with precision tracked at each product (exact for exact input). Otherwise the leading block cw^β of T must have constant $c > 0$ (or $c < 0$ with integer a); with $U = T/(cw^\beta) - 1$, relative precision $P - \beta$ and relative cutoff $c_{\text{rel}} = \min(P - \beta, K_w - a\beta)$, the result is $c^a w^{a\beta} \sum_k \binom{a}{k} U^k$ with precision $a\beta + c_{\text{rel}}$. For a nonempty U , its series-truncation degree is bounded by $\lceil c_{\text{rel}} / \text{val } U \rceil \deg_{\text{max}} U$, combined with the input error degree D ; an empty U has no computed series tail;
- **Log**: the leading block must be cw^β with constant $c > 0$; the result is $\log c + \beta L + \log(1 + U)$;
- **Exp**: T must have no negative exponents and its exponent-0 block must be $c + kL$ with exact numeric k ; the result is $e^c w^k \exp(U)$;
- an analytic head ϕ (**Sin**, **ArcTan**, **BesselJ**[0, ·], ...) applied to $T = c + U$ with constant c : $\sum_k \phi^{(k)}(c) U^k / k!$ with the Taylor coefficients obtained from **Series**[$\phi[c + \mathbf{t}]$, {**t**, 0, N}];
- any other subexpression is expanded by **Series**[**expr**, {**w**, 0, **n**}, **Assumptions** → **w** > 0]; the coefficients (which may contain $\log w$) and the irrational-power monomials that **Series** leaves outside its **SeriesData** are parsed as jets, the precision is the order of the 0 term, and the logarithmic degree D is read from the first omitted block obtained by expanding one order further.

The whole expression is re-expanded with a larger K_w until the final precision reaches the requested cutoff. The result is reported as the retained blocks plus $\mathcal{O}_{\sigma_*, k_*}$, where σ_* is the exponent of the first omitted block if one lies below the final precision (its complete coefficient having been computed), and otherwise the precision pair (P, D) itself.

Proposition 3.5 (Soundness). *If every leaf and every Taylor, Laurent, Puiseux or fallback expansion supplied by **Series** satisfies its asserted branch and remainder contract, then the root satisfies (3.3), and the reported remainder class is valid.*

Proof. Induct over the expression tree. For sums and products, use (2.1) for propagated errors and add the explicit discarded blocks to the remainder; in particular, their logarithmic degrees at the precision boundary cannot be ignored. For unit series, apply Lemma 3.2 to the finite jet's positive generators and the mean value estimate $\Phi(1 + U) - \Phi(1 + U_{\text{known}}) = \mathcal{O}(U - U_{\text{known}})$ for its unknown part. Laurent and Puiseux heads are reduced to these same operations on a known signed increment. The final remainder statement follows from the refined tail bound whenever a complete first omitted block has been computed below the precision limit. $\square$

When **Series** supplies a Taylor, Laurent, Puiseux or fallback expansion, its branch and remainder assertions remain inputs to the soundness proposition. Observing the degree of one omitted block alone does not prove that an arbitrary unknown tail satisfies the resulting bound. The structural operations preserve valid assertions, but the package does not prove every analytic contract of a special function or a user-supplied remainder. If cancellation has hidden a leading increment needed by a pole or branch point, the working precision must increase before composition.

3.3 The second question

For $F(x) = (1 + x + x^{\sqrt{2}})^{\sqrt{2}}$ at $x \rightarrow +\infty$ the local variable is $w = 1/x$ and

$$F = w^{-2}(1 + w^{\sqrt{2}-1} + w^{\sqrt{2}})^{\sqrt{2}} = w^{-2} \sum_{k \geq 0} \binom{\sqrt{2}}{k} (w^{\sqrt{2}-1} + w^{\sqrt{2}})^k,$$

an instance of item (1) of Theorem 3.3 with gaps $\sqrt{2} - 1$ and $\sqrt{2}$. The exponents of x are $2 - (\sqrt{2} - 1)k_1 - \sqrt{2}k_2$; ordering the semigroup elements gives (1.5), and the first omitted exponent is $7 - 5\sqrt{2}$ (multi-index $(5, 0)$). The workaround in the question, which peels off one leading term per limit computation, produces the same seven terms; the evaluator supplies them with the remainder $\mathcal{O}(x^{7-5\sqrt{2}})$ and can continue to larger finite cutoffs within its resource budget.

4 The model and the real branch

4.1 The normal form

Let $u \rightarrow 0^+$ be the local variable (later $u = x - x_0$, $x_0 - x$, $1/x$ or $-1/x$). The central object is a finite *power-log germ*

$$f(u) = y_0 + a u^p \left(1 + \sum_{i=1}^m u^{\delta_i} B_i(\log u) \right), \quad a \neq 0, \quad p \neq 0, \quad \delta_i > 0, \quad B_i \in \mathbb{R}[L] \setminus \{0\}, \quad (4.1)$$

with $y_0 \in \mathbb{R}$ when $p > 0$ and $y_0 = 0$ when $p < 0$ (a constant term is then a correction with gap $-p$). Equal gaps are merged. The *leading block* au^p has a constant coefficient; a logarithmic factor in the leading block ($u \log u$, $u(\log u)^{-1}$) changes the problem (Section 12). Everything below is stated for this exact finite germ; functions that are only known through a truncated expansion are handled in Section 9.

The *uniformizer* is

$$z = \left(\frac{v}{a} \right)^{1/p}, \quad v = y - y_0 \quad (p > 0), \quad v = y \quad (p < 0), \quad L = \log z = \frac{\log(v/a)}{p}, \quad (4.2)$$

so that $z \rightarrow 0^+$ in every case: for $p > 0$ because $v \rightarrow 0$ with $v/a > 0$, for $p < 0$ because $v \rightarrow \pm\infty$ with the sign of a . The role of z is that the inverse satisfies $u \sim z$.

4.2 Existence and uniqueness of the local inverse

Theorem 4.1 (Positive local branch). *For the germ (4.1) there is $\varepsilon > 0$ such that f is strictly monotone on $(0, \varepsilon)$, with*

$$f(u) - y_0 \sim a u^p, \quad f'(u) \sim a p u^{p-1} \quad (u \rightarrow 0^+). \quad (4.3)$$

Its inverse g , defined on the image of $(0, \varepsilon)$, is the unique function with $g > 0$, $g \rightarrow 0$ and $f(g(y)) = y$ there, and it satisfies

$$g(y) \sim z, \quad g'(y) \sim \frac{z}{pv} \quad (z \rightarrow 0^+). \quad (4.4)$$

Proof. Write $R(u) = \sum_i u^{\delta_i} B_i(\log u)$. By Lemma 2.2, $R(u) \rightarrow 0$ and, since $uR'(u) = \sum_i u^{\delta_i} (\delta_i B_i + B'_i)(\log u)$ by Lemma 2.6, also $uR'(u) \rightarrow 0$. Hence

$$f'(u) = a p u^{p-1} \left(1 + R(u) + \frac{uR'(u)}{p} \right) = a p u^{p-1} (1 + o(1)),$$

which has the sign of ap on some $(0, \varepsilon)$; this gives monotonicity and (4.3). The inverse exists on the image, tends to 0 because f is a homeomorphism of $(0, \varepsilon)$ onto its image, and is the only solution of $f(u) = y$ in $(0, \varepsilon)$. Substituting $u = g(y)$ into $f(u) - y_0 \sim au^p$ gives $v \sim ag(y)^p$, i.e. $g \sim z$; and $g' = 1/f'(g) \sim 1/(apg^{p-1}) = g/(pag^p) \sim z/(pv)$. $\square$

The theorem is local. Whether the local branch extends to a global inverse is a separate question, answered for the two examples of [1] as follows.

Proposition 4.2 (Global invertibility of the examples). $f_1(x) = x + x^2(1 + \log x)$ and $f_2(x) = x + x^{\sqrt{2}}$ are strictly increasing bijections of $(0, \infty)$ onto itself, with

$$f'_1(x) = 1 + x(3 + 2 \log x) \geq m_1 := 1 - 2e^{-5/2} > 0.835, \quad f'_2(x) = 1 + \sqrt{2}x^{\sqrt{2}-1} > 1. \quad (4.5)$$

Their inverses g_1, g_2 are C^1 up to $y = 0$ with $g_i(0) = 0$, $g'_i(0) = 1$, but g_1 is not C^2 at 0: $g''_1(y) = -f''_1(g_1)/f'_1(g_1)^3 \sim -(5 + 2 \log y) \rightarrow +\infty$.

Proof. $x(3 + 2 \log x)$ has derivative $5 + 2 \log x$, minimum at $x = e^{-5/2}$ with value $-2e^{-5/2}$; the limits of f_1 at 0 and ∞ are 0 and ∞ . The rest is immediate. $\square$

So for both examples the branch of Theorem 4.1 is the real inverse, and the expansions (1.3)–(1.4) describe it near 0; the absence of a second-order Taylor expansion of g_1 at 0 is exactly why $\log y$ must appear at y^2 .

4.3 The complex germ returned by InverseFunction

The equation $f_1(x) = 0$, $x \neq 0$, reads $1 + x(1 + \log x) = 0$. With $t = 1 + \log x$ it becomes $te^t = -e$, so $t = W_k(-e)$ and

$$c_k = e^{W_k(-e)-1} = -\frac{1}{W_k(-e)}, \quad f'_1(c_k) = c_k - 1, \quad f''_1(c_k) = 3 + 2W_k(-e), \quad (4.6)$$

where the derivative identities follow from $c_k W_k(-e) = -1$, provided the logarithm near c_k is chosen to have value $W_k(-e) - 1$ there. This branch qualification matters: with the fixed principal logarithm used by `Log`, the relevant roots are the conjugate pair $k = 0, -1$; other k require other local logarithm branches. For each compatible choice the derivative is nonzero, and the analytic inverse function theorem gives an inverse germ through c_k . The reported `Series[InverseFunction[...][y], {y, 0, 2}]` result chooses $k = -1$:

$$c + \frac{y}{c-1} - \frac{f''_1(c)}{2f'_1(c)^3}y^2 + \dots = c + \frac{y}{c-1} - \frac{3 + 2W_{-1}(-e)}{2(c-1)^3}y^2 + \mathcal{O}(y^3),$$

which is (1.2). Nothing is wrong with this answer: it is the inverse of a different branch of f_1^{-1} , the one through a complex root, and adding `Assumptions -> y > 0` cannot repair it, because the reality of y says nothing about which preimage is intended. The real branch is characterised by $g(y) \rightarrow 0$, and that condition, not an assumption on y , is what an inversion routine must be given; the package builds the branch of Theorem 4.1 directly and never asks `InverseFunction` to choose.

5 Lagrange–Bürmann inversion with a marker

The classical Lagrange inversion theorem expands the inverse of an analytic map at a point where its derivative does not vanish. Our maps are not analytic at 0, and the inverse has no

Taylor series there; but at every *positive* point the map is analytic, and the trick is to invert a perturbation $u + \eta h(u) = q$ around the positive point q in the perturbation parameter η . The logarithms and irrational powers in h are then analytic functions of u near $q > 0$, and the derivatives that appear are computed by the Euler identity without ever evaluating anything at 0.

5.1 The marker theorem

Theorem 5.1 (Lagrange–Bürmann with a marker). *Let $q > 0$ and let h and Φ be analytic in a complex neighbourhood of q . For $|\eta|$ sufficiently small the equation*

$$u + \eta h(u) = q \tag{5.1}$$

has a unique solution $u(\eta)$ near q with $u(0) = q$, it is analytic in η , and

$$\Phi(u(\eta)) = \Phi(q) + \sum_{N \geq 1} \frac{(-\eta)^N}{N!} \frac{d^{N-1}}{dq^{N-1}} [\Phi'(q) h(q)^N]. \tag{5.2}$$

The same identity holds coefficientwise in η for arbitrary formal power series h, Φ with coefficients in any commutative $\mathbb{Q}$ -algebra, i.e. as an identity of formal power series in η whose coefficients are polynomials in finitely many derivatives of h and Φ at q .

Proof. Choose $\rho > 0$ so small that h and Φ are analytic on the closed disc $|s - q| \leq \rho$, and let $M = \max_{|s-q|=\rho} |h(s)|$. For $|\eta| < \rho/M$ we have $|\eta h(s)| < |s - q|$ on the circle $C = \{|s - q| = \rho\}$, so by Rouché’s theorem $F_\eta(s) = s - q + \eta h(s)$ has exactly one zero $u(\eta)$ in the disc, and by the argument principle

$$\Phi(u(\eta)) = \frac{1}{2\pi i} \oint_C \Phi(s) \frac{F'_\eta(s)}{F_\eta(s)} ds = \frac{1}{2\pi i} \oint_C \Phi(s) \frac{1 + \eta h'(s)}{s - q + \eta h(s)} ds.$$

On C the geometric series $\frac{1}{s - q + \eta h(s)} = \sum_{N \geq 0} \frac{(-\eta)^N h(s)^N}{(s - q)^{N+1}}$ converges uniformly, so

$$\Phi(u(\eta)) = \sum_{N \geq 0} (-\eta)^N \left[\frac{1}{2\pi i} \oint_C \frac{\Phi(s) h(s)^N}{(s - q)^{N+1}} ds + \frac{1}{2\pi i} \oint_C \frac{\eta \Phi(s) h'(s) h(s)^N}{(s - q)^{N+1}} ds \right].$$

By Cauchy’s formula the first integral is $\frac{1}{N!} \frac{d^N}{dq^N} [\Phi h^N](q)$ and the second, after shifting the index $N \rightarrow N - 1$, contributes $-\frac{1}{(N-1)!} \frac{d^{N-1}}{dq^{N-1}} [\Phi h' h^{N-1}](q)$ to the coefficient of $(-\eta)^N$ for $N \geq 1$. Since $\frac{d}{dq} [\Phi h^N] = \Phi' h^N + N \Phi h' h^{N-1}$, the coefficient of $(-\eta)^N$ is

$$\frac{1}{N!} \frac{d^{N-1}}{dq^{N-1}} [\Phi' h^N + N \Phi h' h^{N-1}] - \frac{1}{(N-1)!} \frac{d^{N-1}}{dq^{N-1}} [\Phi h' h^{N-1}] = \frac{1}{N!} \frac{d^{N-1}}{dq^{N-1}} [\Phi' h^N],$$

which is (5.2); the constant term is $\Phi(q)$. Analyticity of $u(\eta)$ follows from the implicit function theorem at $(q, 0)$, where $\partial_u F_0 = 1$. For the formal statement, both sides of (5.2) have coefficients that are universal polynomials in the Taylor coefficients of h and Φ at q ; two polynomials that agree for all analytic data agree identically (they agree on a Zariski-dense set), so the identity holds for arbitrary formal coefficients. $\square$

Remark 5.2. Other proofs exist [4, 5, 3]; the one above is chosen because it exhibits the analytic content: the solution is the unique root in a disc around q , a fact reused in Section 7 to obtain explicit constants. The formal statement is what makes the coefficient formula of Section 6 an identity, independent of any convergence question.

5.2 Consequences

Near-identity form. With $\Phi(q) = q$ and $\eta = 1$,

$$f(u) = u + h(u) \implies g(y) = y + \sum_{N \geq 1} \frac{(-1)^N}{N!} \frac{d^{N-1}}{dy^{N-1}} [h(y)^N], \quad (5.3)$$

formally; Sections 7 and 9 say when setting $\eta = 1$ is legitimate as an asymptotic statement. The formula (5.3) is a complete answer for the log example: with $h(y) = y^2(1 + \log y)$, the N -th term is $\frac{(-1)^N}{N!} \partial_y^{N-1} [y^{2N}(1 + \log y)^N]$, computed by Lemma 2.6; for $N = 2$, $\frac{1}{2} \partial_y [y^4(1 + L)^2] = y^3(2(1 + L)^2 + (1 + L)) = y^3(2L^2 + 5L + 3)$.

Several markers. If $h = \sum_i \eta_i h_i$ with independent markers, expanding h^N multinomially in (5.2) gives the coefficient of $\prod \eta_i^{k_i}$, $|\mathbf{k}| = N$, as

$$\frac{(-1)^N}{\mathbf{k}!} \frac{d^{N-1}}{dq^{N-1}} \left[\Phi'(q) \prod_i h_i(q)^{k_i} \right], \quad (5.4)$$

because the multinomial coefficient $N!/\mathbf{k}!$ cancels the $1/N!$.

Transport of observables. Formula (5.2) gives at once the expansion of any analytic function of the inverse. With $\Phi(u) = u^r$ one obtains powers of the inverse; with $\Phi = \log$ its logarithm; differentiating the identity for u^r with respect to r at $r = 0$ gives the logarithm again. These are used in Section 6.

General core. Suppose $F = F_0 + R$ where F_0 has a known inverse ϕ on the branch of interest and R is a perturbation. Put $t = F_0(x)$, so that $F(x) = y$ becomes $t + \eta R(\phi(t)) = y$ with $\eta = 1$; Theorem 5.1 with $\Phi = \phi$ yields

$$x = \phi(y) + \sum_{N \geq 1} \frac{(-1)^N}{N!} \frac{d^{N-1}}{dy^{N-1}} \left[\phi'(y) R(\phi(y))^N \right]. \quad (5.5)$$

This is the “phase-coordinate” inversion of [11]; it is implemented as `PerturbativeInverse` $[\phi, R, \{\mathbf{x}, \mathbf{y}\}, \mathbf{n}]$ and is the natural tool when the dominant part of F is not a monomial (Section 12). For the monomial core $F_0(x) = ax^p$, $\phi(y) = (y/a)^{1/p}$, it is equivalent to the substitution $t = x^p$ used in the next section.

6 The coefficient theorem

6.1 Statement

Let f be the germ (4.1), z the uniformizer (4.2), $L = \log z$, $\mathcal{D} = d/dL$. For a multi-index $\mathbf{k} \in \mathbb{N}^m$ write $n = |\mathbf{k}|$, $\sigma = \mathbf{k} \cdot \boldsymbol{\delta}$ and $B^{\mathbf{k}} = \prod_i B_i^{k_i}$.

Theorem 6.1 (Euler-operator coefficient formula). *For every real r , the branch g of Theorem 4.1 satisfies, formally and (Theorem 7.1) as a convergent expansion for small z ,*

$$g(y)^r = z^r + \sum_{\mathbf{k} \neq 0} z^{r+\mathbf{k} \cdot \boldsymbol{\delta}} C_{\mathbf{k}}^{(r)}(L), \quad C_{\mathbf{k}}^{(r)}(L) = \frac{(-1)^n r}{p^n \mathbf{k}!} \prod_{j=1}^{n-1} (\mathcal{D} + r + \sigma + pj) B^{\mathbf{k}}(L). \quad (6.1)$$

The empty product ($n = 1$) is the identity. Multi-indices of equal weight are summed: the block of g^r at $z^{r+\sigma}$ is $C_\sigma^{(r)} = \sum_{\mathbf{k} \cdot \delta = \sigma} C_{\mathbf{k}}^{(r)}$. In particular

$$\begin{aligned} g(y) &= z \left(1 + \sum_{\mathbf{k} \neq 0} z^{\mathbf{k} \cdot \delta} C_{\mathbf{k}}(L) \right), \quad C_{\mathbf{k}} = C_{\mathbf{k}}^{(1)}, \\ \log g(y) &= L + \sum_{\mathbf{k} \neq 0} z^{\mathbf{k} \cdot \delta} \frac{(-1)^n}{p^n \mathbf{k}!} \prod_{j=1}^{n-1} (\mathcal{D} + \sigma + pj) B^{\mathbf{k}}(L). \end{aligned} \quad (6.2)$$

Proof. Put $t = u^p$ and $q = v/a$ (so $q = z^p$). The equation $f(u) = y$ becomes

$$q = t + \sum_i t^{1+\delta_i/p} B_i \left(\frac{\log t}{p} \right),$$

which is (5.1) with markers: $t + \sum_i \eta_i h_i(t) = q$, $h_i(t) = t^{1+\delta_i/p} B_i(\log t/p)$, at $\eta_i = 1$. The observable is $\Phi(t) = t^{r/p} = u^r$. All h_i and Φ are analytic near any $q > 0$. By (5.4) the coefficient of $\prod \eta_i^{k_i}$ is

$$\frac{(-1)^n}{\mathbf{k}!} \frac{d^{n-1}}{dq^{n-1}} \left[\frac{r}{p} q^{r/p-1} \prod_i (q^{1+\delta_i/p} B_i(\frac{\log q}{p}))^{k_i} \right] = \frac{(-1)^n r}{p \mathbf{k}!} \frac{d^{n-1}}{dq^{n-1}} \left[q^{n-1+(r+\sigma)/p} B^{\mathbf{k}}(\frac{\log q}{p}) \right].$$

Apply the repeated Euler identity (2.2) in the variable q , with $\Lambda = \log q$ and the polynomial $\tilde{P}(\Lambda) = B^{\mathbf{k}}(\Lambda/p)$: the derivative of order $n-1$ of $q^b \tilde{P}(\Lambda)$, $b = n-1+(r+\sigma)/p$, is $q^{b-n+1} \prod_{j=0}^{n-2} (b-j + d/d\Lambda) \tilde{P}$. Since $L = \Lambda/p$, $d/d\Lambda = \mathcal{D}/p$, and $b-j = (r+\sigma+p(n-1-j))/p$; hence each factor is $\frac{1}{p}(\mathcal{D} + r + \sigma + p(n-1-j))$, and with $j' = n-1-j$ running from 1 to $n-1$ the product is $p^{-(n-1)} \prod_{j'=1}^{n-1} (\mathcal{D} + r + \sigma + pj') B^{\mathbf{k}}(L)$. Finally $q^{b-n+1} = q^{(r+\sigma)/p} = z^{r+\sigma}$, and the prefactor r/p combines with $p^{-(n-1)}$ into r/p^n . The formal validity is Theorem 5.1's formal clause, applied to the finitely many markers; setting all $\eta_i = 1$ is legitimate coefficientwise because by Lemma 2.4 only finitely many multi-indices contribute to any exponent. The formula for $\log g$ is the derivative of (6.1) with respect to r at $r = 0$, where the factor r vanishes and only the term with the factor differentiated survives; equivalently it is Theorem 5.1 with $\Phi = \log$. $\square$

Remark 6.2. The formula (6.1) is the one on which all nine reports of Section 14 agree, with the same normalization: sign $(-1)^n$, the multi-index factorial $\mathbf{k}!$ (not $n!$: the multinomial coefficient from expanding h^n has been absorbed), the factor p^{-n} (one $1/p$ from Φ' and one from each of the $n-1$ derivatives, through $d/d\Lambda = \mathcal{D}/p$), and the operator constants $r + \sigma + pj$. Forgetting the substitution $L = \log(v/a)/p$ in the coefficients, or the factor p^{-n} , are the two classical mistakes for nonunit p .

6.2 First orders and cross terms

For a single block ($m = 1$, gap δ , polynomial B):

$$g = z - \frac{z^{1+\delta}}{p} B + \frac{z^{1+2\delta}}{2p^2} ((1+2\delta+p)B^2 + 2BB') - \frac{z^{1+3\delta}}{6p^3} (\mathcal{D} + 1 + 3\delta + p)(\mathcal{D} + 1 + 3\delta + 2p)B^3 + \dots \quad (6.3)$$

For two blocks the mixed multi-index $e_i + e_j$ contributes

$$\frac{z^{1+\delta_i+\delta_j}}{p^2} (\mathcal{D} + 1 + \delta_i + \delta_j + p)(B_i B_j) = \frac{z^{1+\delta_i+\delta_j}}{p^2} ((1 + \delta_i + \delta_j + p)B_i B_j + B'_i B_j + B_i B'_j), \quad (6.4)$$

which shows why the inverse of a sum of perturbations is not the sum of the separately computed inverse corrections.

6.3 Degrees and leading logarithms

Proposition 6.3 (Degrees). *Let $d_i = \deg B_i$ and $D(\mathbf{k}) = \sum_i k_i d_i$. Then $\deg C_{\mathbf{k}}^{(r)} \leq D(\mathbf{k})$, with coefficient at the largest possible degree*

$$[L^{D(\mathbf{k})}] C_{\mathbf{k}}^{(r)} = \frac{(-1)^{nr}}{p^n \mathbf{k}!} \prod_{j=1}^{n-1} (r + \sigma + pj) \prod_i ([L^{d_i}] B_i)^{k_i}, \quad (6.5)$$

which is nonzero exactly when $r \neq 0$ and $r + \sigma + pj \neq 0$ for $1 \leq j \leq n-1$ (always true for $r > 0$, $p > 0$). Under those conditions the degree is exactly $D(\mathbf{k})$. Zero operator constants differentiate and may lower the degree or annihilate the polynomial, while $r = 0$ annihilates every correction. An aggregated block $C_{\sigma}^{(r)}$ has degree at most $\max\{D(\mathbf{k}) : \mathbf{k} \cdot \boldsymbol{\delta} = \sigma\}$; equality requires a nonzero sum of its top-degree coefficients.

Proof. Each factor $\mathcal{D} + c$ with $c \neq 0$ preserves the degree and multiplies the top coefficient by c ; a factor with $c = 0$ differentiates. The product $B^{\mathbf{k}}$ has degree $D(\mathbf{k})$, and the prefactor contains r . $\square$

For the log example ($a = p = \delta = 1$, $B = 1 + L$, single index n) the formula reads

$$\begin{aligned} P_n(L) &= \frac{(-1)^n}{n!} \prod_{j=n+2}^{2n} (\mathcal{D} + j) (1 + L)^n, \\ [L^n] P_n &= \frac{(-1)^n}{n!} \prod_{j=n+2}^{2n} j = (-1)^n \frac{(2n)!}{n! (n+1)!} = (-1)^n \text{Cat}_n, \end{aligned} \quad (6.6)$$

the signed Catalan numbers. The next coefficient is also explicit: selecting the derivative in exactly one factor of the product gives

$$[L^{n-1}] P_n = (-1)^n \text{Cat}_n \, n (1 + H_{2n} - H_{n+1}), \quad H_k = \sum_{j \leq k} 1/j, \quad (6.7)$$

because $\prod_{j=n+2}^{2n} (\mathcal{D} + j) (1 + L)^n$ has L^{n-1} -coefficient $(\prod j) (n + \sum_{j=n+2}^{2n} n/j)$ (the term n from expanding $(1 + L)^n$ and the terms n/j from applying $\mathcal{D}$ in the factor j). Every P_n is divisible by $1 + L$: in the variable $t = 1 + L$ the operator product $\prod_{j=n+2}^{2n} (\mathcal{D} + j)$ consists of $n-1$ factors, each of which either differentiates (lowering the power of t by one) or multiplies by a constant, so every monomial of $\prod (\mathcal{D} + j) t^n$ has degree at least $n - (n-1) = 1$ in t ; the package confirms the factorisation through $n = 8$. Thus

$$\begin{aligned} P_1 &= -(1 + L), & P_2 &= (1 + L)(2L + 3), & P_3 &= -\frac{1}{2}(1 + L)(10L^2 + 31L + 23), \\ P_4 &= \frac{1}{6}(1 + L)(84L^3 + 398L^2 + 607L + 299), \end{aligned}$$

which are the blocks of (1.3).

6.4 Constant coefficients: Fuss–Catalan numbers

When every B_i is a constant b_i , the operators act as scalars and

$$C_{\mathbf{k}}^{(r)} = \frac{(-1)^{nr}}{p^n \mathbf{k}!} \prod_{j=1}^{n-1} (r + \sigma + pj) \prod_i b_i^{k_i} = \frac{r}{r + \sigma} \binom{-(r + \sigma)/p}{n} \frac{n!}{\mathbf{k}!} \prod_i b_i^{k_i}. \quad (6.8)$$

The binomial expression in (6.8) is read by its removable limit when $r + \sigma = 0$; the operator-product expression has no division by $r + \sigma$ and remains directly valid. For $f(x) = x + bx^\alpha$

($p = 1$, $\delta = \alpha - 1$, $r = 1$) this is the generalised Fuss–Catalan law

$$g(y) = \sum_{n \geq 0} c_n b^n y^{1+n(\alpha-1)},$$

$$c_n = \frac{(-1)^n}{n!} \prod_{j=0}^{n-2} (n\alpha - j) = \frac{(-1)^n}{1+n(\alpha-1)} \binom{n\alpha}{n} = (-1)^n \frac{\Gamma(n\alpha+1)}{\Gamma(n+1)\Gamma(n(\alpha-1)+2)}, \quad (6.9)$$

valid for every real $\alpha > 1$ (no rationality is used), with $c_0 = 1$ understood separately in the product formula. For $\alpha = 2$ the c_n are the signed Catalan numbers and, when $b \neq 0$, $g = (\sqrt{1+4by} - 1)/(2b)$; the limit at $b = 0$ is $g = y$. For $b = 1$ and $\alpha = \sqrt{2}$,

$$c_1 = -1, \quad c_2 = \sqrt{2}, \quad c_3 = -\frac{6-\sqrt{2}}{2},$$

$$c_4 = \frac{17\sqrt{2}}{3} - 4, \quad c_5 = \frac{51\sqrt{2}}{4} - \frac{305}{12}, \quad c_6 = \frac{322\sqrt{2}}{5} - 77,$$

which gives (1.4). In this one-gap case $g(y) = yV(by^{\alpha-1})$ with V the solution of $V + tV^\alpha = 1$, $V(0) = 1$, analytic at $t = 0$ with radius of convergence

$$R_\alpha = \frac{(\alpha-1)^{\alpha-1}}{\alpha^\alpha}, \quad |c_n| \sim \frac{\sqrt{\alpha}}{\sqrt{2\pi}(\alpha-1)^{3/2}} n^{-3/2} R_\alpha^{-n}, \quad (6.10)$$

by Stirling's formula, so the expansion (1.4) converges (and equals the positive inverse, by analytic continuation of the implicit equation) for $0 < y < R_{\sqrt{2}}^{1/(\sqrt{2}-1)} \approx 0.12686$. The singularity at $t = -R_\alpha$ is the fold $V = \alpha/(\alpha-1)$ of the implicit equation on the negative axis.

6.5 Formal existence, uniqueness, and the two iterative engines

Theorem 6.1 produces the coefficients; the following independent argument shows that they are the only possible ones and provides two alternative algorithms. Write $g = z(1+U)$ with U an unknown jet of positive valuation. Then $f(g) = y$ is equivalent to

$$E(U) := (1+U)^p \left(1 + \sum_i z^{\delta_i} (1+U)^{\delta_i} B_i(L + \log(1+U)) \right) - 1 = 0, \quad (6.11)$$

an equation in the jet algebra (Lemma 2.7 evaluates every term).

Proposition 6.4 (Triangular structure). *Equation (6.11) has exactly one formal solution U of positive valuation with support in the semigroup S generated by the gaps. Its coefficient at z^σ is determined by the coefficients at smaller weights through one division by p .*

Proof. $E(U) = pU + \sum_i z^{\delta_i} B_i(L) + (\text{terms of weight } \geq \min(2 \text{ val } U, \text{ val } U + \delta_*))$, because $(1+U)^p = 1 + pU + \mathcal{O}(U^2)$ and every correction carries a factor z^{δ_i} . Order the elements of S increasingly (Lemma 2.4); at a new weight σ the coefficient of z^σ in $E(U)$ is pU_σ plus an expression in the $U_{\sigma'}$ with $\sigma' < \sigma$, which fixes U_σ . If U, V are two solutions with $U - V$ of least weight σ , then $E(U) - E(V) = (U - V)(p + \text{positive weight})$ has a nonzero block at z^σ , a contradiction. $\square$

Picard iteration. Rewriting (6.11) as $U = \Psi(U) := (1 + \sum_i z^{\delta_i} (1+U)^{\delta_i} B_i(L + \log(1+U)))^{-1/p} - 1$, the map Ψ gains one gap per step: $\text{val}(\Psi(U) - \Psi(V)) \geq \text{val}(U - V) + \delta_*$. Hence $U_{k+1} = \Psi(U_k)$, $U_0 = 0$, agrees with the solution below weight $(k+1)\delta_*$, and $\lceil H/\delta_* \rceil$ iterations give all blocks below H .

Newton iteration. With $E'(U) = p(1+U)^{p-1} + \sum_i z^{\delta_i}(1+U)^{p+\delta_i-1}((p+\delta_i)B_i + B'_i)(L + \log(1+U))$, a unit of the jet algebra with constant leading coefficient p , the iteration $U_{k+1} = U_k - E(U_k)/E'(U_k)$ doubles the valuation of the error: if $\text{val}(U_k - U) \geq s$ then $\text{val}(U_{k+1} - U) \geq 2s$, because $E(U_k) = E'(U)(U_k - U) + \mathcal{O}((U_k - U)^2)$ and $E'(U_k)^{-1} = E'(U)^{-1}(1 + \mathcal{O}(U_k - U))$. Starting from $U_0 = 0$ with $s = \delta_*$, about $\log_2(H/\delta_*)$ iterations suffice. The package offers this engine as `Method -> "Newton"`; it and the multi-index formula must agree, which is a strong internal check.

7 Convergence

A formal identity such as (6.1) says nothing, by itself, about the function g : the series might diverge, or converge to something else. This section proves that the expansion converges absolutely for small z and equals the branch of Theorem 4.1, first by an implicit-function argument that also yields the coefficient bounds needed for the remainder theory, then by a Rouché argument that yields explicit constants.

7.1 Convergence in finitely many small parameters

Theorem 7.1 (Convergence). *Let f be the germ (4.1) and g its branch. For every real r the series (6.1) converges absolutely for all sufficiently small $z > 0$, its sum is $g(y)^r$, and there are constants M, K such that*

$$|C_{\mathbf{k}}^{(r)}(L)| \leq M K^{|\mathbf{k}|} M(z)^{D(\mathbf{k})} \quad (L \leq \log z_1), \quad D(\mathbf{k}) = \sum_i k_i \deg B_i. \quad (7.1)$$

In particular $g^r \in \mathcal{E}$ (Definition 3.1) with support $r + S$.

Proof. Write $g = zw$ with $w \rightarrow 1$; by (6.11) the unknown w solves

$$w^p \left(1 + \sum_i z^{\delta_i} w^{\delta_i} B_i(L + \log w) \right) = 1.$$

Expand each $B_i(L + \log w) = \sum_{j=0}^{d_i} L^j A_{ij}(\log w)$ with $A_{ij}(v) = \sum_{s \geq j} b_{is} \binom{s}{j} v^{s-j}$, where $B_i = \sum_s b_{is} L^s$. Introduce independent complex parameters T_{ij} in place of $z^{\delta_i} L^j$ and consider

$$\mathcal{E}(w, T) = w^p \left(1 + \sum_{i,j} T_{ij} w^{\delta_i} A_{ij}(\log w) \right) - 1. \quad (7.2)$$

On the disc $|w - 1| < 1$ with the principal logarithm, $\mathcal{E}$ is analytic in (w, T) ; $\mathcal{E}(1, 0) = 0$ and $\partial_w \mathcal{E}(1, 0) = p \neq 0$. The analytic implicit function theorem gives an analytic $W(T)$ on a polydisc $|T_{ij}| < \rho$ with $W(0) = 1$ and $\mathcal{E}(W(T), T) = 0$, and W^r is analytic there too. Its Taylor series $W(T)^r = \sum_{\nu} c_{\nu} T^{\nu}$ converges absolutely on the smaller polydisc $|T_{ij}| \leq \rho/2$, with the Cauchy estimates $|c_{\nu}| \leq M(2/\rho)^{|\nu|}$.

Now specialise $T_{ij} = z^{\delta_i} L^j$. Since $z^{\delta_i} M(z)^{d_i} \rightarrow 0$, the point lies in the polydisc for small z , and $W(T(z))$ is real (all data are real and the solution is unique) and close to 1, so $zW(T(z))$ is a solution of $f(u) = y$ with $u \sim z$; by Theorem 4.1 it is $g(y)$. Group the absolutely convergent series $\sum_{\nu} c_{\nu} \prod_{i,j} (z^{\delta_i} L^j)^{\nu_{ij}}$ by $k_i = \sum_j \nu_{ij}$. There are at most $\prod_i (d_i + 1)^{k_i}$ terms for a given $\mathbf{k}$, with total logarithmic degree $\sum_{i,j} j \nu_{ij} \leq D(\mathbf{k})$. Their sum is $z^{\mathbf{k} \cdot \delta} \tilde{C}_{\mathbf{k}}(L)$, and the Cauchy bounds give both (7.1) and $\|\tilde{C}_{\mathbf{k}}\|_1 \leq M K^{|\mathbf{k}|}$ for $K = (2/\rho) \max_i (d_i + 1)$. To identify the individual multi-index coefficients even when weights coincide, retain independent markers η_i by substituting $T_{ij} = \eta_i z^{\delta_i} L^j$. Both constructions solve the same formal implicit equation in the markers, whose linearization in w is $p \neq 0$, so its solution is unique coefficientwise. Thus

$\tilde{C}_k = C_k^{(r)}$. Finally, aggregation by equal weights preserves the bounds of Definition 3.1: all representations of weight σ have length between $\sigma/\delta_{\max}$ and σ/δ_* , so the shortest representation controls their lengths up to a fixed factor; their count is at most exponential in that length. $\square$

Remark 7.2. The theorem gives convergence of a multivariate power series after the substitution $T_{ij} = z^{\delta_i}(\log z)^j$. The normalized series g^r/z^r converges uniformly after shrinking z_1 so all these parameters remain in a smaller polydisc; for $r < 0$ this does not assert absolute uniform convergence of g^r itself on $(0, z_1)$. A single-variable radius exists in special cases such as (6.10). Nothing is uniform in a, p, δ_i, B_i : for $x + x^\alpha$ as $\alpha \downarrow 1$, the radius R_α in $t = y^{\alpha-1}$ tends to 1, whereas the corresponding endpoint $R_\alpha^{1/(\alpha-1)}$ in y tends to 0.

7.2 An explicit conditional majorant

The implicit-function argument has no explicit constants. Rouché's theorem in the marker equation gives them, at the price of a hypothesis that must be checked at each y .

Theorem 7.3 (Explicit tail bound). *Consider $f(t) = t + h(t)$, $h(t) = \sum_i t^{1+\alpha_i} P_i(\log t)$ with real α_i and $P_i(L) = \sum_s p_{is} L^s$. The general germ reduces to this form by $t = u^p$, $q = v/a$, with $\alpha_i = \delta_i/p$; these are positive at $t \rightarrow 0$ when $p > 0$ and negative at $t \rightarrow \infty$ when $p < 0$. Fix $0 < r < 1$, put $c_r = -\log(1-r)$,*

$$A_r(b) = \max\{(1-r)^b, (1+r)^b\},$$

and

$$\mathcal{Q}_r(t) = \frac{1}{r} \sum_i t^{\alpha_i} A_r(1 + \alpha_i) \sum_s |p_{is}| (|\log t| + c_r)^s. \quad (7.3)$$

Let $t > 0$ satisfy $\mathcal{Q}_r(t) < 1$. Then:

1. for every complex λ with $|\lambda| < 1/\mathcal{Q}_r(t)$ the equation $u + \lambda h(u) = t$ has exactly one solution $u(\lambda)$ in the disc $|u - t| < rt$; it is analytic in λ , $u(0) = t$, and $u(\lambda)$ is real for real λ ;
2. the Lagrange terms $u_N = \frac{(-1)^N}{N!} \frac{d^{N-1}}{dt^{N-1}} [h(t)^N]$ satisfy $|u_N| \leq \frac{rt}{N} \mathcal{Q}_r(t)^N$, and $u(\lambda) = t + \sum_{N \geq 1} u_N \lambda^N$ converges absolutely for $|\lambda| < 1/\mathcal{Q}_r(t)$, in particular at $\lambda = 1$;
3. $u(1)$ is the positive solution of $f(u) = t$ in $(t(1-r), t(1+r))$. It coincides with any local branch $g(t) \sim t$ near the endpoint under consideration, once t is sufficiently close to that endpoint, and

$$\left| u(1) - t - \sum_{N=1}^{N_0} u_N \right| \leq \frac{rt}{N_0 + 1} \frac{\mathcal{Q}_r(t)^{N_0+1}}{1 - \mathcal{Q}_r(t)}. \quad (7.4)$$

Proof. On the circle $|s - t| = rt$ write $s = t(1 + \zeta)$ with $|\zeta| = r < 1$. For a real exponent b , $|s^b| = |s|^b \leq t^b A_r(b)$; using only the upper bound on $|s|$ would be incorrect for $b < 0$. The principal logarithm satisfies $|\log s| \leq |\log t| + |\log(1 + \zeta)| \leq |\log t| + c_r$. Hence $|h(s)| \leq rt \mathcal{Q}_r(t)$. Rouché's theorem gives exactly one zero, counted with multiplicity, whenever $|\lambda| \mathcal{Q}_r(t) < 1$. That zero is simple, analytic in λ , and real for real λ by conjugation and uniqueness. Cauchy's estimate in the marker formula gives

$$|u_N| \leq \frac{1}{N!} \frac{(N-1)! (rt \mathcal{Q}_r)^N}{(rt)^{N-1}} = \frac{rt}{N} \mathcal{Q}_r^N.$$

The root at $\lambda = 1$ is positive and lies in the indicated interval. Any branch with $g(t) \sim t$ eventually lies in that interval too, so the roots agree. Finally sum the bound over $N > N_0$, using $1/N \leq 1/(N_0 + 1)$. $\square$

Corollary 7.4 (Exponent cutoff). *Assume $\alpha_i > 0$ and $B \geq 0$. Let G_B retain all multi-indices of weight $\mathbf{k} \cdot \boldsymbol{\alpha} \leq B$. Every multi-index with $|\mathbf{k}| \leq N_0 := \lfloor B/\alpha_{\max} \rfloor$ is retained, so under $\mathcal{Q}_r(t) < 1$*

$$|g(t) - G_B(t)| \leq \frac{rt}{N_0 + 1} \frac{\mathcal{Q}_r(t)^{N_0+1}}{1 - \mathcal{Q}_r(t)}.$$

Proof. The omitted part of the multi-marker expansion is a subseries of $\sum_{N > N_0} u_N$ (the expansion of u_N into multi-indices is absolutely convergent by the same Cauchy estimates, since $\sum_{|\mathbf{k}|=N} \frac{\prod M_i^{k_i}}{\mathbf{k}!} = \frac{(\sum M_i)^N}{N!}$). $\square$

Example 7.5. For $f_1(x) = x + x^2(1 + \log x)$ and $r = \frac{1}{2}$: $\mathcal{Q}_{1/2}(y) = \frac{9}{2} y (1 + |\log y| + \log 2)$, so for $y \leq \frac{1}{100}$ (where $\mathcal{Q}_{1/2} \leq 0.284$) the expansion converges and the error after N_0 blocks is at most $\frac{y}{2(N_0+1)} \mathcal{Q}_{1/2}(y)^{N_0+1} / (1 - \mathcal{Q}_{1/2}(y))$; at $y = 10^{-2}$ and $N_0 = 3$ this is $1.1 \cdot 10^{-5}$ while the actual error is $1.5 \cdot 10^{-7}$. The bound is conditional (it says nothing when $\mathcal{Q}_r \geq 1$) and not tight, but it is rigorous, cheap, and turns the asymptotic statements into inequalities at specific points.

8 Remainders after a cutoff

8.1 Exponent cutoff: the first omitted block

Let the target cutoff be an exponent q of the local variable of the result, exclusive. In terms of weights, with r the observable power, a block $z^{r+\sigma} C(L)$ has exponent $(r + \sigma)/|p|$ in the local variable of y (Section 10), so the retained set is

$$I_H = \{\mathbf{k} \in \mathbb{N}^m : \mathbf{k} \cdot \boldsymbol{\delta} < H\}, \quad H = |p|q - r, \quad (8.1)$$

a finite downward-closed set (Lemma 2.4). Set $C_0^{(r)} = 1$. When $H > 0$ its *one-step boundary* is

$$\partial I_H = \{\mathbf{k} + e_i : \mathbf{k} \in I_H, 1 \leq i \leq m\} \setminus I_H, \quad (8.2)$$

also finite. When $H \leq 0$, I_H is empty and we define its boundary to be $\{0\}$, since the leading monomial is already omitted. With at least one correction gap, put

$$\sigma_* = \min\{\boldsymbol{\nu} \cdot \boldsymbol{\delta} : \boldsymbol{\nu} \in \partial I_H\}, \quad C_{\sigma_*}^{(r)} = \sum_{\mathbf{k} \cdot \boldsymbol{\delta} = \sigma_*} C_{\mathbf{k}}^{(r)}, \quad k_* = \deg C_{\sigma_*}^{(r)} \quad (k_* = 0 \text{ if } C_{\sigma_*}^{(r)} = 0). \quad (8.3)$$

Lemma 8.1. $\sigma_* = \min(S \setminus [0, H])$, and every multi-index of weight σ_* lies in ∂I_H . Thus the complete coefficient at the first omitted weight is obtained from the boundary alone.

Proof. For $H \leq 0$ the assertion is immediate, with $\sigma_* = 0$. For $H > 0$, let $\boldsymbol{\nu} \notin I_H$ have weight σ . Along a path from 0 to $\boldsymbol{\nu}$ that increases one coordinate at a time, the weights increase strictly, and the first point outside I_H is a boundary point of weight $\leq \sigma$; hence $\sigma_* \leq \sigma$, so σ_* is the least weight outside $[0, H)$. If $\sigma = \sigma_*$ the first point outside I_H has weight exactly σ_* by minimality, and no later point on the path can have the same weight, so $\boldsymbol{\nu}$ is that boundary point. $\square$

Theorem 8.2 (Remainder after an exponent cutoff). *Let $G_H^{(r)} = z^r \sum_{\mathbf{k} \in I_H} z^{\mathbf{k} \cdot \boldsymbol{\delta}} C_{\mathbf{k}}^{(r)}(L)$. Then, as $z \rightarrow 0^+$,*

$$g(y)^r - G_H^{(r)}(y) = z^{r+\sigma_*} C_{\sigma_*}^{(r)}(L) + o(z^{r+\sigma_*}) = \mathcal{O}(z^{r+\sigma_*} M(z)^{k_*}). \quad (8.4)$$

If $C_{\sigma_}^{(r)} \neq 0$ the bound is attained by its top logarithm; if $C_{\sigma_*}^{(r)} = 0$ (a resonant cancellation) the remainder is $o(z^{r+\sigma_*})$, and the true order is that of the next nonzero block, which the theorem does not identify.*

Proof. By Theorem 7.1, $g^r \in \mathcal{E}$ with the bound (7.1), and the omitted multi-indices are those of weight $\geq \sigma_*$ by Lemma 8.1; Lemma 3.2 (refined form) gives the claim. $\square$

Corollary 8.3 (In the variable of the result). *With w the local variable of y ($w = |y - y_0|$ for finite y_0 , $w = 1/|y|$ for infinite y_0),*

$$g(y)^r - G_H^{(r)}(y) = \mathcal{O}(w^{(r+\sigma_*)/|p|} M(w)^{k_*}), \quad \frac{r + \sigma_*}{|p|} \geq q,$$

since $z = (w/|a|)^{1/p}$ for $p > 0$, $z = (|a|w)^{1/|p|}$ for $p < 0$, and $M(z) \asymp M(w)$.

For a pure monomial ($m = 0$), the inverse observable is exactly z^r : the tail is zero if this term is retained and equals z^r otherwise. The package reports the pair $((r + \sigma_*)/|p|, k_*)$ as `PowerLogRemainder[w,  $\beta$ , k]` and the block $z^{r+\sigma_*} C_{\sigma_*}^{(r)}$ as `"FrontierTerm"`. If that coefficient vanishes, a sharper order requires computing further weights or proving that the remaining series terminates.

Example 8.4 (Resonances). For $f = x + x^2 + x^3$ the gaps are 1 and 2; at weight 2 the indices (2, 0) and (0, 1) contribute 2 and -1 , at weight 3 the indices (3, 0) and (1, 1) contribute -5 and $+5$, so

$$g(y) = y - y^2 + y^3 + 0 \cdot y^4 - 4y^5 + 14y^6 - 30y^7 + 33y^8 + \dots,$$

and after the cutoff 4 the reported remainder is $\mathcal{O}(y^4)$ while the truth is $\mathcal{O}(y^5)$. For $f = x + x^2 + 2x^3$ the block at y^3 cancels instead. Exact arithmetic on the exponents is what makes such cancellations detectable at all: with floating-point exponents, (2, 0) and (0, 1) might not be recognised as the same weight.

Example 8.5 (The log example). Here $m = 1$, $\delta = 1$, $B = 1 + L$, so $I_H = \{0, \dots, N\}$ for $H = N + 1$, the boundary is $\{N + 1\}$, and by (6.6) the frontier block $y^{N+2} P_{N+1}(\log y)$ has degree exactly $N + 1$ with top coefficient $(-1)^{N+1} \text{Cat}_{N+1}$. Hence

$$g_1(y) = y + \sum_{n=1}^N y^{n+1} P_n(\log y) + \mathcal{O}(y^{N+2} M(y)^{N+1}),$$

$$\frac{g_1(y) - y - \sum_{n \leq N} y^{n+1} P_n(\log y)}{y^{N+2} (\log y)^{N+1}} \rightarrow (-1)^{N+1} \text{Cat}_{N+1}, \quad (8.5)$$

and in particular after $y - y^2(1 + \log y)$ the error is $y^3(2 \log^2 y + 5 \log y + 3) + \dots$, not $\mathcal{O}(y^3)$. The ratio in (8.5) converges slowly (like $1 + c/\log y$): for $N = 1$ it is 1.89 at $y = 10^{-20}$ and 1.98 at $y = 10^{-100}$.

8.2 Depth truncation

When the exponents are symbolic (Section 12) their order may be undecidable, but the total degree $|\mathbf{k}|$ of a multi-index is always known.

Theorem 8.6 (Remainder after a depth cutoff). *Let $G_{\leq N}^{(r)} = z^r \sum_{|\mathbf{k}| \leq N} z^{\mathbf{k} \cdot \delta} C_{\mathbf{k}}^{(r)}(L)$, $\delta_* = \min_i \delta_i$ and $d_* = \max_i \deg B_i$. Then*

$$g(y)^r - G_{\leq N}^{(r)}(y) = \mathcal{O}(z^{r+(N+1)\delta_*} M(z)^{(N+1)d_*}). \quad (8.6)$$

Proof. The omitted terms have $|\mathbf{k}| \geq N + 1$; use (7.1) and the count m^n of multi-indices of degree n exactly as in the proof of Lemma 3.2. $\square$

Depth truncation retains every block of exponent below $r + (N + 1)\delta_*$, but it may also retain blocks above that bound while omitting smaller ones when several gaps are present: for $f = x + x^2 + x^4$ depth 1 gives $y - y^2 - y^4$ although the omitted depth-two term $2y^3$ dominates y^4 . The remainder (8.6) is nevertheless correct as stated.

8.3 Numerical illustration

Table 1 shows, for the two examples of [1] and for the inverse of (1.5) at infinity, the actual error $|g(y) - G(y)|$ of the truncated expansion (the exact inverse computed by `FindRoot` at 80 digits) and its ratio to the reported remainder scale $w^{\beta_*} M(w)^{k_*}$. The ratios stay bounded, and in the log example they grow slowly with N because the leading coefficient of the frontier block, Cat_{N+1} , does.

	$y = 10^{-2}$		$y = 10^{-4}$		$y = 10^{-8}$	
	error	ratio	error	ratio	error	ratio
$f_1 = x + x^2(1 + \log x)$, N corrections						
$N = 1$	$2.4 \cdot 10^{-5}$	0.77	$1.3 \cdot 10^{-10}$	1.22	$5.9 \cdot 10^{-22}$	1.56
$N = 2$	$1.8 \cdot 10^{-6}$	1.03	$2.4 \cdot 10^{-13}$	2.26	$2.5 \cdot 10^{-28}$	3.38
$N = 3$	$1.5 \cdot 10^{-7}$	1.51	$5.1 \cdot 10^{-16}$	4.69	$1.2 \cdot 10^{-34}$	8.18
$N = 5$	$1.2 \cdot 10^{-9}$	3.75	$2.7 \cdot 10^{-21}$	24.0	$3.1 \cdot 10^{-47}$	57.4
$f_2 = x + x^{\sqrt{2}}$, N corrections						
$N = 1$	$2.5 \cdot 10^{-4}$	1.14	$6.6 \cdot 10^{-8}$	1.37	$3.3 \cdot 10^{-15}$	1.41
$N = 2$	$6.0 \cdot 10^{-5}$	1.82	$2.4 \cdot 10^{-9}$	2.21	$2.6 \cdot 10^{-18}$	2.29
$N = 3$	$1.5 \cdot 10^{-5}$	3.16	$9.1 \cdot 10^{-11}$	3.86	$2.2 \cdot 10^{-21}$	4.01
$N = 5$	$1.2 \cdot 10^{-6}$	10.9	$1.5 \cdot 10^{-13}$	13.5	$1.8 \cdot 10^{-27}$	14.1
	$y = 10^4$		$y = 10^8$		$y = 10^{16}$	
	error	ratio	error	ratio	error	ratio
inverse of $(1 + x + x^{\sqrt{2}})^{\sqrt{2}}$, n blocks						
$n = 2$	0.72	0.33	1.9	0.39	9.3	0.40
$n = 4$	$9.6 \cdot 10^{-2}$	0.65	$1.5 \cdot 10^{-2}$	0.69	$3.4 \cdot 10^{-4}$	0.71
$n = 7$	$2.8 \cdot 10^{-4}$	0.039	$1.2 \cdot 10^{-6}$	0.023	$4.0 \cdot 10^{-11}$	0.015

Table 1: Recorded errors of truncated expansions and their ratio to the remainder scale. In the third group the local variable is $1/y$ and the leading inverse term is $\sqrt{y}$; for the two-block approximation the absolute error grows while its relative error decreases.

9 Transport of remainders and finite smoothness

The theory so far concerns an exact finite germ. In practice the forward function is often known only through a truncated expansion, and one wants to know how the omitted part affects the inverse; conversely, one wants to turn a small residual $f(G(y)) - y$ of an approximation G into a bound for $G - g$. Both are mean-value arguments, but the scaling factor z^{1-p} of the derivative must not be forgotten.

9.1 From the forward function to the inverse

Theorem 9.1 (Transport of a forward remainder). *Let f be the germ (4.1) with $p \neq 0$ and let $F = f + R$ on $(0, u_0)$, where*

$$R(u) = \mathcal{O}(u^\rho M(u)^k), \quad R'(u) = \mathcal{O}(u^{\rho-1} M(u)^k), \quad \rho > p, \quad k \in \mathbb{N}. \quad (9.1)$$

Then F has a local inverse G on the same branch as g , and

$$G(y) - g(y) = \mathcal{O}(z^{\rho-p+1} M(z)^k) = \mathcal{O}(w^{(\rho-p+1)/|p|} M(w)^k), \quad (9.2)$$

where $w = |y - y_0|$ for $p > 0$ and $w = 1/|y|$ for $p < 0$. For any fixed real observable power r ,

$$G(y)^r - g(y)^r = \mathcal{O}(z^{\rho-p+r} M(z)^k) = \mathcal{O}(w^{(\rho-p+r)/|p|} M(w)^k). \quad (9.3)$$

For $r = 0$ this difference is identically zero. The same conclusions hold if the derivative hypothesis is replaced by the assumption that F is continuous and strictly monotone near 0.

Proof. From (9.1) and Lemma 2.2, $F(u) - y_0 \sim au^p$ and $F'(u) \sim apu^{p-1}$, so F is strictly monotone near 0 and its inverse G satisfies $G \sim z$ (Theorem 4.1 applies verbatim). Put $u = G(y)$, $v = g(y)$; both are $\sim z$. Then $F(v) - y = R(v) = \mathcal{O}(z^p M(z)^k)$, while $F(u) - y = 0$. On the interval between u and v , which lies in $[z/2, 2z]$ for small z , F' is at least $\frac{1}{2}|ap|z^{p-1}2^{-|p-1|}$ in absolute value. The mean value theorem gives (9.2). With continuous strict monotonicity instead of the derivative bound, the leading equivalence still supplies a continuous local inverse with $G \sim z$, and $f(u) - f(v) = -R(u)$ together with $|f'| \asymp z^{p-1}$ gives the same estimate. Finally the derivative of s^r on $[z/2, 2z]$ is $\mathcal{O}(z^{r-1})$, giving (9.3). $\square$

Remark 9.2. The derivative hypothesis is not a consequence of the value bound: $R(u) = u^2 \sin(u^{-2})$ is $\mathcal{O}(u^2)$, but $R'(u) = 2u \sin(u^{-2}) - 2u^{-1} \cos(u^{-2})$ destroys the monotonicity of $u + R(u)$. Nor does non-strict monotonicity alone ensure an inverse, since it permits plateaus. Expansions obtained by the convergent structural operations above admit corresponding derivative estimates, but a formal **Series** result is not itself a proof of them. A declared "**InputRemainder**" $\rightarrow \{\rho, k\}$ carries the derivative or continuous strict monotonicity hypothesis of this theorem.

Corollary 9.3 (Cutoff cap). *The data justify every block of G^r strictly below the exponent $(\rho - p + r)/|p|$ in w . Thus a requested cutoff q is certified from those data when $q \leq (\rho - p + r)/|p|$; the reported remainder is the coarser of the truncation remainder and (9.3), compared by exponent first and by logarithmic degree at equal exponents. A larger requested cutoff requires a sharper forward remainder or additional information. For powers of the original variable at an infinite source endpoint, replace r here by the local-coordinate power $-r$.*

For example $\sin x + x^2 \log x = x + x^2 \log x - x^3/6 + \mathcal{O}(x^5)$ is a germ with $p = 1$, gaps 1, 2 and $\rho = 5$, $k = 0$; the inverse can be requested up to the cutoff $q = 5$ and is $y - y^2 \log y + y^3(\frac{1}{6} + \log y + 2 \log^2 y) + \mathcal{O}(y^4 M^3)$ for $q = 4$.

9.2 From a residual to the error of the inverse

Theorem 9.4 (Residual bracket). *Let F be continuous on $[A, B]$ and differentiable inside with $F' \geq \mu > 0$. Let $x_0 \in (A, B)$, $y \in \mathbb{R}$, $\varrho = F(x_0) - y$. If $[x_0 - |\varrho|/\mu, x_0 + |\varrho|/\mu] \subseteq [A, B]$, then $F(x) = y$ has exactly one solution x_* in $[A, B]$, it lies in that interval, and*

$$|x_* - x_0| \leq \frac{|\varrho|}{\mu}. \quad (9.4)$$

If also $F' \leq M_1 < \infty$ on (A, B) , then $|x_ - x_0| \geq |\varrho|/M_1$.*

Proof. $F(x_0 + |\varrho|/\mu) \geq F(x_0) + |\varrho| \geq y$ and $F(x_0 - |\varrho|/\mu) \leq y$; the intermediate value theorem gives a root in the interval, monotonicity its uniqueness in $[A, B]$, and the mean value theorem gives the error bounds. A finite derivative maximum was not part of the hypotheses, so the lower error bound is stated separately. $\square$

This gives an a posteriori certificate that needs no expansion theory, only a lower bound on the derivative. For the two examples the global bounds (4.5) give, for every $y > 0$ and every $G > 0$,

$$|G - g_1(y)| \leq \frac{|G + G^2(1 + \log G) - y|}{1 - 2e^{-5/2}}, \quad |G - g_2(y)| \leq |G + G^{\sqrt{2}} - y|. \quad (9.5)$$

For a general germ, $|f'| \asymp z^{p-1}$. For an approximation $G \sim z$, apply the bracket on $[z/2, 2z]$ with $\mu = \frac{1}{2}|ap|z^{p-1}2^{-|p-1|}$, after replacing (f, y) by $(-f, -y)$ if the derivative is negative. A residual small enough that the bracket lies in this interval yields $|G - g| \leq \frac{2^{1+|p-1|}}{|ap|}z^{1-p}|f(G) - y|$.

The residual of the truncated expansion. Combining Theorem 8.2 with the mean value theorem, the truncated inverse G_H satisfies

$$\frac{f(G_H(y)) - y_0}{a z^p} - 1 = -p z^{\sigma*} C_{\sigma*}(L) + o(z^{\sigma*}), \quad (9.6)$$

Here $y_0 = 0$ for $p < 0$, and the truncation retains the leading term so $G_H \sim z$. Indeed $f'(s) = apz^{p-1}(1 + \mathcal{O}(z^{\delta*}M^D))$ between g and G_H , with a positive power margin that absorbs the logarithms in the error. Multiplying by $G_H - g$ proves the stated little- o term. The normalized residual has the first error block multiplied by $-p$; `InverseResidual` computes it in the jet algebra. If the inverse is retained to cutoff q in w , its unnormalized forward residual cutoff is $q + (p - 1)/|p|$; the normalized cutoff subtracts $p/|p|$ from that value.

9.3 Finite smoothness suffices

The finite germ is analytic away from 0, and the convergence proof used that. For perturbations outside the finite power-log class (a coefficient $\sin(\log x)$, a term $e^{-1/x}$, a function known only through derivative bounds) the marker formula (5.3) can still be applied, and the following theorem, which appears in report 05 with a sketched proof, shows that it is a genuine asymptotic expansion whenever the perturbation and finitely many of its derivatives are small in the power-log scale. We give the complete proof.

Theorem 9.5 (Tame near-identity inversion). *Let $N \geq 0$, $\delta > 0$ and $M \geq 0$. Let φ be real-valued and C^{N+1} on $(0, x_0)$ with*

$$\varphi^{(j)}(x) = \mathcal{O}(x^{1+\delta-j}M(x)^M) \quad (x \rightarrow 0^+), \quad j = 0, 1, \dots, N+1. \quad (9.7)$$

Then $f = x + \varphi$ has a unique local inverse g on $(0, y_0)$ with $g(y) \sim y$, and

$$g(y) = y + \sum_{n=1}^N \frac{(-1)^n}{n!} \frac{d^{n-1}}{dy^{n-1}} [\varphi(y)^n] + \mathcal{O}(y^{1+(N+1)\delta}M(y)^{(N+1)M}). \quad (9.8)$$

Proof. Step 1: the homotopy. For $\lambda \in [0, 1]$ put $F_\lambda(x) = x + \lambda\varphi(x)$. By (9.7) with $j = 1$, $\varphi'(x) \rightarrow 0$, so there is x_1 with $F'_\lambda \geq \frac{1}{2}$ on $(0, x_1)$ for all $\lambda \in [0, 1]$; each F_λ is an increasing bijection of $(0, x_1)$ onto an interval containing $(0, y_1)$ for some y_1 independent of λ . Let $G(\lambda, y)$ be its inverse, so that

$$G + \lambda\varphi(G) = y, \quad G(0, y) = y. \quad (9.9)$$

Since $|\varphi(x)| \leq x/2$ for $x < x_1$ (after shrinking x_1), $F_\lambda(y/2) < y < F_\lambda(2y)$ and therefore $G(\lambda, y) \in [y/2, 2y]$ for all λ and all small y . The map $(\lambda, x) \mapsto x + \lambda\varphi(x) - y$ is C^{N+1} with x -derivative $\geq \frac{1}{2}$, so by the implicit function theorem G is C^{N+1} in (λ, y) , with

$$\partial_y G = \frac{1}{1 + \lambda\varphi'(G)} \in [\frac{2}{3}, 2], \quad \partial_\lambda G = -\frac{\varphi(G)}{1 + \lambda\varphi'(G)} = \psi(G) \partial_y G, \quad \psi := -\varphi. \quad (9.10)$$

Step 2: the transport identity. For every C^n function Ψ and $1 \leq n \leq N+1$,

$$\partial_\lambda^n [\Psi(G)] = \partial_y^{n-1} [\psi(G)^n \partial_y \Psi(G)]. \quad (9.11)$$

For $n = 1$: $\partial_\lambda \Psi(G) = \Psi'(G) \partial_\lambda G = \Psi'(G) \psi(G) \partial_y G = \psi(G) \partial_y [\Psi(G)]$. Assume (9.11) for n and all Ψ . Then, using it once more for the functions ψ^n and Ψ in the second step,

$$\begin{aligned} \partial_\lambda^{n+1} [\Psi(G)] &= \partial_y^{n-1} \partial_\lambda [\psi(G)^n \partial_y \Psi(G)] \\ &= \partial_y^{n-1} [\psi(G) \partial_y (\psi(G)^n \partial_y \Psi(G))] \end{aligned}$$

$$\begin{aligned}
& + \psi(G)^n \partial_y(\psi(G) \partial_y \Psi(G)) \\
& = \partial_y^{n-1} \partial_y [\psi(G)^{n+1} \partial_y \Psi(G)],
\end{aligned}$$

where the last equality is the product rule read backwards: $\partial_y[\psi^{n+1} \Psi_y] = (n+1)\psi^n \psi_y \Psi_y + \psi^{n+1} \Psi_{yy}$ and the bracket equals $n\psi^n \psi_y \Psi_y + \psi^n(\psi_y \Psi_y + \psi \Psi_{yy})$, the same thing. (Here ψ_y stands for $\partial_y[\psi(G)]$.)

Step 3: the coefficients. At $\lambda = 0$, $G = y$, so (9.11) with $\Psi = \text{id}$ gives $\partial_\lambda^n G(0, y) = \partial_y^{n-1}[\psi(y)^n] = (-1)^n \partial_y^{n-1}[\varphi(y)^n]$, the Lagrange terms. Taylor's formula in λ with integral remainder gives

$$g(y) = G(1, y) = y + \sum_{n=1}^N \frac{(-1)^n}{n!} \partial_y^{n-1}[\varphi^n] + \frac{1}{N!} \int_0^1 (1-\lambda)^N \partial_\lambda^{N+1} G(\lambda, y) d\lambda.$$

Step 4: the remainder. Put $\kappa(y) = y^\delta M(y)^M$. Since $G \in [y/2, 2y]$ and $M(G) \leq 2M(y)$, hypothesis (9.7) gives $\varphi^{(j)}(G) = \mathcal{O}(\kappa y^{1-j})$ uniformly in λ , for $j \leq N+1$. We claim that for $1 \leq j \leq N+1$,

$$\partial_y^j G = \mathcal{O}(y^{1-j}) \quad (j \geq 1), \quad \partial_y^j [\varphi(G)] = \mathcal{O}(\kappa y^{1-j}) \quad (j \geq 0), \quad (9.12)$$

uniformly in $\lambda \in [0, 1]$. Differentiating (9.9) j times in y and solving for $\partial_y^j G$: $(1 + \lambda \varphi'(G)) \partial_y^j G$ equals δ_{j1} minus λ times a sum of terms $\varphi^{(m)}(G) \prod \partial_y^{i_s} G$ with $2 \leq m$, $\sum i_s = j$, $i_s \geq 1$ (Faà di Bruno). By induction on j each such term is $\mathcal{O}(\kappa y^{1-m}) \prod \mathcal{O}(y^{1-i_s}) = \mathcal{O}(\kappa y^{1-j})$, so $\partial_y^j G = \mathcal{O}(\kappa y^{1-j}) \subseteq \mathcal{O}(y^{1-j})$ for $j \geq 2$ (and $\partial_y G = \mathcal{O}(1)$). The same Faà di Bruno expansion with $m \geq 1$ gives the second claim. By Leibniz, $\partial_y^j [\psi(G)^{N+1}] = \mathcal{O}(\kappa^{N+1} y^{N+1-j})$, and $\partial_y^{N-j} [\partial_y G] = \partial_y^{N-j+1} G = \mathcal{O}(y^{-(N-j)})$. Applying (9.11) with $\Psi = \text{id}$ and $n = N+1$,

$$\partial_\lambda^{N+1} G = \partial_y^N [\psi(G)^{N+1} \partial_y G] = \sum_{j=0}^N \binom{N}{j} \partial_y^j [\psi(G)^{N+1}] \partial_y^{N-j+1} G = \mathcal{O}(\kappa^{N+1} y)$$

uniformly in λ , and the integral remainder is $\mathcal{O}(y \kappa(y)^{N+1}) = \mathcal{O}(y^{1+(N+1)\delta} M(y)^{(N+1)M})$. $\square$

Remark 9.6. Applied to $\varphi(x) = x^2(1 + \log x)$ ($\delta = 1$, $M = 1$) the theorem re-proves the remainder $\mathcal{O}(y^{N+2} M^{N+1})$ of (8.5) with a proof that uses no analyticity at all; applied to $\varphi(x) = x^2 \sin(\log x)$ ($\delta = 1$, $M = 0$) it justifies $g(y) = y - y^2 \sin L + y^3 \sin L(2 \sin L + \cos L) + \mathcal{O}(y^4)$, $L = \log y$, a perturbation whose coefficients oscillate and which no power-log expansion describes. The package's `PerturbativeInverse` produces the terms; the theorem is the user's tool for asserting the remainder.

10 Finite endpoints, infinity, and powers of the inverse

The branch theorem, coefficient formula and parametric convergence proof only need $p \neq 0$ and $u \rightarrow 0^+$: the substitution $t = u^p$ and the marker theorem are applied at positive points regardless of the endpoint approached by t . Statements that use positivity explicitly, such as the exponent-cutoff corollary to the Rouché majorant, retain their stated restrictions. This section makes the endpoint bookkeeping explicit.

10.1 Local variables

Let f be given near x_0 , and let $u \rightarrow 0^+$ be the local variable of x :

$$u = \begin{cases} x - x_0 & x \rightarrow x_0^+ \text{ (Direction } \rightarrow \text{ "FromAbove")}, \\ x_0 - x & x \rightarrow x_0^-, \\ 1/x & x \rightarrow +\infty, \\ -1/x & x \rightarrow -\infty. \end{cases} \quad (10.1)$$

Substituting into f and expanding (Section 3) gives a germ $y_0 + au^p(1 + \sum u^{\delta_i} B_i(\log u))$ in the sense of (4.1), where $y_0 = \lim f$ is finite when $p > 0$. For $p < 0$ the target endpoint is $\pm\infty$ with the sign of a , while the additive offset y_0 in (4.1) is set to 0; any constant term is a correction with gap $-p$. The local variable of y is

$$w = \begin{cases} y - y_0 & a > 0, p > 0, \\ -(y - y_0) & a < 0, p > 0, \\ 1/|y| & p < 0. \end{cases} \quad (10.2)$$

The uniformizer $z = (v/a)^{1/p}$ tends to 0^+ in all cases. Its relation to w is $w = |a|z^p$ for $p > 0$ and $w = z^{|p|}/|a|$ for $p < 0$.

10.2 The theorem at an arbitrary endpoint

Theorem 10.1 (Signed leading power). *Theorems 4.1, 6.1, 7.1 and 8.2 hold verbatim for the germ (4.1) with $p < 0$: the unique branch with $u \rightarrow 0^+$ is $u = g(y)$ with $g \sim z$, g^r has the expansion (6.1) convergent for small z (i.e. for $|y|$ large), and the remainder after the exponent cutoff q in $w = 1/|y|$ is $\mathcal{O}(w^{(r+\sigma_*)/|p|} M(w)^{k_*})$.*

Proof. The derivative $f'(u) = apu^{p-1}(1 + o(1))$ has a constant sign, so f is monotone on $(0, \varepsilon)$ with $f \rightarrow \pm\infty$; the inverse exists on a neighbourhood of $\pm\infty$ and $u \sim z$. The proof of Theorem 6.1 is a computation at a fixed positive $q = v/a$ and never uses the sign of p . The parametric equation (7.2) has $\partial_w \mathcal{E}(1, 0) = p \neq 0$ regardless of sign, and the specialisation $T_{ij} = z^{\delta_i} L^j \rightarrow 0$ as $z \rightarrow 0^+$ is the same. The remainder theorem is Lemma 3.2. $\square$

In the variable y the blocks of g^r read $(v/a)^{(r+\sigma)/p} C_\sigma^{(r)}(\frac{1}{p} \log(v/a))$, with negative exponents of y when $p < 0$; the exponent of w is $(r + \sigma)/|p|$, which is the quantity the cutoff refers to.

10.3 Powers of the inverse and the original variable

The observable power r in (6.1) gives directly the expansion of u^r . In terms of x :

$$\begin{aligned} (x - x_0)^r &= g(y)^r (x \rightarrow x_0^+), & (x_0 - x)^r &= g(y)^r (x \rightarrow x_0^-), \\ x^r &= g(y)^{-r} (x \rightarrow +\infty), & x^r &= (-1)^r g(y)^{-r} (x \rightarrow -\infty), \end{aligned} \quad (10.3)$$

The last equality in (10.3), involving $(-1)^r$, is used here only for integer r . Thus x^r at infinity is obtained from (6.1) with local-coordinate power $-r$. Both signed distances $x - x_0$ from above and $x_0 - x$ from below are positive and admit arbitrary real powers. For $r = 1$ and a finite endpoint, the original inverse is $x = x_0 \pm g(y)$. A non-integer power of a negative original variable needs an additional branch convention and is outside the positive-real-power convention of this article.

10.4 Examples

Example 10.2 (The Wright omega function). $f(x) = x + \log x$ as $x \rightarrow +\infty$: $u = 1/x$, $f = u^{-1} - \log u = u^{-1}(1 - u \log u)$, so $a = 1$, $p = -1$, one gap $\delta = 1$ with $B = -L$, and target endpoint $+\infty$. The uniformizer is $z = 1/y$, $L = -\log y$, and (6.1) with $r = -1$ gives the original variable $x = 1/g(y)$:

$$x = y - \log y + \frac{\log y}{y} + \frac{\frac{1}{2} \log^2 y - \log y}{y^2} + \frac{\frac{1}{3} \log^3 y - \frac{3}{2} \log^2 y + \log y}{y^3} + \mathcal{O}\left(\frac{\log^4 y}{y^4}\right), \quad (10.4)$$

the classical expansion of the Wright omega function $\omega(y) = W(e^y)$ [10, 11]; the blocks are the Lambert polynomials $P_n(\log y)$ of [11].

Example 10.3 (Reciprocal corrections). $f(x) = x + 1/x$ at $+\infty$: $u = 1/x$, $f = u^{-1}(1 + u^2)$, $p = -1$, gap 2, constant coefficient. The inverse is the larger root of $x^2 - yx + 1 = 0$, $x = (y + \sqrt{y^2 - 4})/2 = y - y^{-1} - y^{-3} - 2y^{-5} - \dots$ (the coefficients are Catalan numbers), which (6.8) reproduces.

Example 10.4 (Half-integer gaps). $f(x) = x + \sqrt{x}$ at $+\infty$: $f = u^{-1}(1 + u^{1/2})$, gap $\frac{1}{2}$. The inverse is $x = (\sqrt{4y + 1} - 1)^2/4 = y - \sqrt{y} + \frac{1}{2} - \frac{1}{8\sqrt{y}} + \frac{1}{128y^{3/2}} + \mathcal{O}(y^{-5/2})$; the constant term $\frac{1}{2}$ is the block of weight $2 \cdot \frac{1}{2}$, i.e. exponent 0 in $w = 1/y$, despite the infinite target endpoint.

Example 10.5 (A finite endpoint from below). $f(x) = x - x^2$ near $x_0 = 0$ with $x \rightarrow 0^-$: $u = -x$, $f = -u - u^2$, $a = -1$, $p = 1$, $y_0 = 0$, and $w = -y$ (the image is $y < 0$). The inverse is $u = -y - y^2 - 2y^3 - \dots$, i.e. $x = y + y^2 + 2y^3 + \mathcal{O}(y^4)$, the branch of $x = (1 - \sqrt{1 - 4y})/2$ continued to negative y .

Example 10.6 (A pole). $f(x) = 1 + 1/x$ near $x_0 = 0^+$: $f = u^{-1}(1 + u)$, $p = -1$, target endpoint $+\infty$, and $x = 1/(y - 1) = y^{-1} + y^{-2} + \dots$ with remainder $\mathcal{O}(y^{-3})$ after two terms.

Example 10.7 (The second question, inverted). For $F(x) = (1 + x + x^{\sqrt{2}})^{\sqrt{2}}$ at $+\infty$ the forward germ is $u^{-2}(1 + \sqrt{2}u^{\sqrt{2}-1} + \sqrt{2}u^{\sqrt{2}} + \dots)$ from (1.5), with gaps generated by $\sqrt{2} - 1$ and $\sqrt{2}$, $p = -2$, target endpoint $+\infty$, and $z = y^{-1/2}$. The inverse is

$$\begin{aligned} x = \sqrt{y} - \frac{1}{\sqrt{2}} y^{(2-\sqrt{2})/2} + \left(\frac{3}{4} - \frac{1}{2\sqrt{2}}\right) y^{(3-2\sqrt{2})/2} \\ + \left(1 - \frac{5}{3\sqrt{2}}\right) y^{(4-3\sqrt{2})/2} - \frac{1}{\sqrt{2}} y^{(1-\sqrt{2})/2} + \mathcal{O}(y^{(5-4\sqrt{2})/2}), \end{aligned} \quad (10.5)$$

the terms being ordered by decreasing exponent $-(r + \sigma)/2$ of y , with local-coordinate power $r = -1$. The forward remainder is transported by (9.3) with that value of r ; the package raises the forward order until the cap reaches the requested cutoff. An independent derivation removes the outer power exactly: $1 + x + x^{\sqrt{2}} = y^{1/\sqrt{2}}$, then applies the finite-germ coefficient formula. Numerically, with seven blocks at $y = 10^{16}$ the recorded error is $4 \cdot 10^{-11}$ against a leading term 10^8 (Table 1).

11 Worked examples

All expansions in this section were produced by the package and checked by the residual composition of Section 9 and numerically; the displayed logarithms are log of the indicated argument.

11.1 The two examples of the first question, once more

For $f_1 = x + x^2(1 + \log x)$ the package call `AsymptoticInverse[x + x^2 (1 + Log[x]), {x, 0}, {y, 6}]` returns (1.3) with the remainder `PowerLogRemainder[y, 6, 5]`; the sixth block is

$$\begin{aligned} P_5(L) &= -42L^5 - \frac{3727}{12}L^4 - \frac{5365}{6}L^3 - 1256L^2 - \frac{5177}{6}L - \frac{2789}{12} \\ &= -\frac{1}{12}(1+L)(504L^4 + 3223L^3 + 7507L^2 + 7565L + 2789), \end{aligned}$$

and the seventh $P_6 = \frac{1}{60}(1+L)(7920L^5 + 63852L^4 + 200833L^3 + 308537L^2 + 231883L + 68307)$. With rational exponents the same object is available as `SeriesData[y, 0, {1, -1 - Log[y], 3 + 5 Log[y] + 2 Log[y]^2, ...}, 1, 6, 1]`, on which the built-in series arithmetic works (with the caveat that its $0[y]^6$ hides the factor $M(y)^5$).

For $f_2 = x + x^{\sqrt{2}}$, `SeriesTermGoal -> 4` returns exactly the four terms requested in the question with remainder `PowerLogRemainder[y, 4 Sqrt[2] - 3, 0]`; the exponent cutoff that produces the same result is $4\sqrt{2} - 3$.

11.2 Nonunit leading power

$f(x) = 3x^2(1 + x(1 + \log x))$: $a = 3$, $p = 2$, $\delta = 1$, $B = 1 + L$. With $z = \sqrt{y/3}$ and $L = \log z = \frac{1}{2}\log(y/3)$, (6.3) gives

$$g(y) = z - \frac{z^2}{2}(1+L) + \frac{z^3}{8}(5L^2 + 12L + 7) - \frac{z^4}{48}(\mathcal{D} + 6)(\mathcal{D} + 8)(1+L)^3 + \mathcal{O}(z^5 M^4), \quad (11.1)$$

where $\frac{1}{48}(\mathcal{D} + 6)(\mathcal{D} + 8)(1+L)^3 = (1+L)^3 + \frac{7}{8}(1+L)^2 + \frac{1}{8}(1+L)$. Two ingredients are easy to forget: the factor $p^{-n} = 2^{-n}$ and the logarithm $L = \frac{1}{2}\log(y/3)$, not $\log y$. In the variable y the expansion reads

$$\begin{aligned} g(y) &= \frac{\sqrt{y}}{\sqrt{3}} - \frac{y}{3}\left(\frac{1}{2} + \frac{\log(y/3)}{4}\right) + \frac{y^{3/2}}{3\sqrt{3}}\left(\frac{7}{8} + \frac{3}{4}\log\frac{y}{3} + \frac{5}{32}\log^2\frac{y}{3}\right) \\ &\quad - \frac{y^2}{9}\left(2 + \frac{39}{16}\log\frac{y}{3} + \frac{31}{32}\log^2\frac{y}{3} + \frac{1}{8}\log^3\frac{y}{3}\right) + \mathcal{O}(y^{5/2} M^4). \end{aligned}$$

11.3 Several gaps

Irrational and logarithmic gaps together. $f(x) = x + x^{\sqrt{2}} + x^2(1 + \log x)$ has gaps $\delta_1 = \sqrt{2} - 1$ and $\delta_2 = 1$. Below the exponent 3 the support is $\{1, \sqrt{2}, 2\sqrt{2} - 1, 2, 3\sqrt{2} - 2, 1 + \sqrt{2}, 4\sqrt{2} - 3, 2\sqrt{2}\}$ and

$$\begin{aligned} g(y) &= y - y^{\sqrt{2}} + \sqrt{2}y^{2\sqrt{2}-1} - y^2(1+L) - \frac{6-\sqrt{2}}{2}y^{3\sqrt{2}-2} + y^{1+\sqrt{2}}((\sqrt{2}+2)L + \sqrt{2}+3) \\ &\quad + \left(\frac{17\sqrt{2}}{3} - 4\right)y^{4\sqrt{2}-3} - y^{2\sqrt{2}}\left((3\sqrt{2}+5)L + 5\sqrt{2} + \frac{13}{2}\right) + \mathcal{O}(y^3 M^2), \quad L = \log y, \end{aligned}$$

the mixed blocks coming from (6.4): the index $(1, 1)$ gives $(\mathcal{D}+2+\sqrt{2})(1+L) = (\sqrt{2}+2)(1+L)+1$. The remainder is the block $y^3(2L^2 + 5L + 3)$ of the index $(0, 2)$, of degree 2.

Two irrational gaps. $f(x) = x + x^{\sqrt{2}}(1 + \log x) + 2x^{\sqrt{3}}$ has gaps $\sqrt{2} - 1$ and $\sqrt{3} - 1$, whose sums are the numbers $k_1(\sqrt{2} - 1) + k_2(\sqrt{3} - 1)$; the package represents them by `RootReduce` (for instance $\sqrt{2} + \sqrt{3} = \text{Root}[1 - 10 \#^2 + \#^4 \&, 4]$) and displays them with `ToRadicals`. Below the exponent 3 there are eleven blocks; the residual check passes only when the coefficients, which involve $\sqrt{5 + 2\sqrt{6}} = \sqrt{2} + \sqrt{3}$, are canonicalised with `RootReduce`: this is precisely the zero-recognition defect that made one test of report 01 fail (Section 14).

11.4 Automatically expanded forward functions

$\sin x + x^2 \log x$ is not a finite power-log sum. The evaluator of Section 3.2 expands it to $x + x^2 \log x - \frac{1}{6}x^3 + \frac{1}{120}x^5 + \mathcal{O}(x^7)$, records $(\rho, k) = (7, 0)$, and the inverse below the exponent 4 is

$$g(y) = y - y^2 \log y + y^3 \left(\frac{1}{6} + \log y + 2 \log^2 y \right) + \mathcal{O}(y^4 M^3);$$

here the y^3 block combines $(\mathcal{D} + 4)L^2/2 = 2L^2 + L$ from the index $(2, 0)$ (gap 1, $B_1 = L$) with $+\frac{1}{6}$ from the index $(0, 1)$ (gap 2, $B_2 = -\frac{1}{6}$). Similarly $\tan x$ gives $\arctan$: **SeriesTermGoal** $\rightarrow$ 4 yields $y - \frac{1}{3}y^3 + \frac{1}{5}y^5 - \frac{1}{7}y^7$, the forward order being raised automatically until the transported remainder allows the fourth term.

11.5 Symbolic data

With symbolic coefficients the exponents remain numeric and everything goes through under assumptions that make the leading coefficient's sign and the reality of the coefficients provable: $f = ax + bx^2$ with $a > 0$, $b \in \mathbb{R}$ gives $g = y/a - by^2/a^3 + 2b^2y^3/a^5 + \mathcal{O}(y^4)$. With a symbolic exponent the ordering of the support is undecidable in general and depth truncation (Theorem 8.6) is used: $f = x + x^\alpha$ with $\alpha > 1$ gives, at depth 3,

$$g = y - y^\alpha + \alpha y^{2\alpha-1} - \frac{\alpha(3\alpha-1)}{2} y^{3\alpha-2} + \mathcal{O}(y^{4\alpha-3}),$$

which is (6.9) with a symbolic α .

11.6 Powers and logarithms of the inverse

For $f = x + x^2$, (6.1) with $r = 2$ gives $g^2 = y^2 - 2y^3 + 5y^4 - 14y^5 + \dots$ (Catalan numbers again), with $r = -1$ it gives $1/g = y^{-1} + 1 - y + 2y^2 - 5y^3 + \dots$, and (6.2) gives $\log g = \log y - y + \frac{3}{2}y^2 - \frac{10}{3}y^3 + \dots$. These expansions are exact consequences of the same coefficient formula and cost nothing extra; expanding **Normal**[g]^2 instead would silently lose the remainder information.

12 Lambert coordinates and boundaries of the theory

The convergent power-log class is closed under the operations of Theorem 3.3 and inversion of germs with a monomial leading block. A leading logarithm or exponential changes the natural coordinate. For exact Lambert cores, that coordinate change reduces the problem to the same polynomial-logarithmic arithmetic. More general perturbations require additional bookkeeping, described below.

12.1 A leading logarithm: Lambert W and the log log scale

Let $a \neq 0$, $p \neq 0$, $b \neq 0$, and $x \rightarrow 0^+$ with $x < 1$. The exact core $ax^p(-\log x)^b = y$ has $Y = y/a > 0$. With $\tau = -\log x \rightarrow +\infty$, it reads $\tau^b e^{-p\tau} = Y$, hence

$$x = \exp \left[\frac{b}{p} W_k \left(-\frac{p}{b} Y^{1/b} \right) \right], \quad k = \begin{cases} -1, & p/b > 0, \\ 0, & p/b < 0. \end{cases} \quad (12.1)$$

The branch is forced by $\tau \rightarrow \infty$: when $p/b > 0$ the argument approaches 0 from below and $W_{-1} \rightarrow -\infty$; when $p/b < 0$ it tends to $+\infty$ and $W_0 \rightarrow +\infty$. For $p, b > 0$, the other real root

on W_0 approaches $x = 1$ as $Y \rightarrow 0$, while the two branches meet at the maximum $x = e^{-b/p}$. For instance the inverse of $x \log(1/x)$ near 0^+ is $-y/W_{-1}(-y)$, and that of $x \log x$ on the same source branch is $y/W_{-1}(y)$ for $y \rightarrow 0^-$. For $p > 0$, with $\Lambda = \log(a/y) \rightarrow \infty$ and $\Lambda_2 = \log \Lambda$, the expansion begins

$$x = \left(\frac{y}{a}\right)^{1/p} \left(\frac{\Lambda}{p}\right)^{-b/p} \left[1 - \frac{b^2}{p\Lambda} \log \frac{\Lambda}{p} + \mathcal{O}\left(\frac{\Lambda_2^2}{\Lambda^2}\right)\right], \quad (12.2)$$

a series in $1/\Lambda$ with polynomial coefficients in Λ_2 , multiplying the leading factor. The following construction computes all its logarithmic orders, with convergence and a remainder bound.

12.2 A convergent universal Lambert expansion

Let Z approach $+\infty$ on the principal real branch or 0^- on the lower real branch. Set

$$s = \begin{cases} 1, & W_0(Z), \ Z \rightarrow +\infty, \\ -1, & W_{-1}(Z), \ Z \rightarrow 0^-, \end{cases} \quad A = |\log |Z||, \quad t = A^{-1}, \quad \ell = \log t.$$

Then $\log |W| + W = \log |Z| = sA$. Writing $W = s(1 + V)/t$ gives the normalized equation

$$V + st \log(1 + V) - st\ell = 0. \quad (12.3)$$

Theorem 12.1 (Lambert expansion in a logarithmic coordinate). *Equation (12.3) has a unique solution $V \rightarrow 0$ of the form $\sum_{n \geq 1} t^n P_n(\ell)$ with $P_1 = s\ell$ and $\deg P_n \leq n-1$ for $n \geq 2$. This series converges for sufficiently small $t > 0$. For $N \geq 1$, truncating V after t^N leaves $\mathcal{O}(t^{N+1}M(t)^N)$, and multiplying by s/t gives the corresponding remainder for W . In particular,*

$$W = \frac{s}{t} + \ell - s\ell t + \left(\ell + \frac{\ell^2}{2}\right)t^2 - s\left(\ell + \frac{3\ell^2}{2} + \frac{\ell^3}{3}\right)t^3 + \mathcal{O}(t^4 M(t)^4). \quad (12.4)$$

Proof. Replace $t\ell$ by an independent parameter T . The function $V + st \log(1 + V) - sT$ is analytic near $(V, t, T) = (0, 0, 0)$ and its V -derivative there is 1. The analytic implicit function theorem gives a convergent Taylor series $V(t, T)$. Moreover $V = sT + tH(t, T)$ for an analytic H with $H(0, 0) = 0$. After $T = t \log t$, every term of total degree $n \geq 2$ therefore has logarithmic degree at most $n-1$, and $P_1 = s\ell$. Cauchy estimates bound the coefficient norms exponentially. The tail starting at degree $N+1$ is bounded by a geometric series with first term $t^{N+1}M(t)^N$, proving the remainder. Coefficients can be computed by the iteration $V \leftarrow st\ell - st \log(1 + V)$, which gains one power of t each time, or by Newton iteration. The real solution has $W \sim sA$ and hence belongs to the specified real Lambert branch. Expanding the equation gives (12.4). $\square$

An explicit sufficient convergence condition is

$$q_L = 4t(1 + |\log t|) < 1. \quad (12.5)$$

Indeed, for complex $|t|, |T| \leq 1/4$, the map $V \mapsto sT - st \log(1 + V)$ sends $|V| \leq 1/2$ into itself, since $(1 + \log 2)/4 < 1/2$, and has derivative bounded by $1/2$. Its fixed point is jointly analytic and bounded by $1/2$, so its Taylor coefficients are bounded by $4^{a+b}/2$. Beyond the leading term sT , every coefficient has $a \geq 1$ because $V(0, T) = sT$. If V_N retains total degrees through $N \geq 1$, the geometric majorant therefore gives

$$|V - V_N| \leq \frac{2t q_L^N}{1 - q_L}, \quad \left|W - \frac{s}{t}(1 + V_N)\right| \leq \frac{2q_L^N}{1 - q_L}. \quad (12.6)$$

The construction follows the two-small-parameter method in the supplied transseries notes [11]; it establishes a sufficient domain and does not claim convergence from a fixed-logarithm implicit-function argument alone.

For the leading-log core, put $Z = -(p/b)Y^{1/b}$ and choose s from its branch. Since $-bs/p > 0$, an arbitrary real observable power is best assembled as

$$x^r = Y^{r/p} \left(-\frac{bs}{p} \right)^{-br/p} t^{br/p} (1+V)^{-br/p}. \quad (12.7)$$

This uses a positive real power of the unit $1+V$. Directly exponentiating a truncated W would shift the required precision by one order; (12.7) makes that precision shift explicit.

Example 12.2 (The inverse of $x \log x$). For $y \rightarrow 0^-$ let $A = \log(1/(-y))$ and $B = \log A$. The branch with $x \rightarrow 0^+$ is

$$x = \frac{-y}{A} \left[1 - \frac{B}{A} + \frac{B^2 - B}{A^2} + \frac{-B^3 + \frac{5}{2}B^2 - B}{A^3} + \mathcal{O}\left(\frac{(1+B)^4}{A^4}\right) \right].$$

The endpoint $x \rightarrow 1$ selects a different inverse germ and is not represented by this expansion.

12.3 Exponential cores at infinity

For a positive source coordinate x with $X = x^p \rightarrow +\infty$, consider $y = ax^b e^{cx^p}$, where $a \neq 0$, $p \neq 0$, $c \neq 0$, and $Y = y/a > 0$. If $b \neq 0$, taking the power p/b gives

$$X = \frac{b}{cp} W_k \left(\frac{cp}{b} Y^{p/b} \right), \quad k = \begin{cases} 0, & cp/b > 0, \\ -1, & cp/b < 0. \end{cases} \quad (12.8)$$

These branch choices ensure $X > 0$ and $X \rightarrow \infty$: the argument tends to $+\infty$ in the first case and 0^- in the second. If $b = 0$, inversion is elementary, $x = (\log Y/c)^{1/p}$. In the Lambert cases, with s, t, V as above,

$$x^r = \left(\frac{bs}{cp} \right)^{r/p} t^{-r/p} (1+V)^{r/p},$$

whose positive prefactor and unit series again have explicit precision tracking.

Example 12.3 (The inverse of xe^x at infinity). Here $x = W_0(y)$. With $A = \log y$ and $B = \log A$,

$$x = A - B + \frac{B}{A} + \frac{B(B-2)}{2A^2} + \frac{B(2B^2 - 9B + 6)}{6A^3} + \mathcal{O}\left(\frac{(1+B)^4}{A^4}\right).$$

Both this expansion and the $x \log x$ expansion use $t = 1/A$ as the local variable. An exponent cutoff in t specifies logarithmic accuracy; it is not a cutoff in y or $1/y$. A nonconstant external factor such as $-y$ must also multiply the remainder descriptor.

12.4 Perturbations of a Lambert core

Relative perturbations $F(x) = ax^p(-\log x)^b(1 + \sum_i x^{\delta_i} B_i(\log x))$ can be approached with the general-core formula (5.5). If ϕ is the exact core inverse and $\tau = -\log \phi(y)$, then

$$\phi'(y) = \left[ax^{p-1} \tau^{b-1} (p\tau - b) \right]_{x=\phi(y)}^{-1}.$$

The marker coefficients have powers of $\phi(y)$ and rational functions of τ , with denominators arising from $p\tau - b$. A convergent marker expansion requires analytic control on the chosen branch, and re-expanding its coefficients introduces a second order of truncation: perturbation power and logarithmic depth. The exact-core Lambert construction proves the logarithmic

expansion. Sections 18 and 27 retain the exact core and compute separate perturbation corrections for the admitted finite families.

The leading logarithmic polynomial itself can also be generalized before introducing a second perturbation scale.

Corollary 12.4 (An arbitrary polynomial at the leading exponent). *Let $F_0(u) = au^p Q(\log u)$, where $p \neq 0$ and Q is a nonconstant real polynomial of degree $q \geq 1$, on the real branch $u \rightarrow 0^+$. Its inverse has a convergent relative power-log expansion in a reciprocal-logarithm coordinate, even when Q is not a power of an affine polynomial.*

Proof. Absorb the leading coefficient of Q into a , so Q is monic, and set $b = [L^{q-1}]Q/q$ and $\tau = -\log u - b$. Then

$$Q(-\tau - b) = (-\tau)^q R(1/\tau), \quad R(v) = 1 + \sum_{j=2}^q r_j v^j.$$

With $Y = y/(a(-1)^q) > 0$ (after subtracting any constant target offset), $k = -p/q$, $s = \text{sign } k$ and $Z = kY^{1/q}e^{pb/q}$, take $A = s \log |Z|$, $t = 1/A$, and $\tau = (1 + V)/(|k|t)$. The inverse equation becomes

$$V + st \left[-\ell + \log(1 + V) + \frac{1}{q} \log R \left(\frac{|k|t}{1 + V} \right) \right] = 0. \quad (12.9)$$

Replacing $t\ell$ by an independent parameter again gives an analytic equation with V -derivative 1 at the origin. Its solution is jointly analytic, and $R(0) = 1$ makes its real powers analytic too. For any real local-coordinate observable power r ,

$$u^r = Y^{r/p} \left(\frac{A}{|k|} \right)^{-qr/p} (1 + V)^{-qr/p} R \left(\frac{|k|t}{1 + V} \right)^{-r/p}.$$

This yields convergence and the same type of finite-order logarithmic remainder as before. The sufficient numerical bound (12.6) was proved for $R = 1$ and is not asserted unchanged for a general R . $\square$

The polynomial correction modifies the logarithmic coefficients and must be included in the equation; it is not beyond all logarithmic orders. In general this inverse has no claimed exact **ProductLog** expression. The construction also applies to logarithmic cores at either infinite source endpoint after $u = 1/x$ or $u = -1/x$, using the observable and sign rules of Section 10.

Proposition 12.5 (Perturbations beyond every logarithmic order). *Suppose a finite leading-log core $F_0(u) = au^p Q(\log u)$ from Corollary 12.4, or an affine-log power core, is supplemented by finitely many polynomial-logarithmic terms of strictly larger u -exponent. Alternatively, suppose a growing exponential core $F_0(x) = ax^b e^{cx^p}$ with $c, p > 0$ at $x \rightarrow +\infty$ is supplemented by a finite power-log sum. Let g_0 and g be the corresponding inverse branches, with any common target offset subtracted, and let t be the Lambert coordinate of g_0 . For every fixed real r and every $N > 0$,*

$$g(y)^r - g_0(y)^r = o(g_0(y)^r t^N).$$

Consequently every finite relative logarithmic expansion of the core inverse is also an expansion of the full inverse, with the same finite-order remainder class.

Proof. In the leading-log case $F'_0(u) \sim pF_0(u)/u$, and the perturbation and its derivative are smaller by a factor $\mathcal{O}(u^\delta M(u)^D)$ for some $\delta > 0$ and finite D . A mean-value bracket near g_0 yields $(g - g_0)/g_0 = \mathcal{O}(g_0^\delta M(g_0)^D)$. Here $1/t$ is a positive constant times $|\log g_0|$ plus lower logarithmic terms, so this bound is $o(t^N)$ for every N .

In the exponential case $F'_0/F_0 \sim cpx^{p-1}$, whereas the additive perturbation R and its derivative have only power-log growth. A bracket of width $\mathcal{O}(|R(g_0)|/|F'_0(g_0)|)$ has slope comparable to $F'_0(g_0)$: the logarithm of that derivative changes by $\mathcal{O}(R/F_0) = o(1)$ across the bracket. Thus $(g - g_0)/g_0$ is exponentially small in g_0^p , while t is comparable to g_0^{-p} . Finally apply the mean value theorem to the positive real power u^r . Continuous local inverse branches exist because the perturbation derivative is smaller than the nonzero core derivative. $\square$

The package marks this situation with `"LeadingCoreOnly" -> True`, stores the discarded expression in `"BeyondLogarithmicOrders"`, and leaves `"ExactInverseExpression"` unavailable. Those terms affect an exponentially smaller scale even though they do not change any finite coefficient in the reported logarithmic expansion.

12.5 Zero gaps: corrections of relative size $1/\log$

For $f(x) = x + x/\log x$ the correction has power gap 0 and relative size $1/\log x$. All corrections of the inverse then sit at the same power y with increasing negative powers of $\log y$, the index set I_H is infinite, and the exponent cutoff has no meaning. The right coordinate is $\tau = 1/\log y$: with $g = yW$ the equation becomes $W(1 + \tau/(1 + \tau \log W)) = 1$, analytic near $(W, \tau) = (1, 0)$, so W is an ordinary convergent power series in τ , $W = 1 - \tau + \tau^2 - 2\tau^3 + \frac{7}{2}\tau^4 - \dots$, i.e. $g(y) = y(1 - 1/\log y + 1/\log^2 y - 2/\log^3 y + \dots)$. The general-core formula produces these terms, but the scale is the reciprocal-logarithm scale, not the power-log scale, and Section 21 implements the required reciprocal-logarithm valuation.

12.6 Flat terms

$e^{-1/x}$ is smaller than every power of x ; a germ $x + e^{-1/x}$ has the same asymptotic power-log expansion as x . Its inverse obeys $g = y - e^{-1/g}$, so $y - e^{-1/y} \leq g < y$ and $(y - g)/y^2 \rightarrow 0$. Consequently $e^{-1/g}/e^{-1/y} \rightarrow 1$, proving $g - y = -e^{-1/y}(1 + o(1))$. The marker formula begins $y - e^{-1/y} + y^{-2}e^{-2/y} + (y^{-3} - \frac{3}{2}y^{-4})e^{-3/y} + \dots$. These exponential sectors are invisible to a pure power-log expansion. Section 22 computes them with a separate sector cutoff and a complete-tail bound; this differs from the exact exponential cores above, whose inverses reduce to Lambert coordinates.

12.7 Oscillatory coefficients

$x^2 \sin(\log x)$ is $\mathcal{O}(x^2)$ with all derivatives of the right size, so Theorem 9.5 applies and the marker formula gives a valid expansion $y - y^2 \sin L + y^3 \sin L(2 \sin L + \cos L) + \mathcal{O}(y^4)$ whose “coefficients” are bounded oscillating functions of $L = \log y$. They are not polynomials in L , the coefficient algebra is different, and the sharp remainder theory of Section 8 (which uses degrees) does not apply; only the tame theorem does. Section 23 supplies a finite Fourier coefficient algebra and a nonoscillatory remainder bound. The lower-level `PerturbativeInverse` remains a formula generator.

12.8 Symbolic exponents and nonuniformity

When a gap δ is a symbol, the order of the weights $k \cdot \delta$ depends on the parameter: the order of $2\delta_1$ and δ_2 changes at $\delta_2 = 2\delta_1$. Exponent truncation needs provable comparisons; depth truncation (Theorem 8.6) only needs positive gaps after a valid leading term and real branch have been selected. Nothing is uniform as a gap tends to 0: the y -radius $R_\alpha^{1/(\alpha-1)}$ of the $x + x^\alpha$ inverse tends to 0 as $\alpha \downarrow 1$, although $R_\alpha \rightarrow 1$, and at $\alpha = 1$ the inverse is $y/2$.

12.9 Nested powers of the logarithm

Non-integer powers $(-\log x)^s$, $\log(-\log x)$, and negative powers of $\log x$ in a correction leave the polynomial coefficient algebra. For example $\varphi(x) = x^{1+\delta}(-\log x)^{1/2}$ is real for small positive x and satisfies Theorem 9.5 with $M = 1$, but repeated differentiation produces negative half-integer as well as positive half-integer powers of $-L$. A suitable coefficient algebra therefore contains $(-L)^{1/2}$ and its inverse. The Euler identity and composition recurrence still apply when the coefficient functions are differentiable or analytic as needed; closure and remainder bounds in the enlarged scale must also be established.

The finite hierarchy of Section 21 establishes these closure and remainder contracts for generalized Laurent monomials in finitely many positive logarithmic levels. More general coefficient functions still need an independently justified algebra and ordering.

13 The package

The package `AsymptoticInverse` (directory `AsymptoticInverse/` of the repository; usage in its `README.md`) implements the power-log engines of Sections 2–10, the Lambert-coordinate construction of Section 12.2, and the logarithmic target transformation of Section 15. Its interface is

```
AsymptoticInverse[f, {x, x0}, {y, cutoff}]      (* or ..., y, SeriesTermGoal -> n
*)
AsymptoticExpansion[f, {x, x0, cutoff}]        (* or {x, x0}, SeriesTermGoal -> n
*)
```

with x_0 a real number, `Infinity` or `-Infinity`, options `Assumptions`, `Direction`, `Method` ("Lagrange" or "Newton" for the power-log engine), `"Power"` (r), `"Truncation"` ("Exponent" or "Depth"), `"InputRemainder"` and `"MaxTerms"`. Both return a `PowerLogSeries` object: `Normal` gives the finite expression, `s["Remainder"]` its inert descriptor, `s["FrontierTerm"]` the computed first omitted block, `s["Terms"]` exponent/coefficient pairs, `s["SeriesData"]` a `SeriesData` object when the scale permits it, and `s["Properties"]` the full property list.

The auxiliary functions provide residual checking through `InverseResidual`, numerical comparison through `InverseNumericalCheck`, marker formulas through `PerturbativeInverse`, individual coefficients through `InverseExpansionCoefficient`, and normalized model data through `PowerLogModel`.

Expansion objects support explicit arithmetic, composition, observables, truncation and refinement through the `Series` operations of Section 17. The operation `AsymptoticCoreInverse` retains an exact inverse while computing higher-power perturbation corrections with the complete-tail contract of Section 18. `InverseCertificate` uses rational intervals and explicit elementary-function tail estimates to enclose a numerical root; Section 19 states the existence, uniqueness and accuracy conditions.

Recognized Lambert cores dispatch to `"Method" -> "Lambert"`. Their objects carry `"Scale" -> "Logarithmic"`, an exact `"Prefactor"`, and `"LogarithmicVariable" -> t`. A pair $\{n, C\}$ contributes `Prefactor` $t^n C(\log t)$, and a cutoff q retains $n < q$ in this normalized bracket. The remainder is `|Prefactor| PowerLogRemainder[t,  $\beta$ , k]`; the prefactor must not be lost when estimating errors. `"LambertBranch"` and `"TargetDomain"` record the selected real branch. Such objects have no ordinary `SeriesData` in y , and depth truncation in the original power gaps does not apply. For example,

```
AsymptoticInverse[x Log[x], {x, 0}, {y, 4}]
AsymptoticInverse[x Exp[x], {x, Infinity}, {y, 5}]
```


produce the displayed logarithmic orders in Section 12.2. The first uses $y \rightarrow 0^-$ and the lower Lambert branch; the second uses $y \rightarrow +\infty$ and the principal branch.

13.1 Data structures and exact comparisons

A jet is a sorted list of pairs $\{\sigma, C_\sigma\}$, with exact exponents and polynomial coefficients. Algebraic exponents have `RootReduce` canonical forms, so $\sqrt{2} + \sqrt{3}$ and $\sqrt{5 + 2\sqrt{6}}$ are recognized as equal. For other exact exponents, symbolic simplification must prove the comparisons used by the computation; it is not a complete equality algorithm for arbitrary transcendental expressions. Undecidable order comparisons produce a **Failure**. Inexact **Real** input is rejected because a rounded exponent does not identify the exact semigroup: choosing a rational decimal approximation to $\sqrt{2}$ changes both the problem and its possible resonances.

Coefficients are canonicalised too: algebraic numbers with `RootReduce`, logarithms of positive rationals into sums of logarithms of primes, symbolic parameters with `Simplify` under the assumptions. Without this, exactly vanishing residual coefficients such as $-2\sqrt{2} - 2\sqrt{3} + 2\sqrt{5 + 2\sqrt{6}}$ or $-\frac{1}{16} \log^2 4 - \frac{1}{64} \log 2 \log 32 + \frac{1}{64} \log 8 \log 128$ are not recognised as zero by `Simplify`, which is what caused the two test failures among the nine reports (Section 14).

13.2 The engines

Forward. The evaluator of Section 3.2 is a structural recursion returning triples (T, P, D) . Finite power-log sums are recognized as exact. The working order grows when cancellation hides the leading term or prevents the requested precision from being reached. Analytic heads are expanded in their small argument, so expressions such as $\sin(x \log x)$ are handled without asking `Series` to order their full composition. Laurent and Puiseux expansions of supported heads are composed in the same way, with the local increment's sign checked at real branch points. Absolute values use the eventual sign of the leading logarithmic polynomial.

Inverse, Lagrange. The forward germ is normalised (Section 4); the multi-indices of weight $< H$ are enumerated by a depth-first search with exact comparisons ((8.1); the one-step boundary is formed at the same time); for each index the polynomial loop

```
q = Expand[Times @@ MapThread[Power, {polys, k}]];
Do[q = Expand[D[q, ell] + (r + w + p j) q], {j, 1, n - 1}];
q = (-1)^n r q / (p^n Times @@ (Factorial /@ k))
```

evaluates (6.1); equal weights are merged and the blocks of weight σ_* on the boundary are summed into the frontier term. For `SeriesTermGoal -> n`, a persistent frontier visits successive complete weights until n nonzero blocks exist, reusing previous coefficients. The independent "GroupedLagrange" method combines collisions before applying the Euler operators; Section 16 derives both changes.

Inverse, Newton. The normalised equation (6.11) and its derivative are evaluated with the composition recurrence (Lemma 2.7), and the derivative's reciprocal by the geometric series of a unit with constant leading term. The working cutoff doubles through exact weights until it reaches H ; each recorded precision has a verified formal residual. The three ordinary engines are compared in the test suite.

Lambert. The separate `Kernel/LambertInverse.wl` module recognizes affine-log powers, arbitrary leading logarithmic polynomials and exponential-power cores. It solves (12.3) or its

polynomial correction (12.9) using unit-series arithmetic, then forms the observable bracket. Logarithmic cores at finite endpoints and either infinity use the same positive source coordinate. For an arbitrary leading polynomial, metadata does not claim an exact `ProductLog` inverse. The original forward expression is retained for numerical checks. The restricted perturbations of Proposition 12.5 carry metadata identifying the leading-core approximation.

Assembly. Blocks $z^{r+\sigma}C(L)$ are converted to $(v/a)^{(r+\sigma)/p}C(\log(v/a)/p)$, and the sign and shift of (10.3) are applied. The forward remainder is transported with exponent $\rho - p + r$ in the local source variable; at infinite source endpoints, the requested original-variable power changes the local observable to $-r$. This observable dependence is essential for a valid cap on the requested cutoff.

13.3 Checks

`InverseResidual` composes the stored finite forward model with the stored blocks and checks whether $(f(G) - y_0)/(az^p) - 1$ vanishes below the relative cutoff. A larger cutoff can expose the first omitted block and its factor $-p$ from (9.6). This is independent of the coefficient formula, but its scope is the finite model. Lambert objects use the normalized Lambert equation and identify that scope separately. `InverseNumericalCheck` solves the original $f(x) = y_1$ at high precision, checks the selected real source side, and reports the approximation error and remainder scale; numerical agreement does not prove an asymptotic bound. The test runner `Tests/RunTests.wl` combines the original examples with review regressions, expanded-input cases, Lambert cases and performance regressions.

13.4 Cost

The retained multi-indices number at most $\prod_i(1 + \lfloor H/\delta_i \rfloor)$, and each coefficient uses $|\mathbf{k}| - 1$ bidiagonal polynomial operations. Sparse multiplication now skips product pairs already beyond the cutoff; nonnegative integer powers use repeated squaring; and depth enumeration visits the simplex $|\mathbf{k}| \leq N$ and its boundary directly instead of constructing a Cartesian box. Resource limits are checked while indices are generated. Exact algebraic canonicalization and sparse composition remain significant costs, and Newton’s logarithmic iteration count alone does not predict which method is faster.

The reproducible `Tests/BenchmarkPerformance.wl` compares these changes with frozen reference algorithms and checks equality of their results. A Windows/Wolfram run recorded the following seconds; these are fixture measurements, not portable speed guarantees.

Fixture	Reference	Updated
200 × 200 product, cutoff 5	2.90063	0.02723
Exact monomial to power 2048	2.08708	0.01991
Depth 2 in 10 dimensions	4.14531	0.03753
Weights {1, 2, 3, 5}, cutoff 20	0.23768	0.41552

The last fixture is slower after bounded allocation; the improvements target wasted product work, repeated multiplication and Cartesian depth enumeration rather than guaranteeing a speedup for every input.

14 The nine reports

Nine independently written reports on the first question [12] preceded this article; each contains an article of 22–30 pages, a Wolfram Language package, native tests, and a Python verification record, and none of them had been run in a Wolfram kernel by its authors. Their detailed comparison is in `reports/COMPARISON.md`; this section summarises what they agree on, what each adds, and what was wrong.

14.1 Agreement

All nine treat the finite germ (4.1) (report 07 with $p = 1$ only), select the real branch by $x/(y/a)^{1/p} \rightarrow 1$, prove the coefficient formula (6.1) with exactly the same normalization through the marker form of Lagrange–Bürmann inversion, prove convergence (four by the parametric implicit function theorem, five by Rouché’s theorem with Cauchy estimates), prove a remainder theorem with the logarithmic factor, derive the transport of forward remainders, and compute the same expansions (1.3)–(1.4) with the Catalan and Fuss–Catalan laws. All represent expansions as sparse exact jets, reject floating-point input, and provide a residual check.

14.2 Individual contributions absorbed here

- Report 01: the residue proof of the marker theorem and its formal clause; the g^r and $\log g$ formulas; the `InputRemainder` interface with cutoff cap; the $x + x/\log x$ analysis.
- Report 02: the observable power r built into the main formula and implemented; the Picard fixed-point engine; the explicit marker-disc convergence criterion; the partial resummation in $y \log y$.
- Report 03: the exact degree and leading-coefficient formula; the contraction proof of convergence; the search bound $\lceil B/\delta_* \rceil \delta_*$ for the frontier.
- Report 04: the first omitted block as an asymptotic equality with its complete polynomial; the successive-power engine; the condition number.
- Report 05: the tame near-identity theorem (here Theorem 9.5, with a complete proof); Picard iteration; global bounds for both examples.
- Report 06: the composition recurrence (Lemma 2.7); the Newton engine in the jet algebra; the rising-factorial closed form (6.8).
- Report 07: the explicit conditional majorant $\mathcal{Q}_r$ (Theorem 7.3) and its cutoff form; the negative-side germ by reflection.
- Report 08: the explanation of the complex germ (1.2); the Lambert- W core (12.1) and the $\log \log$ scale; the finite-point slope bound; flat perturbations.
- Report 09: depth truncation with symbolic exponents; the harmonic-number formula (6.7); the explicit tail majorant for the log example.

14.3 Errors found

Report 07 asserts that a nonunit leading power forces an infinite binomial expansion; the substitution $t = x^p$ in Theorem 6.1 shows that the coefficient formula stays finite and exact. Reports 07 and 08 state remainder degrees $(d[\sigma_*/\delta], Nd)$ that are valid but can be far from the

exact degree of the frontier block. Report 05's proof of the tame theorem omits the inductive estimates of Step 4 of Theorem 9.5. Two residual checkers (reports 01 and 04) rely on `Simplify` to recognise zero and fail on denested radicals and on logarithms of composite integers.

14.4 Native execution

Running the nine test suites on Wolfram 15.0.1 gave: 01: 35/36; 02: 95/95; 03: 34/34; 04: 34/35; 05: 38/38; 06: 33/33; 07: 34/34; 08: 40/40; 09: 35/35. The two failures are the zero-recognition defects just mentioned, in residual checkers, not wrong expansions. Every package therefore computes what its article claims, which is a strong independent confirmation of the formula (6.1) and of the two answers. What none of them does, and what motivated the present package, is to treat finite endpoints from both sides and infinity through signed leading powers, to expand arbitrary forward functions with a rigorous remainder (which is the second question), to track precision and logarithmic degree through compositions, and to return `SeriesData` when the exponents allow it.

15 Inversion through explicit coordinate transformations

The leading expression in the original variable need not belong to the scale used by the inverse engine. A change of source or target coordinate can put the equation into a supported scale while preserving the requested real branch. The transformation is part of the mathematical result: it determines the expansion variable, how the cutoff is interpreted, and how an inner remainder becomes an error in the original variable.

15.1 Target transformations

Theorem 15.1 (An exact change of target coordinate). *Let F be continuous and strictly monotone on a source interval I , and let J be continuous and one-to-one on $F(I)$. Put $\Phi = J \circ F$. If h is the inverse of Φ on $J(F(I))$, then the inverse of F on $F(I)$ is*

$$g(y) = h(J(y)).$$

If an observable of that inverse has an expansion $h(v)^r = A(v) + \mathcal{O}(R(v))$ as v approaches the transformed target endpoint, then

$$g(y)^r = A(J(y)) + \mathcal{O}(R(J(y)))$$

on the corresponding source branch, provided the same real power is defined there.

Proof. $\Phi(h(J(y))) = J(y)$ gives $J(F(h(J(y)))) = J(y)$. Injectivity of J gives $F(h(J(y))) = y$, and the source interval fixes the inverse branch. The remainder statement is direct substitution along the transformed endpoint; no output reconstruction is involved. $\square$

The theorem distinguishes two operations that are sometimes conflated. Changing the target of the equation does not amplify the error of its inverse observable. Changing the source variable and then reconstructing it generally does, as discussed below.

15.2 Exponential phases

Suppose an expression can be written on the selected real source branch as

$$F(x) = y_0 + \varepsilon e^{\Phi(x)}, \quad \varepsilon \in \{1, -1\}. \quad (15.1)$$

An eventually signed amplitude is included in the phase: if the original expression is $y_0 + A(x)e^{P(x)}$, prove the eventual sign ε of A , then set $\Phi = P + \log(\varepsilon A)$. The inverse equation becomes

$$\Phi(x) = v, \quad v = \log(\varepsilon(y - y_0)), \quad \varepsilon(y - y_0) > 0. \quad (15.2)$$

The phase may contain several powers and polynomial-logarithmic terms even when the original function has exponential growth or decay. It is inverted on the same source interval. If $\Phi \rightarrow +\infty$, then $F \rightarrow \varepsilon\infty$; if $\Phi \rightarrow -\infty$, the target approaches y_0 from the side with sign ε .

Proposition 15.2 (The original residual and the phase residual). *For a real approximation X on that source interval, put $D = \Phi(X) - v$. Then*

$$\frac{F(X) - y}{y - y_0} = e^D - 1. \quad (15.3)$$

In particular, if $D \rightarrow 0$, the relative original residual is $D + \mathcal{O}(D^2)$. If $|\Phi'| \geq \mu > 0$ between X and the true inverse $g(y)$, then $|X - g(y)| \leq |D|/\mu$.

Proof. Since $y - y_0 = \varepsilon e^v$, divide $\varepsilon(e^{\Phi(X)} - e^v)$ by εe^v . The remaining statements follow from the Taylor expansion of e^D at 0 and the mean value theorem for Φ . $\square$

Thus a phase residual can be numerically useful without forming enormous exponentials. A numerical solve in the phase remains numerical evidence unless the derivative and residual have rigorous bounds; it nevertheless refers to the original equation through the exact identity (15.3).

Example 15.3 (Several powers in an exponential). For $F(x) = e^{x^2+x}$ at $+\infty$, put $v = \log y$. The inverse is

$$x = \frac{-1 + \sqrt{1 + 4v}}{2} = \sqrt{v} - \frac{1}{2} + \frac{1}{8\sqrt{v}} - \frac{1}{128v^{3/2}} + \mathcal{O}(v^{-5/2}).$$

This is an ordinary inverse expansion in the phase variable v . Its local target variable is $1/v = 1/\log y$, and its remainder is measured in that coordinate. For e^{-x^2-x} at the same source endpoint, replace v by $-\log y$ and take $y \rightarrow 0^+$.

Example 15.4 (A nonconstant amplitude). For $F(x) = xe^{x^2+x}$ at $+\infty$, the phase is $\Phi = x^2 + x + \log x$. In terms of $v = \log y$, inversion begins

$$x = \sqrt{v} - \frac{1}{2} + \frac{\frac{1}{8} - \frac{1}{4}\log v}{\sqrt{v}} + \mathcal{O}(v^{-1}(1 + |\log v|)^2).$$

The logarithm of the amplitude contributes at the displayed order. Omitting it would preserve the leading term but give the wrong next coefficient. Translating x to $x - c$ translates the reconstructed inverse by c and leaves the phase calculation in that shifted source coordinate.

Example 15.5 (A flat forward function with an elementary inverse). $F(x) = e^{-1/x}$ on $x \rightarrow 0^+$ has no nonzero power-log forward expansion, but its inverse is exactly

$$x = -\frac{1}{\log y}, \quad 0 < y < 1.$$

The logarithmic target transformation turns the equation into $-1/x = \log y$. This example does not require an expansion into flat perturbation sectors: the exponential is part of the exact coordinate transformation.

Variable powers can enter the same route when their real domain is established. For x^x on $x > 0$, the phase is $x \log x$, so its inverse at $+\infty$ reduces to the Lambert-coordinate inversion of that phase. The identity $A(x)^{B(x)} = \exp(B(x) \log A(x))$ is used only where $A(x) > 0$.

15.3 Source reconstruction and its error

Theorem 15.6 (A change of source coordinate). *Let $x = H(t)$ be a continuously differentiable one-to-one parametrization of the selected source branch, and write $K(t) = F(H(t))$. Let h be its inverse and suppose*

$$h(y) = A(y) + E(y), \quad E(y) = \mathcal{O}(R(y)).$$

If H is defined on the interval between A and $A + E$, and $|H'| \leq B(y)$ there, then

$$F^{-1}(y) = H(A(y)) + \mathcal{O}(B(y)R(y)).$$

The same conclusion applies to an observable after replacing H by that observable composed with H .

Proof. The branch identity gives $F^{-1} = H \circ h$. Apply the mean value theorem between A and $h = A + E$. $\square$

For the coordinate transformations used here, more informative exact identities specify when a small error survives reconstruction.

Corollary 15.7 (Exponential reconstruction). *For a positive source coordinate $u = e^{st}$, $s \neq 0$, if $h = A + E$ and $E \rightarrow 0$, then*

$$u = e^{sA}(1 + \mathcal{O}(E)), \quad u^r = e^{rsA}(1 + \mathcal{O}(E)) \quad (15.4)$$

for every fixed real r . If E does not tend to zero, an additive expansion for h alone does not justify these relative error assertions.

Proof. $u = e^{sA}e^{sE}$, and $e^{sE} = 1 + \mathcal{O}(E)$ when $E \rightarrow 0$. Replace s by rs for the observable. $\square$

Corollary 15.8 (Logarithmic reconstruction). *For $x = -\log t/\lambda$, $\lambda \neq 0$, suppose the positive inner inverse satisfies $h = A + E$ with $A > 0$ and $E/A \rightarrow 0$. Then*

$$x = -\frac{\log A}{\lambda} - \frac{1}{\lambda} \log \left(1 + \frac{E}{A} \right) = -\frac{\log A}{\lambda} + \mathcal{O}(E/A). \quad (15.5)$$

Thus the inner remainder must be divided by the inner leading scale before it becomes an error in x .

Proof. Factor $h = A(1 + E/A)$ and expand the logarithm of the positive unit factor. The hypothesis $E/A \rightarrow 0$ also guarantees $h > 0$ eventually. $\square$

15.4 Pure logarithmic functions

For $x \rightarrow +\infty$, set $t = \log x \rightarrow +\infty$. For a positive source distance $u \rightarrow 0^+$, set $t = -\log u \rightarrow +\infty$, then reconstruct $u = e^{-t}$ and the original point $x = x_0 \pm u$. A pure logarithmic forward expression becomes a power-log expression in t . Its inverse is computed there before exponential reconstruction.

Example 15.9 (The constant term in an exponent matters). For $F(x) = (\log x)^2 + \log x$ at $+\infty$, the transformed equation is $t^2 + t = y$. Hence

$$t = \sqrt{y} - \frac{1}{2} + \frac{1}{8\sqrt{y}} + \mathcal{O}(y^{-3/2}),$$

and one valid reconstruction is

$$x = \exp\left(\sqrt{y} - \frac{1}{2} + \frac{1}{8\sqrt{y}}\right) (1 + \mathcal{O}(y^{-3/2})).$$

In contrast, $t = \sqrt{y} + \mathcal{O}(1)$ does not justify $x = e^{\sqrt{y}}(1 + o(1))$: the omitted constant would change the leading multiplicative factor by $e^{-1/2}$. An implementation must retain enough inner terms to make its omitted additive error tend to zero, even when the initial term request would otherwise stop earlier.

This precision requirement belongs to reconstruction, not to the mere existence of the inner asymptotic expansion. When a user requests a particular relative order in the reconstructed expression, the implementation must translate that request into an inner additive precision and record both meanings.

15.5 Finite sums of exponentials

Let

$$F(x) = \sum_{j=1}^m a_j e^{\mu_j x}, \quad a_j, \mu_j \in \mathbb{R},$$

with equal rates merged. At $x \rightarrow +\infty$ choose $\lambda > 0$, and at $x \rightarrow -\infty$ choose $\lambda < 0$, so $t = e^{-\lambda x} \rightarrow 0^+$ in either case. Then

$$F\left(-\frac{\log t}{\lambda}\right) = \sum_j a_j t^{-\mu_j/\lambda}. \quad (15.6)$$

This is a finite real-power germ. Rational relations among the rates are unnecessary: local finiteness of the positive correction gaps supplies finite exponent truncations. Its leading coefficient and exponent determine the target side, while the chosen sign of λ preserves the source endpoint. After inversion, Corollary 15.8 transports the error back to x .

Example 15.10 (An exact oracle). For $F(x) = e^x + e^{2x}$ at $+\infty$, take $t = e^{-x}$. Then $F = t^{-1} + t^{-2}$, and an exact inverse is

$$x = \log\left(\frac{\sqrt{1+4y}-1}{2}\right) = \frac{1}{2} \log y - \frac{1}{2\sqrt{y}} + \frac{1}{48y^{3/2}} + \mathcal{O}(y^{-5/2}).$$

The inverse of $e^x + e^{\sqrt{2}x}$ uses the same coordinate with irrational exponents. Constant terms, duplicate rates and cancellations are handled before extracting the leading transformed monomial.

15.6 Result contracts and verification

The transformed result retains its inner inverse, coordinate substitution, original function and selected domains. Its scale is identified as **"Transformed"**. The logarithmic target route has coordinate kind **"TargetLog"**; it uses (15.2). The target expression is recorded explicitly. The stored cutoff belongs to the underlying series coordinate after substitution. An ordinary **SeriesData** in the original target variable is unavailable when that representation would hide the coordinate change. The package README lists the corresponding property names.

The transformed residual describes the model actually composed, and its metadata states that scope. Numerical checks solve the transformed equation at high precision and verify the original source side; they remain numerical evidence. A separate certificate can use Proposition 15.2

with rigorous derivative and residual bounds. An additive input error for F must be transported through the target map before it can be treated as an error in Φ ; until that contract is supplied, rejecting an unsupported declaration is mathematically necessary.

These coordinate theorems justify the stated reductions independently of how many recognizers are implemented. The staged implementation plan tracks the accepted families, source reconstruction operations, tests and remaining extensions; the existence of a reduction alone does not establish that every corresponding expression has already passed the package’s parser and native test suite.

16 Incremental and grouped inverse computation

Three changes reduce repeated work without changing the coefficient theorem or the meaning of a complete power-log block. They concern the ordinary engine with finitely many positive exact gaps δ_i . Symbolic depth truncation continues to use the individual multi-index formula.

16.1 A persistent exact-weight frontier

Let I be the set of retained multi-indices, initially empty, and let its candidate frontier initially consist of the zero multi-index. At each step, select the smallest candidate weight σ , promote *every* candidate of that weight into I , and add all their one-step successors $\mathbf{k} + e_i$ to the candidate set. Candidates are deduplicated by their entire integer vector; their weights are canonicalized independently.

Proposition 16.1 (Complete layers and coefficient reuse). *After the layer at weight σ is processed, I is exactly $\{\mathbf{k} \in \mathbb{N}^m : \mathbf{k} \cdot \boldsymbol{\delta} \leq \sigma\}$, and the candidate set is its one-step outer boundary. Each retained multi-index has its coefficient evaluated exactly once. Summing the coefficients of an entire layer before counting a block preserves all resonances and exact cancellations.*

Proof. Every nonzero multi-index has a predecessor of strictly smaller weight because all gaps are positive. Inductively, every index of the next candidate weight has already been generated by a processed predecessor. A successor of an index in the current layer has strictly larger weight, so processing one member of the layer cannot expose a previously unseen member of the same layer. Deduplication prevents a second insertion through another predecessor. The coefficient of each weight is the complete finite sum in Theorem 6.1, whether or not that sum is zero. $\square$

The implementation returns its state by value. It stores the normalized model, assumptions, observable power, retained indices, candidate frontier, and merged blocks. Reusing a state therefore reuses coefficients only for the same model and assumptions. It does not use a global cache that might accidentally mix branches or parameter hypotheses. The resource budget counts distinct retained and candidate indices, as in the original bounded enumeration. Raising a term goal no longer recomputes coefficients of every earlier layer. A cancellation may still require arbitrarily many layers before another nonzero block appears; incremental enumeration does not provide a new termination theorem.

16.2 Newton iteration with growing precision

Write $E(U) = 0$ for (6.11), and let $\delta_* = \min_i \delta_i$. The zero jet is correct below δ_* . If U_s is correct strictly below weight s , one Newton step performed strictly below

$$h = \min(2s, H), \quad U_h = \left[U_s - \frac{E(U_s)}{E'(U_s)} \right]_{<h},$$

is correct strictly below h . Indeed, $E'(U_s)$ has nonzero constant coefficient p , and the Newton error is quadratic in the old error. Products, unit powers, and logarithmic composition do not decrease the weight of the error. The argument uses exponent valuation, with polynomial logarithms as coefficients; logarithmic degree need not stay bounded as h increases.

The working cutoffs are $\min(2\delta_*, H), \min(4\delta_*, H), \dots, H$, rather than H at every step. Their construction uses exact comparisons, with no floating-point estimate of the required iteration count. The implementation checks $[E(U_h)]_{<h} = 0$ after each step before recording the new precision. This is an exact formal consistency check. It is not a numerical error certificate for a particular nonzero value of z .

A reusable Newton state contains its unit jet and certified formal precision. Refinement to a larger cutoff starts from that jet; a request for a smaller cutoff leaves the stronger state intact, and a caller may truncate a copy. The monomial model has the exact unit jet zero and needs no Newton steps. Other models require a finite cutoff.

16.3 Combine resonances before applying Euler operators

Put

$$H_1(z, L) = \sum_i z^{\delta_i} B_i(L), \quad H_1(z, L)^n = \sum_{\sigma} z^{\sigma} A_{n,\sigma}(L).$$

The multinomial theorem gives

$$A_{n,\sigma} = \sum_{\substack{|k|=n \\ k \cdot \delta = \sigma}} \frac{n!}{k!} B^k.$$

At fixed depth n and weight σ , all summands of (6.1) have the same differential operator. By linearity, the unit bracket of g^r can therefore be computed as

$$\frac{g(y)^r}{z^r} = 1 + \sum_{n \geq 1} \sum_{\sigma} z^{\sigma} \frac{(-1)^{n_r}}{p^n n!} \prod_{j=1}^{n-1} (\mathcal{D} + r + \sigma + pj) A_{n,\sigma}(L). \quad (16.1)$$

For cutoff H , successive sparse products compute $[H_1^n]_{<H}$. Only n with $n\delta_* < H$ can contribute, so the computation is finite. Coefficients of equal weight are combined during each sparse product, before differentiation. Contributions at the same weight but different depths are then combined in the result; they too can cancel.

Equation (16.1) is especially useful when the generators have many additive relations. It does not require commensurable gaps, and it does not replace the public individual-coefficient formula: that formula retains the distinct perturbation markers and their provenance. Grouping can reduce coefficient work without eliminating the separate boundary enumeration used for remainder reporting. It is not a claim of uniform speedup on every input; polynomial coefficient growth and sparse multiplication remain material costs.

The independent engine is selected with `Method -> "GroupedLagrange"`. The default `"Lagrange"` method keeps the individual formula, so their agreement remains a direct check of the regrouping. Both use complete exact weights for term goals; grouping does not change which blocks are requested.

16.4 Validation obligations

The regression suite compares every successive incremental region with the independent bounded multi-index enumerator, checks exact algebraic collisions and complete cancellations, and verifies the unique-index resource budget. It compares growing-precision Newton with the individual Lagrange formula for incommensurable logarithmic gaps, including a negative leading power,

and checks that refinement preserves every previously certified coefficient. Grouped coefficients are compared with individual coefficients for several observable powers and with Newton independently. These are formal algebra checks; runtime measurements and analytic error bounds are separate evidence.

16.5 Recorded native measurements

The reproducible script `Tests/BenchmarkIncremental.wl` compares frozen reference control flow with the new helpers and checks exact result equality. For the public grouped comparison it checks the remainder descriptor as well. One fresh-kernel run on Wolfram 15.0.1 for Windows recorded:

Fixture	Reference seconds	Updated seconds
Eight logarithmic blocks, term goal	0.19092	0.05692
Newton through twelve weights	6.62615	3.35497
Seven-power model, target cutoff twelve	0.35947	0.48142

The first fixture reduces coefficient evaluations from 95 to 31. All three results agree exactly. The grouped public method is slower on this small fixture, where normalization, boundary enumeration and sparse products remain part of the cost. It stays an explicit alternative; these observations do not justify replacing the default method or asserting a uniform speedup.

17 Calculus of expansions with transported precision

The operations `SeriesAdd`, `SeriesMultiply` and `SeriesPower` provide arithmetic without discarding remainders. The functions `SeriesLog` and `SeriesExp` handle logarithms and exponentials; `SeriesCompose` and `SeriesObservable` handle composition and other observables. Truncation, refinement and justified differentiation use `SeriesTruncate`, `SeriesRefine` and `SeriesDifferentiate`. All ten operations share a common representation

$$F(y) = b(y) + c(y) \left[J(w(y)) + \mathcal{O}(w(y)^P M(w(y))^D) \right], \quad w(y) \longrightarrow 0^+, \quad (17.1)$$

where b, c, w are exact expressions and J is a finite power–log jet. The offset b and prefactor c are recorded separately. Thus a cutoff in the jet describes relative precision with respect to c , and the absolute remainder contains $|c|$. Ordinary forward and inverse objects, Lambert objects, and explicitly recorded coordinate substitutions can be converted to this representation. An operation returns a derived object and its replay recipe; it does not retain an inverse-model label that would falsely suggest that the original inverse residual checker applies to an arbitrary observable.

For addition and multiplication, the precision rules of Section 2 are used after expressing both operands in the same positive coordinate. An exact prefactor may be absorbed into the jet when it is a finite power–log expression. Distinct prefactors can be aligned when their ratio lies in that algebra. The package rejects incompatible coordinates or independently ordered exponential carriers rather than assigning them an unsupported single cutoff.

Powers and logarithms. For a positive monomial leading term aw^α , write $J = aw^\alpha(1 + U)$. A constant power and a logarithm reduce to the unit series of $(1 + U)^r$ and $\log(1 + U)$. The input uncertainty first becomes $\mathcal{O}(w^{P-\alpha}M^D)$ in U . Multiplication by the output prefactor then transports it to the correct absolute scale. Real noninteger powers require positivity on the selected branch. These operations retain any earlier input precision ceiling; requesting a higher output cutoff does not create unknown coefficients.

Exponentials and exact carriers. Suppose

$$H = H_{\leq 0} + H_{> 0} + \mathcal{O}(w^P M^D), \quad P > 0,$$

where $H_{\leq 0}$ contains the nonpositive-weight blocks. Then

$$e^H = e^{H_{\leq 0}} \left[e^{H_{> 0}} + \mathcal{O}(w^P M^D) \right].$$

Indeed the omitted argument tends to zero and $e^r = 1 + \mathcal{O}(r)$ there. The operation keeps $e^{H_{\leq 0}}$ as an exact carrier and expands only the vanishing part. For example $H = w^{-1} + \log w + w$ yields $we^{1/w}(1 + w + w^2/2 + \dots)$. This also supplies the precision rule needed when a source logarithm is inverted and the source variable must be reconstructed by exponentiation. A nonvanishing argument uncertainty is rejected: it cannot justify a relative error tending to zero after exponentiation.

Composition and regular observables. The outer local coordinate must become $aw^\alpha(1 + o(1))$ with $a, \alpha > 0$ after substitution of the inner expansion. This condition checks both the limiting point and its selected side. An outer remainder $\mathcal{O}(v^P M(v)^D)$ becomes $\mathcal{O}(w^{\alpha P} M(w)^D)$, since $\log v = \alpha \log w + \mathcal{O}(1)$. The uncertainty of the inner argument is propagated independently through every jet operation, and the worse precision is retained. Regular analytic observables at a finite limiting argument use their Taylor series and the same precision arithmetic. Poles, branch points and logarithmic leading coefficient functions require their respective scale normalizations; a general analytic-observable call does not silently assume them to be regular.

Differentiation requires derivative information. An estimate $R(w) = \mathcal{O}(w^P M(w)^D)$ alone does not imply any estimate on R' . For example $w^P \sin(e^{1/w})$ has that magnitude bound and a much larger derivative. Therefore a nonzero remainder can be differentiated n times only with a recorded or explicitly supplied contract

$$R^{(j)}(w) = \mathcal{O}(w^{P-j} M(w)^D), \quad 0 \leq j \leq n.$$

The option `"RemainderDerivativeOrder" -> n` is a declaration of these hypotheses, not a numerical test of them. Each differentiation consumes one order of the contract. The jet derivative uses $D_w[w^\beta Q(\log w)] = w^{\beta-1}(\beta Q + Q')$; the derivatives of the exact coordinate, offset and prefactor in (17.1) are included by the chain and product rules. Exact finite expressions require no remainder hypothesis. A discarded tail of a known finite power-log expression also has derivative bounds of every order. In contrast, a derivative bound declared for one unknown remainder is not automatically promoted to a stronger exponent when the source is refined; that stronger bound must be established separately.

Truncation and refinement. Truncation moves the first discarded block, including its logarithmic degree, into the remainder. It cannot improve uncertainty inherited from the input. Refinement instead replays a retained forward expression, inverse problem, coordinate substitution or operation recipe at a higher working order. The replay still respects transported input precision ceilings. Without a retained source or recipe, the package reports that refinement is unavailable; changing an inert remainder exponent would not be a mathematical refinement.

An automatically computed forward remainder is distinct from a declared input error. The object records the latter separately as `DeclaredInputRemainder`. When refining an analytic input such as $\sin x$, the original function can be expanded further; its previous Taylor remainder must not become a permanent declared accuracy ceiling. A remainder explicitly supplied for an unspecified omitted function remains a binding ceiling. Tests exercise both cases with independent inverse coefficients.

18 Perturbations of a retained exact inverse

An inverse such as a real branch of the Lambert function can be useful as an exact building block even when it also has a convergent inverse–logarithmic expansion. Retaining that building block separates two different refinements: expansion of the core inverse itself, and correction for terms outside the core. The operation `AsymptoticCoreInverse` implements the second refinement. Its truncation parameter counts powers of an auxiliary perturbation marker. The returned coefficients are complete functions of the exact core inverse; they are not individual terms of an exponent-sorted power–logarithmic series.

18.1 The coefficient identity

Choose the positive source coordinate $u \rightarrow 0^+$ and write the original source point as

$$x = H(u), \quad H(u) = x_0 + \sigma u \quad \text{or} \quad H(u) = \frac{\sigma}{u}, \quad \sigma \in \{1, -1\}.$$

Let $F_0(u)$ be the chosen core and $R(u)$ its perturbation in that coordinate. Let $u_0(y)$ be the selected exact inverse of F_0 , and set $\phi(y) = H(u_0(y))$. The option `"Power" -> r` uses the same observable convention as the ordinary inverse engine: $r = 1$ returns the source point itself; at a finite endpoint, $r \neq 1$ returns its displacement raised to r ; at an infinite endpoint it returns its r th power. Thus the local observable is

$$G_r(u) = \begin{cases} x_0 + \sigma u, & r = 1 \text{ at a finite endpoint,} \\ \sigma^r u^r, & r \neq 1 \text{ at a finite endpoint,} \\ \sigma^r u^{-r}, & \text{at an infinite endpoint.} \end{cases} \quad r_* = \begin{cases} r, & \text{finite endpoint,} \\ -r, & \text{infinite endpoint.} \end{cases} \quad (18.1)$$

The exponent r is an exact nonzero real number. For $\sigma = -1$, the operation requires integer r ; for $\sigma = 1$, all displayed bases are positive on the real source branch. Put $\Phi_r(y) = G_r(u_0(y))$; in particular $\Phi_1 = \phi$. Introduce a marker λ by

$$F_0(u_\lambda) + \lambda R(u_\lambda) = y, \quad u_{\lambda=0} = u_0(y). \quad (18.2)$$

At any fixed point with $F'_0(u_0) \neq 0$, the inverse function theorem gives an analytic solution in λ near zero whenever the data are analytic locally. The Lagrange–Bürmann identity gives

$$G_r(u_\lambda) = \Phi_r(y) + \sum_{n \geq 1} \lambda^n C_n(u_0(y)), \quad (18.3)$$

where

$$C_n(u) = \frac{(-1)^n}{n!} \left(\frac{1}{F'_0(u)} \frac{d}{du} \right)^{n-1} \left(\frac{G'_r(u) R(u)^n}{F'_0(u)} \right). \quad (18.4)$$

Indeed, with $t = F_0(u)$, the equation becomes $t = y - \lambda R(u_0(t))$, and $d/dt = (F'_0(u))^{-1} d/du$. Applying the ordinary Lagrange identity to the observable $\Phi_r(t)$ proves (18.4). This form of the identity differentiates the explicit forward core. It avoids repeated symbolic differentiation of a large expression containing `ProductLog`.

For arbitrary data, this identity alone does not justify setting $\lambda = 1$, does not order the coefficients by their size as $u_0 \rightarrow 0$, and does not bound the tail by its first omitted coefficient. The following restricted class supplies those missing conclusions with a complete tail majorant.

18.2 A class with an asymptotic tail bound

Let $M(u) = 1 + |\log u|$. Consider finite real power–logarithmic expressions

$$F_0(u) = y_0 + u^p Q_0(\log u) + \sum_{j=1}^J u^{p+\eta_j} Q_j(\log u), \quad p \neq 0, \quad \eta_j > 0, \quad (18.5)$$

$$R(u) = \sum_{i=1}^I u^{p+\delta_i} P_i(\log u), \quad \delta_i > 0. \quad (18.6)$$

All polynomials have real coefficients, Q_0 is nonzero, and all powers are fixed real numbers. Equal powers are merged before the leading term is selected. Put

$$q = \deg Q_0, \quad \delta = \min_i \delta_i, \quad D = \max \left(0, \max_i \deg P_i - q \right), \quad \chi(u) = u^\delta M(u)^D. \quad (18.7)$$

For a zero perturbation the observable remains exactly Φ_r , so these parameters and the following tail estimate are unnecessary.

Theorem 18.1 (An exact-core expansion with a complete marker tail). *For fixed data satisfying (18.5)–(18.6), the selected real inverse exists sufficiently near the source endpoint. Fix an observable satisfying (18.1), and write $g_r(y) = G_r(u_1(y))$ for that observable of the inverse. There are positive constants C, K and a source threshold such that, for every integer $N \geq 0$ and all sufficiently small $u_0 > 0$ with $K\chi(u_0) < 1$, the marker series converges at $\lambda = 1$ and*

$$\left| g_r(y) - \Phi_r(y) - \sum_{n=1}^N C_n(u_0(y)) \right| \leq C u_0^{r_*} \frac{(K\chi(u_0))^{N+1}}{1 - K\chi(u_0)}. \quad (18.8)$$

In particular, for every fixed N ,

$$g_r(y) = \Phi_r(y) + \sum_{n=1}^N C_n(u_0(y)) + \mathcal{O}\left(u_0^{r_* + (N+1)\delta} M(u_0)^{(N+1)D}\right). \quad (18.9)$$

The estimate concerns the complete omitted marker tail, including possible cancellations in its first coefficient.

Proof. Write $u = u_0 W$ and divide the equation $F_0(u_0 W) + \lambda R(u_0 W) = F_0(u_0)$ by the leading scale $u_0^p (\log u_0)^q$ and its nonzero leading coefficient. Use $\tau = 1/\log u_0$ and finitely many additional parameters for the higher-power terms. For example,

$$\frac{(\log(u_0 W))^k}{(\log u_0)^q} = (\log u_0)^{k-q} (1 + \tau \log W)^k.$$

If $k > q$, include the factor $u_0^\eta (\log u_0)^{k-q}$ in its small parameter; if $k \leq q$, retain the nonnegative power τ^{q-k} as an analytic coefficient. All such parameters tend to zero. Powers of W and $\log W$ are analytic in a fixed complex neighborhood of $W = 1$ after fixing their branches.

At the limiting parameter origin, the normalized equation is $W^p - 1 = 0$, whose derivative at $W = 1$ is $p \neq 0$. The analytic implicit function theorem therefore gives a holomorphic solution in a common polydisc. The core parameters remain fixed when λ varies, whereas every perturbation parameter is multiplied by λ . Each perturbation parameter is bounded in magnitude by a constant times $\chi(u_0)$ for small u_0 . Cauchy estimates in a smaller fixed polydisc consequently bound the total coefficient of marker degree n by $CK^n \chi(u_0)^n$. Equivalently, the solution is holomorphic and uniformly bounded on a marker disc whose radius is a fixed positive multiple of $1/\chi(u_0)$.

The observable in (18.1), after subtracting the constant x_0 when it occurs, is $\sigma^r u_0^{r*} W^{r*}$. The normalized factor W^{r*} is analytic near $W = 1$ for every fixed real r_* , including negative and irrational values, giving the additional factor u_0^{r*} . The branch guard ensures that this analytic continuation agrees with the stated real observable. Summing the coefficient bound gives (18.8). The local real branch is unique, and $F'_0(u) \sim pu^{p-1}Q_0(\log u)$ while $R'/F'_0 \rightarrow 0$, so this analytic solution agrees with the selected real inverse near the endpoint. Finally $\chi(u_0) \rightarrow 0$, which proves (18.9). $\square$

The constants in this theorem are existence bounds for fixed input data. The implementation does not compute a numerical value of C , K , or the source threshold. Its "MajorantContract" therefore identifies the bound as asymptotic existence and sets "NumericCertificate" to **False**. The stated "SourceRadius" controls a supplied inverse's symbolic branch check; it is not asserted to be the unknown uniform threshold in (18.8).

The nonnegative value of D in (18.7) is deliberately conservative. A quotient by an extra power of $|\log u_0|$ can make an individual coefficient smaller than the reported polynomial-log bound. Keeping that improvement would require an additional logarithmic ordering contract. Also, if the perturbation contains several different powers, one complete marker coefficient may contain contributions lying beyond the leading contribution of a later marker coefficient. Marker depth does not promise a globally sorted list of all resulting monomials.

18.3 The Lambert core and its first correction

For the exact core

$$F_0(u) = y_0 + au^p(\log u + b)^q, \quad p \neq 0, \quad q \in \mathbb{Z}_{\geq 0},$$

the parser verifies that the leading logarithmic polynomial is exactly the displayed affine power. For $q = 0$, the positive local inverse is $u_0 = ((y - y_0)/a)^{1/p}$. For $q > 0$, put

$$A = a(-1)^q, \quad Y = \frac{y - y_0}{A} > 0, \quad k = -\frac{p}{q}, \quad z = kY^{1/q}e^{pb/q}.$$

Then

$$u_0 = Y^{1/p} \left(\frac{W_\nu(z)}{k} \right)^{-q/p}, \quad \nu = \begin{cases} -1, & p > 0, \\ 0, & p < 0. \end{cases} \quad (18.10)$$

These are the real branches for which $u_0 \rightarrow 0^+$. On the lower branch, the real argument condition $-1/e \leq z < 0$ is recorded explicitly. The quotient $W_\nu(z)/k$ is positive on the selected eventual branch, so all displayed real powers have positive bases.

Example 18.2 (Correcting $u \log u$ without expanding its exact inverse). For $F(u) = u \log u + u^2$ as $u \rightarrow 0^+$, let $u_0 = y/W_{-1}(y)$. Formula (18.4) gives

$$C_1(u_0) = -\frac{u_0^2}{1 + \log u_0}, \quad (18.11)$$

$$C_2(u_0) = \frac{u_0^3(4 \log u_0 + 3)}{2(1 + \log u_0)^3}. \quad (18.12)$$

Here $p = q = \delta = 1$ and $D = 0$. Thus truncation after C_1 has the proved remainder $\mathcal{O}(u_0^3)$, and truncation after C_2 has remainder $\mathcal{O}(u_0^4)$. The explicit omitted coefficient can exhibit a sharper logarithmic scale, but the package reports the complete-tail bound established above. Each displayed correction retains the exact W_{-1} expression in u_0 .

18.4 A divergent power plus a logarithm

The exact core may itself contain smaller source powers and logarithms. A further automatically recognized class is

$$F_0(u) = y_0 + au^p + b \log u, \quad p < 0, \quad ab \neq 0. \quad (18.13)$$

It belongs to (18.5): the logarithm has the strictly higher source exponent zero. The leading logarithmic degree is $q = 0$. Set $X = u^p$, $d = b/p$, and $k = a/d$. The equation becomes $aX + d \log X = y - y_0$, whence

$$X = \frac{W_\nu(k e^{(y-y_0)/d})}{k}, \quad u_0 = X^{1/p}, \quad \nu = \begin{cases} 0, & k > 0, \\ -1, & k < 0. \end{cases} \quad (18.14)$$

Indeed $kX e^{kX} = k e^{(y-y_0)/d}$; conversely, taking real logarithms of the positive quantities in this identity recovers the core equation. These branches give $X \rightarrow +\infty$, hence $u_0 \rightarrow 0^+$. The signs of the coefficients, positivity of the local inverse, and the lower branch's real argument interval are part of the branch contract. A constant in the zero-power logarithmic polynomial is extracted into y_0 before this recognition, so affine target shifts are preserved.

Example 18.3 (An exact core for $x + \log x + 1/x$). As $x \rightarrow +\infty$, retain the exact inverse $w = W_0(e^y)$ of $x + \log x$. The source coordinate is $u = 1/x$; the core is $u^{-1} - \log u$ and the perturbation is u . Thus $p = -1$, $q = 0$, $\delta = 2$, $D = 0$, and $u_0 = 1/w$. For the default observable $r = 1$, formula (18.4) gives

$$x = w - \frac{1}{w+1} - \frac{2w+1}{2w(w+1)^3} + \mathcal{O}(w^{-5}). \quad (18.15)$$

The two corrections are complete marker coefficients; the exact $W_0(e^y)$ has not been expanded into logarithms. For the reciprocal observable, $r = -1$ and $r_* = 1$, and the corresponding result is

$$\frac{1}{x} = \frac{1}{w} + \frac{1}{w^2(w+1)} + \frac{4w+3}{2w^3(w+1)^3} + \mathcal{O}(w^{-7}). \quad (18.16)$$

The differing exponents 5 and 7 follow from the same complete-tail theorem with the appropriate observable derivative. A direct expansion of $X - w + \log(X/w) + \lambda/X$ verifies the first result through marker degree two independently of the coefficient generator.

18.5 A declared forward remainder

Suppose the actual forward function also contains an unknown error $E(u)$, with the declared bounds

$$E(u) = \mathcal{O}(u^\rho M(u)^k), \quad E'(u) = \mathcal{O}(u^{\rho-1} M(u)^k), \quad \rho > p. \quad (18.17)$$

The derivative bound is a hypothesis of this declaration, not a consequence of the first bound alone. Since the retained forward derivative has scale $u^{p-1} M(u)^q$, the mean value argument of Section 9 gives an additional error in the selected observable of

$$\mathcal{O}\left(u_0^{r_*+\rho-p} M(u_0)^{\max(0, k-q)}\right). \quad (18.18)$$

The gap $\rho > p$ ensures that this error is relatively small in the positive source coordinate. Both the perturbed solution and the core solution are asymptotic to the same u_0 , so they can be used interchangeably in these scale estimates. The output combines (18.18) with the marker-tail bound by taking the smaller source power, and the larger logarithmic degree if those powers coincide. Increasing marker depth cannot improve this inherited accuracy ceiling. The displayed marker coefficients still describe the specified finite model; coefficients smaller than the unknown

input error are not asserted to be coefficients of every admissible actual function. The transport formula remains valid when $r_* < 0$ or when the resulting absolute error does not tend to zero: the forward gap $\rho > p$ still makes the error small relative to the observable's leading scale $u_0^{r_*}$. For example, at $x \rightarrow 0^+$, a declared $\mathcal{O}(x^3 M(x)^2)$ error in the model $x + x^2$ gives an $\mathcal{O}(M(u_0)^2)$ ceiling for the observable x^{-2} . At $x \rightarrow +\infty$, a declared $\mathcal{O}(x^{-2} M(1/x))$ error in the model $x + 1/x$ gives an $\mathcal{O}(u_0^4 M(u_0))$ ceiling for x^{-1} .

18.6 The implemented operation and its evidence boundary

```
AsymptoticCoreInverse[
  x Log[x], x^2, {x, 0}, {y, 2}]

AsymptoticCoreInverse[
  x + x^2, x^3, {x, 0}, {y, 1},
  "CoreInverse" -> (Sqrt[1 + 4 y] - 1)/2]

AsymptoticCoreInverse[
  x + Log[x], 1/x, {x, Infinity}, {y, 2},
  "Power" -> -1]

AsymptoticCoreInverse[
  7 + 3 (x - 2) Log[5 (x - 2)], 4 (x - 2)^2,
  {x, 2}, {y, 2}, "Power" -> -2]
```

The first call automatically selects the lower real Lambert branch. The second retains a supplied algebraic inverse after a bounded symbolic check of $\phi(F_0(H(u))) = H(u)$ for $0 < u < \text{"SourceRadius"}$ on the selected side. An unverified inverse is rejected. Supplied parameter assumptions must establish the reality of the coefficients and the eventual leading sign. Exponents in this operation are exact real numeric constants; they need not be rational. The third call returns the reciprocal observable in (18.16); the fourth returns the reciprocal square of $x - 2$ while retaining the exact inverse of the scaled Lambert core. It does not raise the translated source point to that power.

The result identifies its scale as **"ExactCorePerturbation"**. Its kind is **"CoreInverse"**, and **"MarkerTerms"** stores pairs $\{n, C_n\}$, with degree zero holding Φ_r . The full exact inverse of the core, the local core inverse, the separately computed first omitted marker coefficient, the branch certificate, and the tail contract are all retained. The **"CoreInverse"** property remains ϕ for every observable; **"CoreObservableExpression"** holds Φ_r and **"LocalObservablePower"** holds r_* . The ordinary **Normal** operation returns the retained expression. The object does not claim an ordinary **SeriesData** representation or an ordinary exponent cutoff. Its independent regressions compare marker coefficients with direct composition and exact algebraic inverses, check both source orientations, shifted endpoints, real powers, and observable-dependent input-error ceilings, and compare the Lambert corrections with the original forward equations at high precision. Such numerical checks supplement the theorem; they do not turn its unspecified constants into a numerical certificate.

This implemented class requires strictly positive source-power gaps in the perturbation. Perturbations of equal source power but smaller logarithmic size, purely logarithmic leading cores, and general exponential exact cores require additional scale and error-transport contracts. They are not silently accepted by the marker identity. The broader roadmap also includes refinement of the retained core, composition of core objects with other expansion carriers, and computable majorants. The present operation supplies the finite power-logarithmic perturbation theorem and its exact-core implementation as one independently testable stage.

19 Executable residual certificates

An asymptotic remainder describes a limiting scale; it need not supply a constant or a numerical neighborhood. A root obtained by numerical iteration is useful independent evidence, but is not a verified enclosure. The function `InverseCertificate` implements the residual bracket of Theorem 9.4 on an explicit source interval. Its successful output contains rational endpoints enclosing a unique root, a bound for the error of the reported center, and the derivative and residual enclosures used to prove those statements.

19.1 The certificate theorem with interval data

Let $I = [A, B]$, let $c \in (A, B)$, and suppose a structural interval evaluator proves that F is continuous on I and differentiable inside, with

$$F'(I) \subseteq [d_-, d_+], \quad F(c) - y \in [\rho_-, \rho_+].$$

Assume either $d_- > 0$ or $d_+ < 0$, and put

$$\mu = \min(|d_-|, |d_+|), \quad M = \max(|d_-|, |d_+|), \quad \epsilon = \max(|\rho_-|, |\rho_+|), \quad h = \epsilon/\mu.$$

If $[c - h, c + h] \subseteq I$, Theorem 9.4, after changing the sign of F if necessary, proves a unique root x_* in I and

$$|x_* - c| \leq h.$$

The mean-value identity further gives

$$x_* \in I \cap [c - h, c + h] \cap \left(c - \frac{[\rho_-, \rho_+]}{[d_-, d_+]} \right).$$

Division is interval division; its denominator is separated from zero. This last interval is the returned "`RootEnclosure`". If the residual interval does not contain zero, there is also the lower bound

$$|x_* - c| \geq \frac{\min(|\rho_-|, |\rho_+|)}{M}.$$

It can prove that a requested accuracy is impossible for a fixed center.

If the symmetric residual bracket does not fit inside I , the evaluator also encloses the residual at A and B . Proved opposite signs in the orientation of F' establish existence by continuity, and the separated derivative still establishes uniqueness. Once existence in I is known, the mean-value theorem gives the same bound $|x_* - c| \leq h$ and the same intersection enclosure, without requiring the symmetric bracket to be contained in I . The output identifies this alternative as exact endpoint sign evidence. This case matters for a coarse asymptotic seed close to one end of a valid interval: increasing arithmetic precision cannot repair an unsuitable symmetric-bracket geometry.

The theorem asserts uniqueness on the supplied interval, which the evaluator also checks lies on the requested side of the source endpoint. It does not certify global injectivity or identify disconnected monotone intervals as the same continuation of an inverse germ. This scope is recorded in the result.

19.2 Rational arithmetic and elementary enclosures

The proof path uses exact rational endpoints. Addition, multiplication, reciprocals, and integer powers use the usual enclosing interval operations. After arithmetic, an endpoint $q \neq 0$ is rounded outward to a dyadic grid with a prescribed number of significant bits. Its magnitude

exponent is determined from the binary lengths of the exact numerator and denominator, followed by one rational comparison. Rounding uses $\lfloor q/\eta \rfloor \eta$ at a lower endpoint and $\lceil q/\eta \rceil \eta$ at an upper endpoint. Thus even rounding is justified by integer arithmetic. Approximate decimal values and floating-point comparisons are absent from every enclosure and every final sign or containment test.

For $0 \leq t \leq \frac{1}{2}$, the exponential evaluator uses

$$S_N(t) = \sum_{j=0}^N \frac{t^j}{j!}, \quad 0 \leq e^t - S_N(t) \leq \frac{t^{N+1}}{(N+1)!} \frac{1}{1 - t/(N+2)}.$$

The bound follows by majorizing the ratio of successive omitted terms by $t/(N+2)$. A positive argument is reduced by powers of two, and the enclosure is squared back using directed rounding. Negative arguments use the reciprocal of the positive exponential enclosure. An explicit resource bound on the argument prevents unbounded integer growth.

For $1 \leq q \leq 2$, put $v = (q-1)/(q+1) \in [0, \frac{1}{3}]$. Then

$$T_N(q) = 2 \sum_{j=0}^{N-1} \frac{v^{2j+1}}{2j+1}, \quad 0 \leq \log q - T_N(q) \leq \frac{2v^{2N+1}}{(2N+1)(1-v^2)}.$$

Every positive rational argument is written exactly as $2^k q$ with $1 \leq q < 2$, using binary numerator and denominator lengths. The identity $\log(2^k q) = k \log 2 + \log q$ therefore reduces all logarithms to the displayed bound. Both tails are positive rational expressions. Increasing N improves their analytic precision independently of the directed arithmetic precision.

Before enclosing a logarithm, exact structural identities can remove an exponential whose real exponent has been enclosed, or split a product whose factors each have a real logarithm enclosure. Positivity is therefore checked before distributing logarithms. If, for example, a positive product has negative factors, the evaluator retains the direct interval route. This avoids materializing a huge exponential only to take its logarithm while preserving the real-branch conditions.

Noninteger powers on a strictly positive base are evaluated through $a^b = \exp(b \log a)$, including supported exact irrational exponents. A positive rational power also admits a zero lower base endpoint. Expressions are handled by structural recursion through sums, products, powers, exponentials and logarithms. A reciprocal crossing zero, a logarithm reaching a nonpositive argument, or an unsupported head produces a failure. In particular, a numerical value of an unsupported expression is never accepted as a proof. Successful structural evaluation establishes the needed continuity on the whole interval; successful evaluation of the symbolic derivative establishes the derivative enclosure. Symbolic parameters without exact supported values are currently rejected with the unresolved expression recorded.

19.3 Interface, phase coordinates, and adaptive accuracy

The basic call is

```
s = AsymptoticInverse[x + x^2, {x, 0}, {y, 3}];
InverseCertificate[s, 6, "Interval" -> {1, 3},
"Center" -> 21/10]
```

The independently known root is 2. The derivative enclosure is $[3, 7]$; the certified bracket encloses 2 even though the asymptotic expansion was not requested at a target where its truncation is especially accurate. The residual theorem is independent of the formal truncation theorem.

Targets must be exact supported numeric expressions. Interval endpoints and an explicitly supplied center must be rational, with the center strictly inside the interval. With **"Center"** -> **Automatic**, numerical evaluation of the expansion chooses a rational center; an unusable seed is replaced by the interval midpoint. This decimal calculation selects a candidate only. All proofs are subsequently rebuilt with exact rational arithmetic. The returned error bound applies to that recorded center, including its rationalization. It does not silently claim an error bound for a different displayed expression.

The equation being certified is always the object's stored explicit **"Function"**. A declared **"InputRemainder"** describes further possible forward functions, but supplies no computable interval bound for their values or derivatives. The certificate therefore does not cover those unspecified terms. Results retain that descriptor, set **"CertifiesInputRemainderFamily"** -> **False**, and label the scope **"StoredExpressionOnly"** when it is nontrivial. An automatically computed forward truncation can also populate the remainder field, so this label does not guess the descriptor's provenance. The original function and exact target are recorded separately from any phase equation. Asymptotic remainder data are never silently promoted into a numerical error certificate.

For an object using a logarithmic target transformation, or a recognized Lambert exponential core admitting the same exact transformation, the certificate evaluates the exact phase and transformed target. The source-domain condition of the coordinate map is proved on the entire interval first. Since the exponential is strictly monotone there, the resulting root is a root of the original equation. This avoids constructing enormous values such as $\exp(x^2 + x)$ merely to subtract two nearly equal numbers.

The option **"TargetError"** -> **epsilon** requests an exact positive rational tolerance. After an initial certificate, a sharper interval-Newton enclosure can supply a new center and verification interval. The preceding certificate already proves a root exists in that interval, so subsequent error estimates need not re-prove existence by the initial bracket-containment condition. Derivative bounds and residual bounds are nevertheless recomputed. Unless disabled by **"RefineExpansion"** -> **False**, refinement also tries a larger asymptotic term goal as a source of a better seed; every accepted seed is verified independently. A bounded failed or unsupported expansion refinement does not invalidate an already established interval theorem.

An **AsymptoticCoreInverse** object is also accepted as a seed. Its exact Lambert core and marker corrections may be evaluated numerically to select the center, while the certificate evaluates the full explicit **"Core"** + **"Perturbation"** equation with rational enclosures. If expansion refinement is requested, the stored exact core inverse is replayed at a larger marker depth. Its symbolic identity check and computation are time bounded; a failed replay falls back to the already certified interval information. The marker majorant's uncomputed constants are not used as numerical bounds.

With an explicit fixed center, refinement increases enclosure precision without silently changing the center. A certified lower error bound above the target causes **Failure["AccuracyFloor", ...]**. Exhausting the refinement budget causes **Failure["AccuracyNotReached", ...]**, retaining the best established certificate if one exists. Neither outcome is reported as a successful accuracy request. The **"History"** property records the attempted centers and enclosure orders. Resource failures and unresolved derivative signs are reported separately from mathematical impossibility.

Relative accuracy. An optional **"RelativeError"** -> **tau** requests a relative tolerance with an independently proved root magnitude. If the root interval I is separated from zero, set $m_I = \min_{z \in I} |z|$ and $M_I = \max_{z \in I} |z|$. An absolute error bound r proves the requested

combined tolerance when

$$r \leq \max(\epsilon, \tau m_I),$$

where ϵ is an explicitly supplied **TargetError**, or zero when none is supplied. This implies $|A - x_*| \leq \max(\epsilon, \tau|x_*|)$ for the unique enclosed root. The reported relative error bound is r/m_I . A proved zero root requires an explicit positive absolute fallback; no relative quotient at zero is invented. If an interval merely crosses zero, further refinement can establish a nonzero lower magnitude or report that the goal was not reached. For a fixed-center impossibility claim, the opposite bound is needed: the certified lower error must exceed $\max(\epsilon, \tau m_I)$. Using m_I in that last test would incorrectly turn a conservative success criterion into an impossibility assertion.

19.4 Validation scope

The regression cases exercise outward rounding at both signs and widely different magnitudes, explicit exponential and logarithmic bounds, interval powers crossing zero, positive and negative derivatives, and a zero residual. Known rational roots provide independent checks of the final enclosures. Further cases cover the lower Lambert branch, a logarithmic phase, successful adaptive accuracy, a proved fixed-center accuracy floor, and rejection of rounded targets, turning points, poles, invalid logarithm domains, and unsupported heads. These tests validate an executable exact-arithmetic implementation of the theorem; they do not constitute a separate formal verification of the Wolfram evaluator.

The package's two motivating functions $x + x^2(1 + \log x)$ and $x + x^{\sqrt{2}}$ are included, together with a nonunit leading power and coefficient, both real Lambert branches, and exact-core perturbations. A supplied interval immediately below $x = e^{-1}$ tests the small derivative near the turning point of $x \log x$. There the reported derivative lower bound is small but strictly positive in absolute value; the result remains an interval theorem, without relying on a logarithmic asymptotic truncation being accurate near that threshold.

20 Implemented source charts and reconstruction checks

The implementation of Section 15 lives in `Kernel/SourceCoordinates.wl`. The two chart families are distinguished by the function used to reconstruct the source: **"SourceExp"** reconstructs an exponential, while **"SourceLog"** reconstructs a logarithm. The underlying inverse always has observable power 1. Powers of the original inverse are formed only after that chart variable has been computed.

Pure logarithmic source dependence. First choose the positive local source coordinate u from the endpoint and direction. Substitute $u = e^{-h}$, so $h \rightarrow +\infty$. A pure logarithmic expression becomes a function $\Phi(h)$ with no residual exponential dependence on h , and the existing inverse dispatcher solves $\Phi(h) = y$. For a finite endpoint x_0 and side $s \in \{-1, 1\}$,

$$x = x_0 + se^{-h};$$

at the source infinity with sign s , the reconstruction is $x = se^h$. Thus the noninteger real powers $(-\log u)^q$, polynomial and rational combinations of $\log u$, and both finite source directions can use the same chart whenever the transformed inverse belongs to a supported scale.

For example, $F(u) = (\log u)^2 + \log u$ becomes $h^2 - h$, and its small positive inverse begins

$$u = e^{-\sqrt{y}-1/2} \left(1 - \frac{1}{8\sqrt{y}} + \mathcal{O}(y^{-1}) \right).$$

The exponential factor multiplies the absolute remainder. A finite-order approximation to h is usable here only when its absolute error tends to zero. A finite relative Lambert expansion for a growing h need not have this property; increasing its logarithmic depth alone does not resolve that obstruction.

An exact transformed inverse may instead be kept as an exact carrier. Repeated exact charts therefore give, for instance, the exact inverse $u = \exp(-e^y)$ of $\log(-\log u)$. Recursion is bounded to avoid an uncontrolled recognizer loop. Exactness is recorded separately: it can come from an exact inner series, an exact inner inverse, or a bounded symbolic check that a small candidate composes exactly to the target. The asymptotic branch of the inner inverse fixes which exact solution is being reconstructed.

Finite exponential families. At $x \rightarrow s\infty$, the implementation takes $u = e^{-sx}$, so $x = -s \log u$. Finite sums of real exponential rates become finite generalized-power germs in u . Equal rates and cancellations are handled by the ordinary canonicalization before a leading term is selected. Rates need not be commensurable. The recognizer requires a nonzero leading algebraic exponent after removing a constant target offset; additional logarithmic factors are permitted when the resulting finite power-log germ is supported by the inner engine.

For example,

$$e^{-x} + e^{-2x} = y \implies x = -\log y + y - \frac{3}{2}y^2 + \frac{10}{3}y^3 + \mathcal{O}(y^4), \quad y \rightarrow 0^+.$$

The inverse of $e^x + e^{-x}$ at $+\infty$ begins $\log y - y^{-2}$; selecting the source endpoint $-\infty$ gives its negative. The chart preserves this distinction rather than choosing a numerical root without an endpoint condition.

Original observables and term counts. At a finite source endpoint, the power option represents $(x - x_0)^r$ when $r \neq 1$, and the inverse itself when $r = 1$. At infinity it represents x^r . A negative source side therefore requires integer r for these real-valued observable conventions. Exponential reconstruction multiplies h by $\pm r$ before exponentiating; a finite source offset is added separately when $r = 1$.

For a non-polynomial power of a logarithmic reconstruction, the finite approximation can be retained inside a specified observable carrier. If $X \sim c|\log w|$ and $x - X = \mathcal{O}(w^P M^D)$ with $P > 0$, the mean value theorem gives

$$\frac{x^r - X^r}{X^r} = \mathcal{O}\left(\frac{x - X}{X}\right) = \mathcal{O}(w^P M^D).$$

This deliberately conservative relative bound makes the retained carrier X^r usable without claiming a finite polynomial-log expansion of its fractional power. An explicit cutoff is required for that representation; a unit-block term count would not count the corrections retained inside the carrier. In the ordinary reconstructed unit, the term-goal option counts the actual nonzero reconstructed blocks.

Residuals and numerical comparisons. The object records both directions of the chart and an expression that recovers the chart variable from the requested observable. Formal residual checking applies that expression to the *returned finite observable*. It does not reuse the unmodified inner inverse blocks after an exponential or logarithmic reconstruction has introduced a new truncation. The checked residual is

$$\frac{\Phi(h_{\text{returned}}) - y}{y - y_0},$$

with denominator y at an infinite target. If reconstruction requires a coefficient algebra outside the supported jet scale, the exact residual expression is still returned, accompanied by an explicit indication that no jet-order assertion was made.

Numerical comparison solves $\Phi(h) = y$ at high precision, reconstructs the source, checks its original side, and compares the requested observable with the returned approximation. This avoids unnecessarily evaluating a very large or very small original exponential during root finding. It is numerical evidence, not an interval certificate. Refinement replays the original inverse problem and reconstruction, preserving inverse provenance. An externally supplied forward remainder must first be transported through the source chart; such a declaration is rejected until that transport has been specified.

21 Finite logarithmic hierarchies and reciprocal-logarithmic order

A power gap and a logarithmic gap describe different orders of smallness. For example, the correction $u/\log u$ has the same source power as u , whereas $u^2\sqrt{-\log u}$ has a positive source-power gap and a nonpolynomial logarithmic coefficient. The operation `AsymptoticLogarithmicInverse` supplies separate contracts for these two situations and for a leading product of logarithmic factors. Exact source-coordinate reductions, including repeated logarithmic charts, are handled by the coordinate machinery of Section 15.

21.1 The finite coefficient hierarchy

In the positive source coordinate $u \rightarrow 0^+$, put

$$L_1(u) = -\log u, \quad L_{j+1}(u) = \log L_j(u), \quad 1 \leq j < d. \quad (21.1)$$

The depth d is finite. All retained levels are positive and tend to infinity sufficiently near the endpoint. A generalized logarithmic monomial is

$$c \prod_{j=1}^d L_j^{\alpha_j}, \quad c \in \mathbb{R}, \quad \alpha_j \in \mathbb{R}. \quad (21.2)$$

Finite sums of these monomials form an algebra under addition and multiplication. Powers in (21.2) have positive bases, including when an exponent is noninteger or negative. An arbitrary real power of a sum is not silently included in this finite algebra.

Proposition 21.1 (Closure under the source Euler derivative). *For a differentiable coefficient $C(L_1, \dots, L_d)$,*

$$u \frac{d}{du} C(L_1(u), \dots, L_d(u)) = \mathcal{E} C, \quad \mathcal{E} = - \sum_{j=1}^d \frac{1}{L_1 \cdots L_{j-1}} \frac{\partial}{\partial L_j}, \quad (21.3)$$

where the empty denominator is 1. In particular, the finite generalized logarithmic algebra is closed under every derivative used in the inverse coefficient formula.

Proof. The chain rule gives $uL_1'(u) = -1$ and $uL_j'(u) = -(L_1 \cdots L_{j-1})^{-1}$. Each summand in (21.3) decreases finitely many monomial exponents and creates no new logarithmic level. $\square$

This is the needed closure for coefficients such as $L_1^{1/2}$: its derivatives contain $L_1^{-1/2}, L_1^{-3/2}, \dots$. Restricting the coefficient reader to positive logarithmic powers would lose closure even in this simplest example. An expression containing $\sqrt{\log u}$ instead has an eventually negative base on $u \rightarrow 0^+$; the real hierarchy does not admit it.

21.2 Positive source-power gaps with logarithmic coefficients

Consider

$$F(u) = y_0 + au^p \left(1 + \sum_{i=1}^m u^{\delta_i} C_i(L_1(u), \dots, L_d(u)) \right), \quad a \neq 0, \quad p \neq 0, \quad \delta_i > 0, \quad (21.4)$$

where every C_i belongs to the finite algebra above. As before, all exponents are fixed exact real constants in the implementation, and assumptions establish the real coefficients and the sign of a . Set $z = ((y - y_0)/a)^{1/p} > 0$, so $z \rightarrow 0^+$. For the observable u^r , the multi-index coefficient of weight $w = \mathbf{k} \cdot \boldsymbol{\delta}$ is

$$\frac{(-1)^{|\mathbf{k}|_r}}{p^{|\mathbf{k}|} \prod_i k_i!} \prod_{j=1}^{|\mathbf{k}|-1} (\mathcal{E} + r + w + pj) \prod_i C_i^{k_i}, \quad (21.5)$$

evaluated at the logarithmic levels of z and multiplied by z^{r+w} . For $\mathbf{k} = 0$, the coefficient is 1. Formula (21.5) is the same Lagrange identity as in the polynomial case; Proposition 21.1 supplies its enlarged coefficient algebra.

Theorem 21.2 (Convergence and a conservative complete-tail bound). *The inverse of (21.4) on its selected eventual real branch has the convergent multi-index expansion (21.5). Choose integers $D_i \geq 0$ such that every monomial in C_i satisfies*

$$\prod_j L_j(u)^{\alpha_j} = \mathcal{O}(M(u)^{D_i}).$$

For example, the ceiling of the largest sum of the positive monomial exponents is sufficient. Let $\mathcal{B}$ be the finite minimal excluded multi-index boundary of a positive source-weight cutoff, and put

$$\beta = \min_{\mathbf{k} \in \mathcal{B}} \mathbf{k} \cdot \boldsymbol{\delta}, \quad D = \max_{\mathbf{k} \in \mathcal{B}} \sum_i k_i D_i. \quad (21.6)$$

Then the complete omitted expansion of u^r is $\mathcal{O}(z^{r+\beta} M(z)^D)$. At a finite source endpoint the original inverse has $r = 1$; at an infinite source endpoint it has $r = -1$, with the selected source sign applied separately.

Proof. Write $u = zW$ near $W = 1$. For each level, the quotient $L_j(zW)/L_j(z)$ is analytic in W and the finite reciprocal logarithmic parameters near their limiting values. For instance, the first quotient is $1 - \log W/L_1(z)$; subsequent quotients are obtained by taking the logarithm of the preceding quotient and dividing by the next level. Each factor raised to a fixed real power is holomorphic near 1. Introduce a separate small parameter for every finite coefficient monomial $z^{\delta_i} \prod_j L_j(z)^{\alpha_j}$. All tend to zero. The normalized implicit equation has derivative $p \neq 0$ in W at its parameter origin, so the analytic implicit function theorem proves convergence in a fixed polydisc of these parameters.

Cauchy estimates, with the finite coefficient sums grouped by i , give a geometric multi-index majorant in $K_i z^{\delta_i} M(z)^{D_i}$ for suitable constants K_i . The omitted upward-closed index set is covered by the translates of its finite minimal boundary. Sum the majorant over each translate. Its denominators remain bounded for sufficiently small z , and the boundary monomial is bounded by $z^\beta M(z)^D$ using (21.6). Multiplication by z^r completes the estimate. The estimate does not assume that the first boundary coefficient is nonzero. $\square$

The reported logarithmic degree can be larger than the optimal one. The finite boundary supplies a valid envelope for the entire omitted sum, while the retained coefficients remain exact generalized expressions. The cutoff uses the ordinary positive target coordinate: a source weight w contributes target power $(r + w)/|p|$. This convention preserves the nonunit leading-power and infinite-endpoint scaling of the original engine.

Example 21.3 (The two nonpolynomial coefficient acceptance cases). For $T = -\log y$, inversion at 0^+ gives

$$(u + u^2 \sqrt{-\log u})^{-1}(y) = y - y^2 \sqrt{T} + y^3(2T - \tfrac{1}{2}) + \mathcal{O}(y^4 M(y)^3), \quad (21.7)$$

$$(u + u^2 \log(-\log u))^{-1}(y) = y - y^2 \log T + y^3 \left(2 \log^2 T - \frac{\log T}{T} \right) + \mathcal{O}(y^4 M(y)^3). \quad (21.8)$$

Here the superscript -1 denotes function inversion. The displayed remainders are the implementation's conservative complete-tail bounds. The derivative term in each cubic coefficient is essential and follows directly from (21.3).

21.3 Reciprocal-logarithmic units

Now consider a core of the form

$$F_0(u) = y_0 + au^p H(-1/\log u), \quad p \neq 0, \quad H(0) = 1, \quad (21.9)$$

where H is a rational function analytic at 0 with real coefficients. Let $z = ((y - y_0)/a)^{1/p} > 0$, $t = -1/\log z > 0$, and write $u = zW$. The normalized inverse equation is exactly

$$W^p H\left(\frac{t}{1 - t \log W}\right) = 1. \quad (21.10)$$

Its derivative in W is p at $(t, W) = (0, 1)$, so its real solution $W(t)$ is analytic at 0. Iterating

$$W \mapsto H\left(\frac{t}{1 - t \log W}\right)^{-1/p} \quad (21.11)$$

from $W = 1$ determines successively correct Taylor coefficients: dependence on the previous W carries a positive power of t . Exact finite Taylor arithmetic supplies the requested bracket and its first omitted orders. An exclusive cutoff h retains exactly the integer degrees $n < h$. Its remainder has the form $\mathcal{O}(t^\beta)$, where β is the recorded first unretained degree or a conservative bound when cancellation prevents its identification in the bounded lookahead.

For the original source inverse, the bracket is multiplied by the exact prefactor σz at a finite endpoint or σ/z at infinity. An observable uses the corresponding power of W and transports the prefactor and remainder together. This is a logarithmic cutoff inside that prefactor, not a cutoff in y or $1/y$.

Example 21.4 (A zero source-power gap). At $u \rightarrow 0^+$, the inverse of $u + u/\log u$ begins

$$y \left(1 + t + t^2 + 2t^3 + \frac{7}{2}t^4 + \mathcal{O}(t^5) \right), \quad t = -\frac{1}{\log y}. \quad (21.12)$$

Equivalently, its bracket is $1 - 1/\log y + 1/\log^2 y - 2/\log^3 y + 7/(2\log^4 y)$. For the same forward expression at $+\infty$, use $t = 1/\log y$ and the bracket $1 - t + t^2 - 2t^3 + 7t^4/2$.

The expression $u + u/\log u - u/(1 + \log u)$ has a cancelled first reciprocal-logarithmic correction. Its forward unit is $1 + t^2/(1 - t)$ in the source logarithmic coordinate, and its inverse bracket begins $1 - t^2 + \dots$. Coefficients are merged before detecting the leading unit, so a cancelled coefficient does not create a fictitious retained term.

Higher source-power terms can accompany (21.9). If each has a strictly positive power gap and an admitted logarithmic coefficient envelope, their induced inverse correction is smaller than the prefactor times every fixed power of t . For $u(1 + 1/\log u) + u^2$, the logarithmic bracket therefore agrees with (21.12) at every finite logarithmic order. The package separately records the omitted source-power expression and an asymptotic bound for that smaller sector. Its logarithmic residual refers to the retained leading core. Equality of the logarithmic brackets does not assert that the full inverses are equal.

21.4 A leading logarithmic monomial

The same normalization extends to

$$F_0(u) = y_0 + au^p \prod_{j=1}^d L_j(u)^{q_j}, \quad p \neq 0, \quad q_j \in \mathbb{R}. \quad (21.13)$$

Let $Y = (y - y_0)/a > 0$ and define

$$\begin{aligned} T &= -\frac{\log Y}{p}, \quad t = \frac{1}{T}, \quad B_1 = T, \quad B_{j+1} = \log B_j, \\ d_0 &= \frac{1}{p} \sum_j q_j \log B_j, \quad z = Y^{1/p} \prod_j B_j^{-q_j/p}. \end{aligned} \quad (21.14)$$

All B_j are positive sufficiently near the selected endpoint. Put

$$R_1(W) = 1 + t(d_0 - \log W), \quad (21.15)$$

$$R_{j+1}(W) = 1 + B_{j+1}^{-1} \log R_j(W). \quad (21.16)$$

Then $L_j(zW)/B_j = R_j(W)$ exactly, and the normalized inverse equation is

$$W^p \prod_j R_j(W)^{q_j} = 1. \quad (21.17)$$

Theorem 21.5 (A convergent leading-hierarchy bracket). *Equation (21.17) has a unique real solution near $W = 1$ for the selected asymptotic branch. Its expansion in powers of t , retaining the lower-level coefficients exactly, converges sufficiently near the endpoint. A truncation beginning its omitted degrees at β has bracket remainder*

$$\mathcal{O}\left(t^\beta(1 + |\log t|)^\beta\right). \quad (21.18)$$

Proof. Treat t , td_0 , and $B_2^{-1}, \dots, B_d^{-1}$ as independent small parameters. The equation is analytic in these parameters and W near their origin and $W = 1$, with derivative $p \neq 0$ in W . The analytic implicit function theorem gives a common polydisc and a convergent series there. Since $d_0 = \mathcal{O}(\log T)$, all the parameters indeed tend to zero on the actual asymptotic curve.

The total degree in the first two parameters equals the formal degree in t after substituting td_0 . Cauchy estimates, uniform in the other parameters on a smaller polydisc, bound the tail of those degrees by a constant times $[t(1 + |d_0|)]^\beta$ for sufficiently small t . This proves (21.18). The fixed-point expression $W = \prod_j R_j(W)^{-q_j/p}$ has positive formal valuation in its dependence on the previous W , so finite exact iteration computes the same coefficients. Real powers use the positive local factors near 1. $\square$

For $F(x) = x \log \log x$ at $+\infty$, let $T = \log y$, $B = \log T$, and $t = 1/T$. The reconstructed inverse starts

$$F^{-1}(y) = \frac{y}{B} \left(1 + \frac{\log B}{B} t + \mathcal{O}\left(t^2(1 + |\log t|)^2\right) \right). \quad (21.19)$$

Thus an iterated logarithm in the leading coefficient need not be replaced by a constant or rejected as a zero source-power gap. The implementation also admits products of levels and real powers of those levels, with positivity recorded explicitly. Positive source-power perturbations can again be retained as separately recorded smaller sectors.

21.5 Implementation contracts and checks

```
AsymptoticLogarithmicInverse[
  x + x/Log[x] + x^2, {x, 0}, {y, 5}]

AsymptoticLogarithmicInverse[
  x + x^2 Log[-Log[x]], {x, 0}, {y, 4}]

AsymptoticLogarithmicInverse[
  x Log[Log[x]], {x, Infinity}, {y, 4}]
```

The scale names identify three distinct contracts:

- "ReciprocalLogUnit": a relative logarithmic cutoff;
- "GeneralizedLogarithmicCoefficients": a target-power cutoff;
- "LeadingLogMonomial": a relative logarithmic cutoff.

All three cutoffs are exclusive. The positive uniformizer, exact prefactor, coefficient levels, branch conditions, and cutoff meaning are all inspectable. The option "LogarithmicLevels" bounds the admitted finite hierarchy; unused levels are removed from the result's actual domain requirements.

When `SeriesTermGoal -> n` is used instead of a cutoff, the constructor searches for n complete nonzero blocks. Cancelled logarithmic coefficients do not count as terms. Generalized coefficient searches advance across complete source-weight boundaries and, if necessary, rebuild at the next excluded target weight to keep exactly the requested blocks. Their target-power cutoff may be negative, but must exceed the leading power of the requested observable. A unit's reciprocal-logarithmic cutoff remains positive. The search has a finite construction budget and preserves its best expansion in a resource failure.

A finite run of zero coefficients does not establish termination. A sufficient exactness certificate reconstructs the selected positive local source from the returned observable and verifies the original forward equation identically. The model must be exact, with no omitted source-power sectors. The positive branch and local uniqueness then identify this candidate with the inverse germ. In particular, an exact source-unit polynomial does not by itself certify a truncated noninteger power of that unit. The tests include cancellation, negative target-power cutoffs, refinement and bounded failures for term-goal requests.

`LogarithmicInverseResidual` composes the retained normalized unit with the exact equation (21.10) or (21.17). Its scope explicitly identifies a leading-core residual when smaller source-power sectors were omitted. Independent tests compare the generalized coefficients with direct derivatives and hand-derived inverse coefficients. Other tests substitute targets generated by the original forward functions and compare against the known source point at high precision. The regression family covers two logarithmic cutoffs, cancellation, irrational source gaps, both source orientations, finite translations, infinite endpoints and invalid real logarithmic domains.

The reciprocal-logarithmic unit has a single ordinary Taylor bracket and therefore exposes a representation usable by the existing series calculus. The other two scales retain their enlarged coefficient algebra explicitly; operations requiring a polynomial-log coefficient representation must reject them until a compatible conversion or generalized operation is available. The convergence contracts are mathematical existence bounds, with "NumericCertificate" set to `False`. An external unknown input remainder is rejected here until its transport through the chosen logarithmic hierarchy is supplied.

This finite implementation does not claim a general transseries field. Arbitrary real powers of coefficient sums, sums of several incomparable leading logarithmic monomials, unrestricted analytic composition of the enlarged coefficient objects, and independent simultaneous cutoffs at every logarithmic level remain additional contracts. The formulas above give a closed finite coefficient algebra for the derivatives actually computed, convergent inverse expansions for the admitted families, and explicitly different precision meanings for the supported scales.

22 Finite commensurate flat sectors

An exponential that is smaller than every power of the source coordinate needs its own truncation index. The operation `AsymptoticFlatInverse` implements a finite part of this larger scale. In the positive source coordinate $u \rightarrow 0^+$, its admitted exact forward expression is

$$F(u) = y_0 + au^q + \sum_{j=1}^J A_j(u) e^{-k_j c/u^p}, \quad p, c > 0, \quad q \neq 0, \quad a \neq 0, \quad k_j \in \mathbb{N}_{>0}, \quad (22.1)$$

where each A_j is a finite real power–log expression. The phase power and rates are fixed exact real numbers; the rates are integer multiples of the smallest input rate. Coefficients may contain symbolic real parameters when the required signs are proved. Different phase powers and incommensurate rates need more than one exponential coordinate and are rejected here.

The zero sector must be exact. An ordinary finite asymptotic approximation to a nonmonomial core leaves a power-sized error, which dominates every sector in (22.1). The implementation therefore admits the exact shifted monomial core $y_0 + au^q$; it does not use a truncated inverse of a general power–log core as the zero sector. Put

$$u_0 = \left(\frac{y - y_0}{a} \right)^{1/q} > 0, \quad E_0 = e^{-c/u_0^p}.$$

At a finite source endpoint the requested observable is $H(u) = x_0 + \sigma u$ for power one and $H(u) = \sigma^r u^r$ otherwise. At an infinite endpoint it is $H(u) = \sigma^r u^{-r}$. Negative source branches require an integer requested power.

22.1 Computing complete sectors

Introduce a formal symbol Z and the polynomial $R(u, Z) = \sum_j A_j(u) Z^{k_j}$. Differentiation of the actual exponential coordinate is represented by the exact operator

$$\mathcal{D} = \partial_u + cp u^{-p-1} Z \partial_Z, \quad \frac{d}{du} e^{-c/u^p} = cp u^{-p-1} e^{-c/u^p}.$$

The marker identity of Section 18 now gives

$$\frac{(-1)^n}{n!} \left(\frac{\mathcal{D}}{aqu^{q-1}} \right)^{n-1} \left(\frac{H'(u)R(u, Z)^n}{aqu^{q-1}} \right). \quad (22.2)$$

Sum this expression over $1 \leq n \leq N$, group equal powers of Z , and evaluate $u = u_0$, $Z = E_0$. Since every $k_j \geq 1$, no marker degree larger than N contributes to an exponential degree at most N . Also $\mathcal{D}$ preserves exponential degree, so intermediate polynomial truncation discards no contribution to a retained sector. Each returned coefficient is a complete finite power–log function of u_0 .

For example,

```
AsymptoticFlatInverse[x + Exp[-1/x], {x, 0}, {y, 3}]
```

returns

$$y - E + y^{-2}E^2 + \left(y^{-3} - \frac{3}{2}y^{-4}\right)E^3, \quad E = e^{-1/y}. \quad (22.3)$$

For $F(x) = x + x^2e^{-1/x}$, the first two sectors are $y - y^2E + (y^2 + 2y^3)E^2$. Adding $2e^{-2/x}$ to the first example changes the second coefficient to $y^{-2} - 2$. Thus the sector index distinguishes input sectors from powers of the perturbation marker.

22.2 A bound for the complete omitted tail

Write every coefficient as a finite sum $u^\alpha P(\log u)$, and choose

$$b = \max(0, \max_\alpha(q + p - \alpha)), \quad d = \max(0, \max_P \deg P),$$

$$B(u) = u^{-b}M(u)^d, \quad \chi(u) = B(u)e^{-c/u^p}.$$

Let $r_* = r$ at a finite source endpoint and $r_* = -r$ at infinity. For fixed admitted data, there are $C, K > 0$ and a source threshold such that the complete omitted sector tail after degree N satisfies

$$|\text{Tail}_N| \leq Cu_0^{r_*+p} \frac{(K\chi(u_0))^{N+1}}{1 - K\chi(u_0)} \quad \text{when } K\chi(u_0) < 1. \quad (22.4)$$

Consequently its recorded asymptotic class is

$$\mathcal{O}\left(E_0^{N+1}u_0^{r_*+p-(N+1)b}M(u_0)^{(N+1)d}\right).$$

This bound is allowed to be conservative; it is independent of whether the first omitted coefficient cancels.

To justify the estimate, set $u = u_0(1 + u_0^p V)$ and divide the forward equation by u_0^{q+p} . The normalized monomial difference has the analytic extension

$$a \frac{(1 + \epsilon V)^q - 1}{\epsilon}, \quad \epsilon = u_0^p,$$

whose derivative at $V = 0$, $\epsilon = 0$ is $aq \neq 0$. Each exponential factor, after extracting $E_0^{k_j}$, becomes

$$\exp\left(k_j c \frac{1 - (1 + \epsilon V)^{-p}}{\epsilon}\right),$$

which is analytic and uniformly bounded for small ϵ, V . The power-log factors have the same local analytic property after their powers of u_0 and $\log u_0$ are extracted. Each normalized coefficient is bounded by a constant times $B(u_0)$. Since $B \geq 1$ for small u_0 , replacing E_0 by a complex variable Z gives a uniform implicit solution on $|Z| < 1/(KB(u_0))$. One may regard the bounded normalized coefficients as additional parameters in a compact polydisc; the derivative condition at $Z = 0$ is uniform. Cauchy estimates therefore bound the coefficient of Z^k by $C(KB)^k$. The nonconstant observable contributes the factor $u_0^{r_*+p}$. Summing the geometric tail proves (22.4), and $\chi(u_0) \rightarrow 0$ establishes applicability at $Z = E_0$.

The constants and threshold are existence bounds, not computed numerical certificates. The result records this distinction in `MajorantContract`. `SectorDepth` is inclusive: it retains $k \leq N$, whereas an ordinary power cutoff is exclusive. `FirstOmittedSector` records degree $N + 1$ separately from the complete-tail remainder. The resource budget bounds both the requested depth and the intermediate exact expressions.

Independent verification substitutes $g = y + \sum_{k=1}^N C_k(y)Z^k$ into $g + Z \exp(1/y - 1/g) - y$ and uses ordinary Taylor expansion in Z . This checks (22.3) without reusing the coefficient operator.

A high-precision solve of the original equation resolves errors on the fourth exponential scale; machine arithmetic cannot supply such evidence. Further tests exercise polynomial factors, two commensurate input sectors, nonunit monomial cores, shifted branches, infinite source endpoints, observable powers, and the rejected scale and resource cases.

23 Finite Fourier coefficients in a logarithmic variable

The polynomial coefficient algebra extends naturally to finite sums

$$B(L) = \sum_{\omega \in \Omega} P_{\omega}(L) e^{i\omega L}, \quad L = \log u, \quad \Omega \subset \mathbb{R} \text{ finite},$$

where every P_{ω} is a polynomial and the frequencies are exact real numbers. This includes $\sin(\log u)$, several incommensurable frequencies, and polynomial amplitudes multiplying those oscillations. The source powers and the Fourier frequencies have different roles: source weights determine asymptotic truncation, while frequencies label coefficient modes. No Fourier mode is discarded merely because its frequency is large.

23.1 Coefficient arithmetic and reality

Multiplication convolves frequencies, and the Euler derivative is

$$\frac{d}{dL} (P_{\omega}(L) e^{i\omega L}) = (P'_{\omega}(L) + i\omega P_{\omega}(L)) e^{i\omega L}.$$

Thus this algebra is closed under precisely the operations used by Theorem 6.1. Distinct syntactic frequencies are combined when exact comparison proves them equal. Algebraic frequencies are reduced before comparison. No floating-point tolerance is used to identify modes; an undecidable comparison produces a failure.

The real-valued condition is the conjugacy relation $P_{-\omega}(L) = \overline{P_{\omega}(L)}$ on real L . The parser checks it under the supplied parameter assumptions. Intermediate complex coefficients are permitted, and real trigonometric expressions are reconstructed for display. Affine trigonometric phases, such as $\sin(2L + 1)$, contribute their constant phase factor to the polynomial amplitude.

23.2 Inverse coefficients and independent composition

The admitted forward family is

$$f(u) = y_0 + au^p \left(1 + \sum_j u^{\delta_j} B_j(\log u) \right), \quad ap \neq 0, \quad \delta_j > 0,$$

with finitely many real Fourier-polynomial coefficients. The leading core is a monomial with a proved nonzero real coefficient. An oscillatory leading block, even one that happens to remain positive, requires a different inverse normalization and is rejected explicitly by this engine.

For a multi-index $\mathbf{k}$ of depth n and weight $\sigma = \mathbf{k} \cdot \boldsymbol{\delta}$, the unchanged coefficient formula is

$$C_{\mathbf{k}}^{(r)}(L) = \frac{(-1)^n r^{n-1}}{p^n \mathbf{k}!} \prod_{j=1}^{n-1} (\mathcal{D} + r + \sigma + pj) \prod_i B_i(L)^{k_i}.$$

Convolution forms the product, and the Fourier Euler operation applies each factor. The implementation preserves the individual coefficients as well as the complete sum at a source weight.

Residual checking uses a separate composition recurrence. If $u = z(1 + U)$, then

$$(1 + U)^\alpha B(L + \log(1 + U)) = \sum_{k \geq 0} Q_k(L) U^k, \quad Q_0 = B, \quad Q_{k+1} = \frac{(\alpha - k)Q_k + \mathcal{D}Q_k}{k + 1}.$$

This follows by Taylor expansion in the multiplicative increment; it holds for each exponential-polynomial mode as well as for their sum. Its use in the normalized forward equation independently checks the inverse coefficients. For a nonunit observable power, the residual first reconstructs the source unit by raising the observable unit to its reciprocal power.

23.3 Envelopes, convergence, and omitted terms

For real L , Fourier modes have modulus one. If $P_\omega(L) = \sum_j c_{\omega,j} L^j$, then

$$|B(L)| \leq \sum_{\omega,j} |c_{\omega,j}| (1 + |L|)^j.$$

This nonoscillatory envelope is stored with each retained block. Euler differentiation preserves the polynomial degree; frequency factors affect coefficient constants, rather than the power of the logarithmic envelope. Consequently a multi-index coefficient has polynomial degree at most $\sum_i k_i \deg B_i$, just as in the polynomial case.

The convergence argument also extends. After replacing each finite source mode by an independent marker, logarithmic composition uses

$$e^{i\omega(L + \log W)} = e^{i\omega L} W^{i\omega},$$

which is analytic for W near 1 with its local logarithm. The implicit derivative at the origin is still $p \neq 0$. The finite-dimensional analytic implicit function theorem gives an absolutely convergent marker expansion for sufficiently small markers. Along the real source slice these markers are bounded by constants times $z^{\delta_i} (1 + |\log z|)^{d_i}$ and tend to zero. The positive source gaps also make the forward derivative asymptotic to apu^{p-1} , so the selected real inverse branch exists near the endpoint.

The complete one-step outer boundary of the retained multi-index region therefore bounds the omitted tail. Its least weight supplies the remainder power, and the maximum summed polynomial degree over that finite boundary supplies a conservative logarithmic envelope. This bound need not be sharp after cancellation, but remains valid even at zeros of every displayed oscillatory coefficient. A declared forward remainder is transported with the same observable-dependent exponent as Theorem 9.1; its matching derivative hypothesis remains required. Constants and source thresholds are not computed numerical certificates.

23.4 Interface and acceptance examples

```
s = AsymptoticFourierInverse[x + x^2 Sin[Log[x]],
  {x, 0}, {y, 4}];
FourierInverseResidual[s]
FourierInverseCoefficient[s, {2}]
```

The result is

$$y - y^2 \sin L + y^3 \sin L (2 \sin L + \cos L) + \mathcal{O}(y^4), \quad L = \log y.$$

With an additional $x^3 \cos(2 \log x)$, the coefficient of y^3 acquires $-\cos(2L)$, combining a directly supplied frequency with frequencies generated by convolution. The identity

$$\sin(\sqrt{2}L) \sin(\sqrt{3}L) + \frac{1}{2} \cos(\sqrt{5 + 2\sqrt{6}}L) = \frac{1}{2} \cos((\sqrt{2} - \sqrt{3})L)$$

tests exact algebraic frequency equality and cancellation.

The cutoff is exclusive in the positive target power coordinate, with the same finite-endpoint, infinity, and observable conventions as the ordinary engine. This first Fourier engine requires an explicit cutoff and uses the individual Lagrange formula. Frequencies are exact numeric real constants; arbitrary symbolic frequency relations and nonlinear logarithmic phases are outside this interface. "MaxFrequencies" bounds distinct modes across the retained jet and intermediate coefficient operations, and "MaxTerms" bounds source indices and convolution work. Exhaustion returns an explicit failure; there is no undocumented frequency truncation.

Regression cases compare the displayed formulas with direct trigonometric differentiation, independently compose normalized residuals, check source and target translations and an infinite source endpoint, and exercise observable powers, algebraic resonance, polynomial envelopes, remainder transport, and frequency budgets. The Fourier envelope test deliberately evaluates at a zero of the oscillatory factor, where a sign-based or oscillatory error scale would give a misleading result.

24 Special functions, asymptotic forward models, and real thresholds

The operation `AsymptoticSpecialInverse` exposes explicit adapters for "Erfc", "LogGamma", "Gamma", "LambertThreshold", and "QuadraticThreshold". These examples need two different contracts. The tail adapters invert a finite Poincaré forward model and transport its omitted error. The threshold adapters use exact local equations in a ramified coordinate. A convergent inverse expansion for a finite model does not make the original Poincaré forward series convergent.

Every adapter permits a real affine change of target, $y = b + aF(x)$, with a provably real offset and a provably nonzero real scale. It stores the original function, the forward model or exact local equation, the transformed target, the selected real branch, and the meaning of the requested cutoff. The following formulas use $a = 1, b = 0$ unless stated otherwise.

24.1 The complementary-error-function tail

For $x > 0$, define

$$S(x) = \sqrt{\pi} x e^{x^2} \operatorname{erfc} x, \quad S_m(x) = \sum_{k=0}^{m-1} \frac{(-1)^k (1/2)_k}{x^{2k}}, \quad A_m = (1/2)_m. \quad (24.1)$$

The real positive-axis expansion has a remainder bounded by its first neglected term; see [DLMF §7.12\(i\)](#). For inverse-error transport we also require a derivative bound.

Lemma 24.1 (Value and derivative bounds for the normalized tail). *For every integer $m \geq 1$ and every $x > 0$,*

$$|S(x) - S_m(x)| \leq A_m x^{-2m}, \quad |S'(x) - S'_m(x)| \leq 2m A_m x^{-2m-1}. \quad (24.2)$$

Proof. Starting with the integral definition of erfc in [DLMF \(7.2.2\)](#), substitute $s^2 = x^2 + t$ to obtain

$$S(x) = \int_0^\infty e^{-t} (1 + t/x^2)^{-1/2} dt. \quad (24.3)$$

Put $w = x^{-2}$. The derivatives of $(1 + tw)^{-1/2}$ alternate in sign, and their magnitudes for $w \geq 0$ are at most their values at $w = 0$. Taylor's integral remainder, followed by integration in t ,

bounds the remainder of the first m terms by $(1/2)_m w^m$. Applying the same argument to the derivative with respect to w bounds its remainder by $m(1/2)_m w^{m-1}$. For $m = 1$, the latter statement is the direct bound on the first derivative. Finally $|dw/dx| = 2x^{-3}$ gives the second inequality in (24.2). The exponential factor in the integral justifies these finite differentiations and integrations. $\square$

Let $Q_m(w)$ be the Taylor polynomial of degree $m - 1$ of $-\log S_m(w^{-1/2})$ at $w = 0$. The adapter constructs the finite phase

$$\Phi_m(x) = x^2 + \log x + Q_m(x^{-2}), \quad v = -\log(\sqrt{\pi} y). \quad (24.4)$$

The exact phase is $\Phi(x) = -\log(\sqrt{\pi} \operatorname{erfc} x)$. For fixed m , Lemma 24.1 implies

$$\Phi(x) - \Phi_m(x) = \mathcal{O}(x^{-2m}), \quad \Phi'(x) - \Phi'_m(x) = \mathcal{O}(x^{-2m-1}). \quad (24.5)$$

These are value and derivative assertions; the second is not inferred merely by differentiating an unspecified big- $\mathcal{O}$ term.

The implementation records explicit conditional bounds as well. Write $\delta = A_m x^{-2m}$, $\delta_1 = 2mA_m x^{-2m-1}$, and $\tilde{\Phi}_m = x^2 + \log x - \log S_m$. Wherever $S_m > \delta$,

$$|\Phi - \Phi_m| \leq \frac{\delta}{S_m - \delta} + |\tilde{\Phi}_m - \Phi_m|, \quad (24.6)$$

$$|\Phi' - \Phi'_m| \leq \frac{\delta_1}{S_m - \delta} + \frac{|S'_m|\delta}{S_m(S_m - \delta)} + |\tilde{\Phi}'_m - \Phi'_m|. \quad (24.7)$$

The terms involving $\tilde{\Phi}_m$ are exact elementary expressions. The positivity condition is retained in the result; it is satisfied eventually because $S_m \rightarrow 1$ and $\delta \rightarrow 0$.

Now invert $\Phi_m(x) = v$ on $x \rightarrow +\infty$. Its leading source scale is x^2 , so the inverse error transported from (24.5) is $\mathcal{O}(x^{-2m-1}) = \mathcal{O}(v^{-m-1/2})$. The ordinary inverse engine receives the explicit local input remainder $\{2m, 0\}$ in $u = 1/x$. A requested cutoff larger than $m + 1/2$ is rejected unless the forward model is refined. This bound is combined with the truncation error of the model inverse.

For example, with $v = -\log(\sqrt{\pi} y)$,

$$\operatorname{erfc}^{-1} y = \sqrt{v} - \frac{\log v}{4\sqrt{v}} + \frac{-\log^2 v + 4\log v - 8}{32v^{3/2}} + \mathcal{O}\left(v^{-5/2}(1 + |\log v|)^3\right). \quad (24.8)$$

The cubic numerator follows by inserting $x = \sqrt{v} + A/\sqrt{v} + B/v^{3/2}$ into $x^2 + \log x + (2x^2)^{-1} + \mathcal{O}(x^{-4}) = v$. The expansion has a power cutoff in $1/v$, although $y \rightarrow 0^+$.

The negative-source tail uses the exact identity $\operatorname{erfc}(-x) = 2 - \operatorname{erfc} x$ and changes the reconstructed source sign. For $b + a \operatorname{erfc} x$ at $-\infty$, the positive small tail is $2 - (y - b)/a$ and the target endpoint is $b + 2a$. These identities handle reflected, shifted and sign-scaled tails without changing the forward-error proof.

24.2 A Stirling model with a retained Lambert inverse

For $x > 0$, let

$$G_m(x) = x(\log x - 1) - \frac{1}{2} \log x + \frac{1}{2} \log(2\pi) + \sum_{k=1}^{m-1} \frac{B_{2k}}{2k(2k-1)x^{2k-1}}. \quad (24.9)$$

The positive-axis Stirling and digamma error bounds in DLMF §5.11(ii) give

$$|\log \Gamma(x) - G_m(x)| \leq A_m x^{-\rho}, \quad (24.10)$$

$$|\psi(x) - G'_m(x)| \leq \rho A_m x^{-\rho-1}, \quad (24.11)$$

$$\rho = 2m - 1, \quad A_m = \frac{|B_{2m}|}{2m(2m-1)}.$$

The finite value and derivative bounds apply for all positive x . The adapter selects the increasing inverse branch $x > 2$; on this interval $\psi(x) \geq \log x - (2x)^{-1} - (12x^2)^{-1} > 0$. This branch choice excludes the separate behavior near the positive minimum of the gamma function.

Use the exact core

$$F_0(x) = x(\log x - 1), \quad \phi(v) = \frac{v}{W_0(v/e)}. \quad (24.12)$$

The difference $G_m - F_0$ is a finite higher-power perturbation in $u = 1/x$. The operation of Section 18 therefore retains ϕ exactly and computes its complete marker corrections. With $u_0 = 1/\phi(v)$, the omitted Stirling error transports to

$$\mathcal{O}(u_0^{2m-1}) \quad (24.13)$$

in the implementation's conservative polynomial-log bound. The actual derivative of the core provides an additional logarithmic denominator, but discarding that improvement preserves a valid bound in the stated remainder class. The final error combines (24.13) with the complete marker tail.

For $m = 1$, the first marker correction gives

$$(\log \Gamma)^{-1}(v) = \phi(v) + \frac{1}{2} - \frac{\log(2\pi)}{2 \log \phi(v)} + \mathcal{O}(\phi(v)^{-1}). \quad (24.14)$$

Increasing marker depth while holding m fixed cannot improve the inherited Stirling accuracy limit. Marker coefficients beyond that limit still describe the chosen finite model, with that interpretation explicit in the result. Neither the finite marker expansion nor its analytic complete-tail theorem establishes convergence of the infinite Stirling series.

For inversion of $\Gamma(x)$, use the exact target transformation $v = \log y$. The implementation solves and checks the equation $\log \Gamma(x) = \log y$; it can therefore accept a target such as e^{10000} without forming that enormous value in the forward residual. For a real affine target $b + a\Gamma(x)$, replace y by $(y - b)/a > 1$. The analogous log-gamma adapter uses $v = (y - b)/a > 0$ directly.

24.3 The two real Lambert branches at their common threshold

The two real inverses of $F(x) = xe^x$ meet at $x = -1$, $y = -1/e$. The branch description is given in [DLMF §4.13](#). Writing $x = -1 + h$ gives the exact local equation

$$eF(-1 + h) + 1 = \frac{h^2}{2} + \frac{h^3}{3} + \frac{h^4}{8} + \frac{h^5}{30} + \cdots. \quad (24.15)$$

The leading power is two. Ordinary analytic reversion in the positive ramified coordinate

$$t = \sqrt{2(1 + ey)} \quad (24.16)$$

therefore yields a convergent local expansion. Its first coefficients are

$$W_\nu(y) = -1 + \sigma t - \frac{t^2}{3} + \sigma \frac{11t^3}{72} - \frac{43t^4}{540} + \mathcal{O}(t^5), \quad (24.17)$$

where $(\nu, \sigma) = (0, 1)$ or $(-1, -1)$. The implementation computes these coefficients from the exact forward equation, not from a large-argument Lambert asymptotic formula.

The option "LambertBranch" selects 0 or -1 . Its source direction must agree: the principal branch approaches -1 from above and the lower branch from below. Contradictory choices are rejected. The recorded real target neighborhood is $-1/e < y < 0$, transformed by the requested affine target map. A target below the real threshold is rejected by the numerical checker. The cutoff is an exclusive power cutoff in the positive distance $y + 1/e$, so a cutoff $5/2$ retains the terms through t^4 .

24.4 An exact coalescing quadratic branch

For an exact quadratic model

$$F(x) = b + ac(x - x_0)^2, \quad ac \neq 0, \quad (24.18)$$

the two branches are exactly

$$x = x_0 + \sigma \sqrt{\frac{y - b}{ac}}, \quad \sigma \in \{1, -1\}, \quad \frac{y - b}{ac} > 0. \quad (24.19)$$

The source direction selects σ . The remainder is zero because this is an exact inverse identity. For $F(x) = x^2 - \mu$, the branches at target $y = 0$ are $\pm\sqrt{\mu}$ and coalesce as $\mu \rightarrow 0^+$. The same coordinate retains the selected sign throughout this limiting comparison. It does not choose a different branch because the two roots become close.

24.5 API, numerical evidence, and switching policy

```
AsymptoticSpecialInverse["Erfc", {x, Infinity}, {y, 3}]
AsymptoticSpecialInverse["LogGamma", {x, Infinity}, {y, 3}]
AsymptoticSpecialInverse["Gamma", {x, Infinity}, {y, 2}]
AsymptoticSpecialInverse[
  "LambertThreshold", {x, -1}, {y, 5/2},
  "LambertBranch" -> -1]
AsymptoticSpecialInverse[
  "QuadraticThreshold", {x, 0}, {y, 1},
  "TargetOffset" -> -mu, Assumptions -> mu > 0,
  Direction -> "FromBelow"]
```

The gamma calls use an integer marker depth. The error-function call uses a power cutoff in its logarithmic target coordinate. The threshold calls use local target-power cutoffs. Each result identifies this distinction. "ModelTerms" controls the number of retained forward asymptotic terms, with a default sufficient for the requested error-function cutoff or gamma marker depth. A model that cannot support an error-function cutoff is rejected; a deliberately coarse Stirling model retains its explicit accuracy floor.

SpecialInverseNumericalCheck uses the original special-function equation or the exact threshold inverse as its independent reference. Exact target substitutions precede numerical evaluation, preserving $\log(e^v) = v$ and avoiding unnecessary cancellation in reflected tiny error-function tails. If a numerical input has already lost the digits needed by its transformed target, the checker reports insufficient precision. Its output sets "Certified" to **False**: a high-precision solve is not an interval proof.

The regression families compare successive model orders at common target values, test the first-neglected value and derivative bounds against the original special functions, compare local Lambert expansions with exact **ProductLog** values on both real branches, and use the exact quadratic formula as an oracle. These are explicit overlap comparisons. They do not establish a

universal automatic crossover threshold between a local expansion, a tail expansion and another numerical representation.

The current adapters are selected explicitly. Their metadata states this switching policy. Complex branches, automatic certified crossover selection, the gamma minimum and its second positive inverse branch, arbitrary coalescing higher-order critical points, and additional special functions remain outside these admitted families. The four roadmap acceptance types are represented here with their own branch and error contracts; their existence does not make these remaining cases automatic.

25 Refinement by retaining complete coefficient states

The public operation `SeriesRefine[s,h]` reuses an ordinary inverse when its stored forward model remains sufficient at the requested cutoff. The implementation in `Kernel/RefinementState.wl` supports the individual Lagrange and Newton methods. Its state is stored in the returned object; there is no global cache, and refinement leaves the source object unchanged. Coordinate wrappers, grouped coefficient generation, and inputs whose forward expansion needs additional terms use the existing source replay path.

Complete Lagrange prefixes. Write the normalized inverse observable as

$$z^r \sum_{\omega < H} z^\omega Q_\omega(\log z), \quad \omega = \mathbf{k} \cdot \mathbf{d}, \quad \mathbf{k} \in \mathbb{N}^m, \quad d_j > 0.$$

The retained state contains a downward closed set of processed indices, its one-step outside boundary, and the complete coefficient polynomial at every processed weight. Zero polynomials are omitted from the displayed blocks but their indices remain processed. Advancing the state takes *all* boundary indices of the smallest weight simultaneously. Every predecessor of such an index has smaller weight, because the gaps are positive. Consequently, no as yet undiscovered index can have the same weight, and resonant contributions are combined before the new block is used or displayed.

An inverse produced with a term goal may already retain this state. For an ordinary inverse constructed at an explicit cutoff, the complete stored blocks instead seed the state. Enumerating $\{\mathbf{k} : \mathbf{k} \cdot \mathbf{d} < H\}$ and its boundary recovers the index bookkeeping without reevaluating the previously computed Euler coefficients. The aggregate blocks remain authoritative; the seed does not attempt to split a resonant coefficient polynomial into its individual contributions.

The first omitted block is also evaluated by complete weight layers. These frontier coefficients are retained for later refinement, including zero layers skipped after cancellation. Caching a frontier can require more indices than merely evaluating that frontier for a remainder. If the additional state would exceed "MaxTerms", the implementation performs the ordinary independent frontier calculation without caching it. Thus this optimization does not reject an otherwise admissible requested cutoff solely because it tried to retain an extra frontier. As in ordinary construction, the bounded search through cancelled frontier layers may return a conservative remainder at the first possible omitted weight.

A bounded cache of logarithmic polynomial powers. The Lagrange state also retains powers of each nonconstant logarithmic coefficient polynomial. Exponents 0 and 1 refer to 1 and the original model coefficient; larger cached powers form an initial consecutive segment. A new power is obtained by multiplying the last retained power by the original polynomial. Thus the product entering an Euler coefficient can reuse polynomial powers requested by earlier indices or

earlier refinement calls. The Euler differential operators and normalization remain unchanged, and the ordinary uncached coefficient routine provides an independent exact comparison.

Cached powers use only unused "MaxTerms" slots after the processed index region and its boundary. When the region grows, the largest retained powers are evicted until their entry count fits the remaining capacity. If no slot is available, the requested coefficient product is evaluated directly. Eviction never discards an inverse coefficient or changes the remainder calculation. In particular, optional caching cannot turn an otherwise admissible index region into a resource failure. This budget counts retained algebraic objects; it is not a byte bound on a polynomial whose individual coefficients may themselves be large exact expressions.

Recovering and extending a Newton state. Suppose the displayed normalized observable is $(1 + U)^r$ with $r \neq 0$. Its constant coefficient is 1, so the formal power with exponent $1/r$ uniquely recovers $1 + U$ in the filtered unit algebra. This recovers an initial Newton state from an existing observable, including inverse powers arising from a source endpoint at infinity. The recovered unit is checked by substituting it into the normalized implicit equation: the residual must vanish exactly below the claimed precision. Only then is the state admitted for further Newton steps.

If a Newton state is already present, refinement retains its unit jet, precision, and recorded step cutoffs. Each new step increases the precision to the smaller of twice the old precision and the requested precision. The existing exact residual check applies at every step. After a coarser display is requested, a stronger retained state remains available to subsequent calls. The formal branch is selected by constant term 1 throughout; no numerical root selection enters this recovery.

Model identity and input precision. Retained states are bound to the normalized model, observable power, assumptions, original function, and source endpoint and direction. A stored signature and direct comparisons detect incompatible states, and the displayed prefix must agree with the corresponding state prefix. This consistency check prevents accidental reuse after changing a model; it is not a certificate for arbitrarily edited internal associations.

A declared forward error remains a hard cap on usable inverse precision. For a leading exponent p and internal observable exponent r , a forward error $\mathcal{O}(u^\rho |\log u|^k)$ permits target cutoff at most

$$\frac{\rho - p + r}{|p|}.$$

The remainder's logarithmic degree is combined with that transported error at equal power. An error introduced by automatic expansion of an analytic source is distinguished from a declared error: if refinement requires more source terms, the constructor is replayed and the old model state is not reused. No cached inverse coefficients justify going beyond an explicit input-error declaration.

What the counters establish. Each refined object exposes "RefinementStatistics" and appends it to "RefinementHistory". The Lagrange counters have distinct meanings:

"ReusedCoefficientEvaluations"

Available individual index contributions. For a seed reconstructed from stored complete blocks this is the size of the old processed region, not work performed during seeding.

"NewCoefficientEvaluations"

Actual new Euler coefficient evaluations, including any uncached frontier work. Supplying a cached polynomial product still counts as one coefficient evaluation.

"ReusedBlocks"

Old nonzero blocks retained below the requested cutoff. The full old display has its own count, "AvailableSourceBlocks".

"ReusedPolynomialPowerRequests"

Requests satisfied by already retained polynomial powers during this call.

"NewPolynomialPowerEvaluations"

New expanded polynomial-power computations, including uncached requests when the optional cache is full. This count is separate from coefficient evaluations and excludes constant coefficients. It covers incremental coefficient generation; an independent uncached frontier calculation is counted separately at the coefficient level.

"RetainedPolynomialPowers"

The number of powers of degree at least 2 retained in the final state. Capacity and eviction counters record how the resource budget was applied.

Newton statistics separately report the old and new precision, the new step cutoffs, and Euler work used for the remainder frontier. These are exact bookkeeping quantities; they are not a count of all arithmetic operations and do not prove a runtime bound.

The reproducible driver `Tests/BenchmarkRefinement.wl` compares a sequence of fresh public inverse constructions with one construction followed by public refinements. It covers dense rational gaps, two and three irrational gaps, resonant weights, and high logarithmic degrees, using both Lagrange and Newton examples. A deep one-gap case reaches target cutoff 64 and also compares with the independent Catalan-number formula for the inverse of $x + x^2$. Acceptance requires equality of every complete block and remainder at every requested cutoff, together with the retained-state invariants. The driver reports elapsed times and state counters. It measures retained expression and state sizes with `ByteCount` and evaluation peaks with `MaxMemoryUsed[expr]`; these kernel measurements are not total-process memory bounds. No machine-dependent speed or memory threshold determines acceptance.

26 Independent generated validation and reproducible failures

The deterministic tests in `GeneratedAdvancedOracles.wlt` complement the radical inverse pairs in `GeneratedInverseOracles.wlt`. Their oracle is a separate finite symbolic substitution calculation. It does not call the package's coefficient formula, sparse jets, multi-index enumeration, Newton engine, or residual checker. Agreement between package methods remains useful supplementary evidence; it is not the expected-value calculation in these tests.

Write a generated local model as

$$f(u) = au^p \left(1 + \sum_{i=1}^m u^{\delta_i} B_i(\log u) \right), \quad p \neq 0, \quad \delta_i > 0,$$

and put $u = z(1 + V)$. Introduce independent variables t_i and a marker ϵ . The oracle solves the finite Taylor problem

$$(1 + V)^p - 1 + \epsilon \sum_i t_i (1 + V)^{p+\delta_i} B_i(L + \log(1 + V)) = 0, \quad V = \sum_{n=1}^N v_n \epsilon^n.$$

At marker order n , the new coefficient v_n occurs with coefficient p . All other terms are determined by earlier coefficients. Native `Series` and ordinary polynomial coefficient extraction therefore give a triangular scalar recurrence, independently of the Euler–Lagrange formula. The

oracle checks both the coefficient of the new unknown and the vanishing of the substituted forward residual through order N .

For an observable u^r , it then expands $(1 + V)^r$. Each monomial t^k acquires the exact weight $k \cdot \delta$ under the substitution $t_i = z^{\delta_i}$. To retain all weights strictly below H , it suffices to choose

$$N = \left\lceil \frac{H}{\min_i \delta_i} \right\rceil - 1.$$

Indeed, an omitted monomial has total marker degree at least $N + 1$ and hence weight at least $(N + 1) \min_i \delta_i \geq H$. The routine fixtures use small depths; the independent oracle explicitly refuses depths greater than six. Equal algebraic weights are combined by ordinary exact symbolic arithmetic when the expected finite expression is compared with `Normal` of the public inverse object. No numerical tolerance decides equality or ordering.

A fixed sixteen-family cycle ensures coverage of rational logarithmic polynomials, two and three irrational gaps, quadratic logarithmic amplitudes, exact collisions, and cancellation. For example, with $p = r = 1$ the coefficients at the first three relative powers of $u(1 + cu + 2c^2u^2 - u^3)$ include a cancelled second-order block. For two independent irrational gaps α, β , the model

$$u \left(1 + Au^\alpha + Bu^\beta + (2 + \alpha + \beta)ABu^{\alpha+\beta} \right)$$

has cancelling inverse contributions at weight $\alpha + \beta$ from different marker degrees. Both examples test complete weight blocks, rather than counting contributions before they are combined.

Other families exercise source translation, the left branch, reciprocal source coordinates at both infinities, signed leading coefficients, nonunit leading powers, and observable powers. Declared forward remainders are tested at their transported precision boundary and beyond it; the latter must return `Failure["InsufficientInputOrder", ...]`. An intentionally insufficient index budget must return `Failure["ResourceLimit", ...]`. The exact coordinate identities

$$\begin{aligned} y = b + a \log \log x &\iff x = \exp \exp \frac{y - b}{a}, \\ y = b + ae^{-1/x^2} &\iff x = \frac{1}{\sqrt{-\log((y - b)/a)}}. \end{aligned}$$

are checked on their stated positive real branches, including a zero remainder. They do not depend on the recurrence oracle.

The standalone script `RunGeneratedCampaign.wl` uses seed 236369 and the Mersenne Twister by default. Its default twenty-four cases are separate from the sixteen-case routine suite. Environment variables with prefix `ASYMPTOTIC_CAMPAIGN_` select `SEED`, `CASES`, `CASE_SECONDS`, `SHRINK_ATTEMPTS`, `SHRINK_SECONDS`, and `OUTPUT`. Explicit finite bounds are enforced: at most ten thousand cases, six hundred seconds per case or shrink pass, and one thousand shrink attempts per failed case. A requested output directory must be new.

Before testing, the runner writes the complete exact inputs and a manifest containing the seed, generator and oracle revision, actual kernel version, system identifier, resource limits, and SHA-256 digests of the runner and every kernel module. Each case writes its original input, requested operation, expected and observed expressions, messages, and elapsed time before shrinking begins. Timeouts, oracle failures, unexpected package failures, incorrect coefficients, and violated rejection contracts remain distinct outcomes. An interruption after an original case record has been written does not erase that record.

The deterministic shrinker tries deleting perturbations, lowering logarithmic degree, simplifying exact coefficients and powers, decreasing the requested cutoff, and simplifying positive gaps. It accepts a candidate only when an explicit lexicographic tuple of nonnegative integer complexity measures decreases and the same failure class recurs. It preserves the original

case, writes every accepted reduction separately, and stops at its time or attempt budget. Every completed shrink attempt also records its input and result, including a different failure class that cannot be accepted as a reduction. An evidence write failure stops the campaign explicitly. Consequently the result is a smaller reproducer, not a claim of global minimality. Exact coordinate fixtures currently have no specialized shrinking rules; their original evidence is still preserved.

The routine suite also checks repeatability, strict shrink progress, retention of the original case, and rejection of reductions that change the failure class. These are tests of a finite implementation on the kernel version recorded by the native run. They are neither a proof for all expressions nor evidence for an untested Wolfram version. The analytic results elsewhere in this article and numerical or interval certificates have separate evidence roles.

27 Retaining an exact growing exponential core

An inverse-logarithmic expansion of the Lambert function does not resolve the much smaller correction caused by adding x^2 to xe^x . The operation `AsymptoticExponentialCoreInverse` retains the Lambert inverse exactly and expands those corrections in a separate small exponential scale. Its zero sector is an exact inverse, and each retained coefficient is a complete function of that inverse.

27.1 The admitted core and source charts

Use a positive coordinate $v \rightarrow +\infty$ and consider the exact model

$$F(v) = F_0(v) + R(v), \quad F_0(v) = y_0 + av^b e^{cv^p}, \quad c, p > 0, \quad a \neq 0, \quad (27.1)$$

where b is a fixed real number and R is a finite real power-log expression $\sum_\alpha v^\alpha P_\alpha(\log v)$. Every such perturbation, including its finitely many derivatives, is smaller than the growing core sufficiently near the endpoint.

At an infinite original source endpoint, the reconstruction is $x = H(v) = s + \sigma v$ with $\sigma \in \{1, -1\}$. The source translation s can be supplied explicitly or recognized from affine factors in the core. At a finite endpoint, use $x = H(v) = x_0 + \sigma/v$. Thus the same operation covers both $xe^x + x^2$ at infinity and $e^{1/x} + x$ at 0^+ . Define $r_* = 1$ for the first reconstruction and $r_* = -1$ for the second.

Let $Y = (y - y_0)/a > 0$. The selected exact core inverse is

$$v_0 = \begin{cases} (\log Y/c)^{1/p}, & b = 0, \\ \left(\frac{b}{cp} W_\nu \left(\frac{cp}{b} Y^{p/b} \right) \right)^{1/p}, & b \neq 0, \end{cases} \quad \nu = \begin{cases} 0, & b > 0, \\ -1, & b < 0. \end{cases} \quad (27.2)$$

These are the real branches with $v_0 \rightarrow +\infty$. For $b < 0$, the argument is negative and the lower branch selects the increasing large- v part of the core. The target domain records the real Lambert argument condition, positivity of the real-power base, and $b + cpv_0^p > 0$. In all cases the exact inverse in the original source is $\phi(y) = H(v_0(y))$.

An optional supplied `"CoreInverse"` is retained only after a bounded symbolic comparison establishes equality with this selected exact expression. Additional terms inside the stated core are rejected if they would make (27.2) merely an approximation. They can instead be placed in the separately declared finite perturbation.

27.2 Complete coefficients without expanding the core inverse

Put

$$D(v) = av^{b-1}(b + cpv^p), \quad F'_0(v) = e^{cv^p} D(v), \quad E(v) = e^{-cv^p}. \quad (27.3)$$

Introduce the marker λ in $F_0(v) + \lambda R(v) = y$. The exact-core Lagrange formula of Section 18 shows that its degree- n correction has the form $C_n(v_0)E(v_0)^n$. The coefficient C_n can be computed using only power-logarithmic and rational operations on v :

$$Q_{n,1}(v) = \frac{H'(v)R(v)^n}{D(v)}, \quad (27.4)$$

$$Q_{n,j+1}(v) = \frac{Q'_{n,j}(v) - jcpv^{p-1}Q_{n,j}(v)}{D(v)}, \quad 1 \leq j < n, \quad (27.5)$$

$$C_n(v) = \frac{(-1)^n}{n!} Q_{n,n}(v). \quad (27.6)$$

Indeed, differentiating $Q(v)E(v)^j$ gives $(Q' - jcpv^{p-1}Q)E^j$, and division by F'_0 adds one factor of E . The implementation therefore avoids intermediate derivatives of `ProductLog` and avoids mixing the exact exponential factor into each coefficient calculation.

The exact core identity gives a second, often more convenient expression for the small scale:

$$E_0 = e^{-cv_0^p} = \frac{v_0^b}{Y}. \quad (27.7)$$

The result records both expressions and uses the right-hand side in the retained formula. It is positive on the selected branch. Through inclusive sector degree N , the approximation is

$$\phi(y) + \sum_{n=1}^N C_n(v_0(y))E_0(y)^n. \quad (27.8)$$

Here one complete marker degree is also one exponential degree, because the forward perturbation contains no additional exponentials.

Example 27.1 (The inverse of $xe^x + x^2$). Let $v_0 = W_0(y)$ and $E_0 = v_0/y$. The first coefficients are

$$C_1(v) = -\frac{v^2}{v+1}, \quad (27.9)$$

$$C_2(v) = \frac{v^3(-v^2 + 2v + 4)}{2(v+1)^3}. \quad (27.10)$$

Thus the first correction to the exact core inverse is

$$(xe^x + x^2)^{-1}(y) = W_0(y) - \frac{W_0(y)^3}{(W_0(y) + 1)y} + \dots$$

This is much smaller than the error of any fixed inverse-logarithmic truncation of $W_0(y)$. Retaining W_0 exactly makes that correction meaningful without requiring a second, deeper core approximation.

27.3 A uniform complete-tail theorem

For the finite perturbation in (27.1), choose

$$\begin{aligned} d &= \max(0, \max_{\alpha}(\alpha - b)), & k &= \max_{\alpha} \deg P_{\alpha}, \\ B(v) &= v^d(1 + |\log v|)^k, & \chi(v) &= B(v)e^{-cv^p}. \end{aligned} \quad (27.11)$$

The zero perturbation is treated separately and leaves the exact inverse unchanged.

Theorem 27.2 (Complete tail around a growing exact core). *For fixed data in (27.1), there exist positive constants C, K and a source threshold such that the selected real inverse has a convergent expansion (27.8) with its complete omitted tail bounded by*

$$|\text{Tail}_N| \leq C v_0^{r_*-p} \frac{(K\chi(v_0))^{N+1}}{1 - K\chi(v_0)}, \quad K\chi(v_0) < 1. \quad (27.12)$$

In particular, for fixed N the remainder class is

$$\mathcal{O}\left(E_0^{N+1} v_0^{r_*-p+(N+1)d} (1 + |\log v_0|)^{(N+1)k}\right). \quad (27.13)$$

The bound covers the whole omitted series independently of cancellation in the first omitted coefficient.

Proof. Set

$$v = v_0(1 + \epsilon V), \quad \epsilon = v_0^{-p}.$$

The displacement scale v_0^{1-p} keeps the change in the exponential phase bounded. Dividing the forward equation by $av_0^b e^{cv_0^p}$ gives a core difference

$$(1 + \epsilon V)^b \exp\left(c \frac{(1 + \epsilon V)^p - 1}{\epsilon}\right) - 1. \quad (27.14)$$

It extends analytically to $\epsilon = 0$, and its derivative in V at $(\epsilon, V) = (0, 0)$ is $cp \neq 0$.

The normalized perturbation is $E_0 R(v_0(1 + \epsilon V))/(av_0^b)$. After extraction of its powers of v_0 and $\log v_0$, each of its finite coefficient factors is analytic and uniformly bounded near $\epsilon = 0, V = 0$, with $1/\log v_0$ as an additional small parameter. Its absolute envelope is a constant times $B(v_0)E_0$. Replace E_0 temporarily by a complex variable Z and rescale $\zeta = B(v_0)Z$. The remaining normalized coefficient parameters range over a bounded finite polydisc. Since the derivative cp at $Z = 0$ is independent of them, the analytic implicit function theorem gives a uniform solution near $V = 0$ for $|\zeta|$ smaller than a fixed radius. Equivalently, it is analytic on $|Z| < 1/(KB(v_0))$.

For the observable, subtract its exact core value. The nonconstant part is

$$H(v_0(1 + \epsilon V)) - H(v_0) = \sigma v_0^{r_*-p} \frac{(1 + \epsilon V)^{r_*} - 1}{\epsilon}.$$

The quotient is analytic near $(\epsilon, V) = (0, 0)$ and uniformly bounded on a smaller polydisc. Cauchy estimates bound its Z^n coefficient by $C(KB)^n$, with the displayed factor $v_0^{r_*-p}$. The coefficient identity above identifies these coefficients with the exact-core Lagrange coefficients. Summing the geometric tail and setting $Z = E_0$ proves (27.12). Finally $\chi(v_0) \rightarrow 0$, and the eventual derivative of the real forward function has the core's nonzero sign, so this solution is the selected real inverse near the endpoint. $\square$

For $xe^x + x^2$, one can take $d = 1, k = 0, p = 1, r_* = 1$. The complete tail after sector N is consequently $\mathcal{O}((v_0 E_0)^{N+1})$. For $e^{1/x} + x$ at 0^+ , the reconstruction has $r_* = -1$, and the first correction is

$$(e^{1/x} + x)^{-1}(y) = \frac{1}{\log y} + \frac{1}{y \log^3 y} + \mathcal{O}\left(\frac{1}{y^2 \log^2 y}\right). \quad (27.15)$$

The last bound is conservative but follows directly from (27.13) with $d = 0$.

27.4 An unknown input error remains a separate first sector

Suppose the actual perturbation includes a real continuously differentiable unknown term $A(v)$ satisfying

$$A(v) = \mathcal{O}(v^{-\rho}(1 + |\log v|)^\ell), \quad A'(v) = \mathcal{O}(v^{-\rho-1}(1 + |\log v|)^\ell). \quad (27.16)$$

Here ρ may be any fixed real number: the growing exponential core dominates every such finite power scale. Its derivative is asymptotic to $acp v^{b+p-1} e^{cv^p}$. The mean value transport argument, with the derivative of H , gives an additional inverse error

$$\mathcal{O}(E_0 v_0^{r_*-b-p-\rho}(1 + |\log v_0|)^\ell). \quad (27.17)$$

The matching derivative assumption ensures eventual monotonicity and permits comparison on the same local inverse branch. The model inverse and the actual inverse are both asymptotic to the exact core coordinate, so the scale can be expressed in v_0 .

Equation (27.17) has exponential degree one, even when the finite model has been expanded through a much larger sector depth. The implementation therefore stores it separately and adds it to the model-tail remainder. It does not relabel it as a higher sector or claim that additional exact model coefficients overcome unknown input data. The retained terms beyond this error still describe the explicitly specified finite model.

27.5 Implemented behavior and verification

```
AsymptoticExponentialCoreInverse[
  x Exp[x], x^2, {x, Infinity}, {y, 2}]

AsymptoticExponentialCoreInverse[
  7 + 3 (x - 2) Exp[x - 2], 5 (x - 2)^2,
  {x, Infinity}, {y, 1}]

AsymptoticExponentialCoreInverse[
  Exp[1/(2 - x)], 2 - x, {x, 2}, {y, 1},
  Direction -> "FromBelow"]
```

The result identifies scale **"ExactExponentialCoreSectors"**. It records the exact inverse and branch certificate, the divergent source coordinate and reconstruction, the complete sector coefficients, the first omitted coefficient, both representations of E_0 , and the separate model and input-error remainders. Its **"MajorantContract"** is an asymptotic existence theorem, with **"NumericCertificate"** set to **False**; the constants and uniform threshold are not computed numerical bounds.

Independent regression checks insert $v + \sum_{n=1}^N C_n(v)Z^n$ into $Xe^{X-v} - v + ZX^2 = 0$ for the primary example and perform ordinary Taylor expansion in Z . Further checks use original forward targets generated from known source points, compare successive sector depths at high precision, and exercise affine shifts and scales, negative source infinity, the lower real Lambert branch, nonunit phase powers and finite endpoint reconstruction. Rejection cases cover a decaying core, an additional term incorrectly placed in the exact core, an extra exponential rank in the perturbation, an unverified inverse, and unproved coefficient reality.

The admitted phase is a single growing monomial cv^p . General phase polynomials would need an appropriate exact inverse supplied with its own branch contract, while perturbations such as $e^{v/2}$ need more than the single exponential sector used here. The operation does not replace those missing contracts with a finite inverse-logarithmic core approximation.

28 Separate inner precision and calculus of flat sectors

The operations `FlatSeriesTruncate`, `FlatSeriesMultiply`, `FlatSeriesObservable`, and `FlatSeriesDifferentiate` provide a finite calculus for the scale of Section 22. They require the same positive monomial target coordinate

$$u = \left(\frac{y - y_0}{a} \right)^{1/q}, \quad E = e^{-c/u^p}, \quad p, c > 0,$$

for every operand. They do not change the exponential phase or silently identify different monomial charts.

Their representation distinguishes two errors:

$$S(y) = C_0(u) + \sum_{k=1}^N E^k \left(T_k(u) + \mathcal{O}(u^{\rho_k} M(u)^{d_k}) \right) + \mathcal{O}(E^{N+1} u^{\rho_*} M(u)^{d_*}). \quad (28.1)$$

Here C_0 is exact, each T_k is a finite power–log expression, and $\rho_k = +\infty$ means that the coefficient is exact. The displayed **Remainder** is a sum with one term for each nonzero inner error and a separate term for the complete omitted exponential tail. **InnerRemainders** and **SectorRemainder** expose the two parts directly. These are asymptotic existence bounds for fixed input data, with no computed numerical constant or threshold.

28.1 Inner truncation and multiplication

`FlatSeriesTruncate[s,h]` retains powers strictly below h in every positive sector. The zero sector remains exact, including any of its powers at or above h . If a coefficient loses its first term at power α , its inner error is recorded at α with the degree of the corresponding logarithmic polynomial. An existing, less precise bound is preserved. Increasing the requested cutoff after truncation does not recover discarded coefficients.

For example, truncating the inner coefficients of

$$S = y - y^2 E + (y^2 + 2y^3) E^2 + \mathcal{O}(y^2 E^3), \quad E = e^{-1/y},$$

at $h = 3$ gives

$$S = y - y^2 E + y^2 E^2 + E^2 \mathcal{O}(y^3) + \mathcal{O}(y^2 E^3).$$

The error $E^2 \mathcal{O}(y^3)$ cannot be put into the third sector: its ratio to any E^3 times a fixed power–log expression is unbounded as $y \rightarrow 0^+$.

Multiplication convolves exponential degrees and uses the ordinary precision-tracked power–log multiplication inside each degree. If $A = \sum_{i \leq N_A} E^i A_i + R_A$ and $B = \sum_{j \leq N_B} E^j B_j + R_B$, the result retains degrees through $N = \min(N_A, N_B)$. Its tail includes all three kinds of contributions:

$$\sum_{i+j > N} E^{i+j} A_i B_j, \quad R_A \sum_j E^j B_j + R_B \sum_i E^i A_i, \quad R_A R_B.$$

Each coefficient bound includes its own inner uncertainty. Since $0 < E < 1$, a term in degree $k > N$ can conservatively be bounded in degree $N + 1$ while retaining its power–log envelope. This can give a looser algebraic power in the final sector bound; it is always a valid bound. In particular, unknown tails are added by bounds and never cancelled because some displayed coefficients cancel.

`FlatSeriesObservable[s,H,z]` implements polynomial composition in a separate formal variable z . The coefficients of H may be exact finite real power–log expressions in the target

chart. Horner evaluation uses the same addition and multiplication contracts, and computes $H(C_0)$ exactly as the zero sector. A constant observable has no dependence on the unknown input tail; the zero observable therefore has zero remainder. General nonpolynomial composition and changes of phase remain outside this API. The degree limit is explicit and defaults to 32; intermediate sector products and coefficient expressions also consume the resource budget.

28.2 Qualified differentiation

A magnitude bound alone does not justify its differentiated bound. Differentiation is enabled for the exact flat inverse only with the stored `ExactMonomialFlatIFT` provenance and its claim for all fixed derivative orders. The constructor's exact finite forward model is essential here. Addition, multiplication, polynomial observables, and inner truncation preserve this analytic provenance. Removing it from an object with a nonzero remainder causes `UnprovedFlatRemainderDerivative`.

For a retained monomial the target derivative is explicit:

$$\begin{aligned} \frac{d}{dy} \left(E^k u^\alpha P(\log u) \right) &= \frac{E^k}{aq} \left[u^{\alpha-q} (\alpha P(\log u) + P'(\log u)) \right. \\ &\quad \left. + kcp u^{\alpha-p-q} P(\log u) \right]. \end{aligned} \quad (28.2)$$

Thus the exponential degree stays k . For $k > 0$, a qualified inner bound at (ρ_k, d_k) is transported to $(\rho_k - p - q, d_k)$. When known differentiated terms are then discarded at that same boundary power, their logarithmic degree can increase the final inner degree bound. The complete tail has the same power shift $\rho_* \mapsto \rho_* - p - q$. This formula applies also when $q < 0$; the shift is a signed exact number, not its absolute value.

For completeness, the analytic tail estimate supports this derivative claim. Fix a small positive real u and use a complex disk

$$|v - u| < \eta u^{p+1},$$

with fixed sufficiently small $\eta > 0$. On this disk the chosen local power and logarithm branches are holomorphic, their power-log envelopes are comparable to those at u , and

$$v^{-p} - u^{-p} = \mathcal{O}(1).$$

Consequently e^{-c/v^p} differs from e^{-c/u^p} by a bounded factor. The normalized implicit-function construction in Section 22 extends uniformly to these disks. Its geometric tail estimate therefore holds there with changed constants. Cauchy's estimate costs a factor u^{-p-1} for one source derivative. The chain-rule factor $du/dy = (aq)^{-1} u^{1-q}$ gives the asserted loss u^{-p-q} for one target derivative. Every fixed higher derivative is handled on nested disks, or by the corresponding target disk of radius comparable to u^{p+q} ; constants may depend on the derivative order. No differentiation of an arbitrary real-axis Big-O bound is used.

The finitely many discarded inner terms are themselves analytic power-log functions, so the same bound, and the exact phase derivative, apply to their errors. After differentiation, the retained rows are clipped to the transported precision. For the previous example, the derivative of the truncated representation has the valid form

$$S' = 1 - (1 + 2y)E + 2E^2 + E^2\mathcal{O}(y) + \mathcal{O}(E^3),$$

where the last bound is the transported $\mathcal{O}(y^2 E^3)$. The full second-sector derivative contains further powers of y that the truncated data do not determine; the implementation does not invent them.

Regression oracles include ordinary marker-polynomial substitution, direct differentiation of exact retained expressions, nonunit and negative leading powers, logarithmic envelope degrees, scalar precision transport, and errors retained at their original exponential degree. They also check rejection of incompatible charts and phases, nonpolynomial observables, unqualified remainder derivatives, and exhausted resource budgets.

29 Integrated refinement and numerical evidence

The common constructor preserves the ordinary coefficient engine when the input already has polynomial logarithmic coefficients. In particular, a request for Newton inversion of $x + x^2 \log x$ retains that method. A reciprocal or generalized logarithmic coefficient instead selects the larger algebra of Section 21. Explicit sector and special-function constructors retain their own cutoff semantics.

29.1 Three meanings of refinement

A numeric second argument to `SeriesRefine` is an exact exclusive cutoff for exponent-truncated objects, or an integer marker depth when replaying a depth-truncated source. Compatible ordinary inverse objects reuse the states described in Section 25; logarithmic objects replay their retained original function at the requested logarithmic or source-power cutoff. A derived result replays its operation recipe and preserves the precision limits of its operands.

```
s = AsymptoticInverse[x + x^2, {x, 0}, {y, 3}];
SeriesRefine[s, 6]
SeriesRefine[s, <|"AdditionalBlocks" -> 3|>]
SeriesRefine[s, <|"Target" -> 1/100, "TargetError" -> 10^-30,
  "Interval" -> {1/200, 1/50}|>]
```

The additional-block request counts complete nonzero exponent blocks of an ordinary inverse, including all resonant contributions at each weight. It advances the existing cutoff by the smallest normalized positive gap and delegates to the retained-state refinement. A finite exact inverse returns immediately when no further blocks exist, even if the numerical block goal is larger. The result records both `"GoalReached"` and `"ExactTermination"`; these statements are different. A declared input error or a resource limit can end refinement earlier, with the best available expansion retained in the failure data.

The tolerance request returns a numerical certificate association, with `"ResultType" -> "NumericalCertificate"`. It invokes the exact rational interval operation of Section 19, retains the source expansion, and reports a certified rational source-root center. The tolerance concerns that root even if the retained symbolic object represents a power observable. It does not replace an asymptotic Big- O by a numeric tolerance. Absolute and relative goals have the same interval and zero-root conventions as `InverseCertificate`. An unknown request key is rejected rather than silently ignored.

29.2 Checking observables against the original equation

`InverseNumericalCheck` accepts ordinary, logarithmic, Fourier, flat, exact-core and exponential-core inverses. Special-function adapters use their stable original-function check. The source and target chart routes retain their exact transformed equations. For an observable with power

$r \neq 1$, the numerical reference is

$$H_r(x_*) = \begin{cases} (x_* - x_0)^r, & x_0 \text{ finite,} \\ x_*^r, & x_0 = \pm\infty. \end{cases}$$

For $r = 1$, the displayed inverse and the numerical reference both include the source translation. The selected source sign recovers a root seed from the observable, including even powers at negative infinity. The computed root must lie on the selected real source side.

The returned **"ReferenceRoot"** and **"ReferenceObservable"** have different meanings, and **"Error"** compares the observable with its finite expansion. **"ForwardResidual"** is evaluated at the recovered source seed. The old **"ExactInverse"** field remains a compatibility alias for the numerical reference root; its name is not evidence of an exact or certified value. Every such comparison is explicitly numerical evidence. Unknown forward remainders cannot be evaluated and remain outside the stored equation's numerical scope.

Exact target substitution precedes decimal evaluation. Thus a target $10^{100} + 1/100$ near the limiting value 10^{100} is reduced to its exact distance before finite-precision arithmetic. Rounding the target first would erase that distance at ordinary working precisions and select the wrong branch or seed. Transformed targets follow the same rule.

The integration tests compare observable powers with an independent quadratic radical, retain the lower Lambert branch, resolve flat-sector differences at high precision, check the original Fourier and special function equations, and exercise failed precision and branch contracts. These checks complement the independently generated coefficient and composition oracles; they do not replace the remainder theorems.

30 Composition and differentiation in a reciprocal logarithm

The reciprocal-logarithmic unit of Section 21 has more regularity than an arbitrary function satisfying a magnitude Big- O estimate. For an exact rational logarithmic forward unit, its implicit inverse is holomorphic in the small reciprocal logarithm. This regularity supplies public composition and differentiation rules without discarding the exact power carrier or asking for an unsupported derivative of an unknown remainder.

The bounded operation implemented in `Kernel/ReciprocalLogOperations.wl` uses the common operations `SeriesCompose`, `SeriesDifferentiate`, and their refinement recipes. It also exposes `ReciprocalLogCompose` and `ReciprocalLogDifferentiate` as explicitly scoped entry points. An inverse input must have scale `ReciprocalLogUnit`, a positive monomial prefactor, zero finite source and target offsets, and an exact rational logarithmic model. Known additional positive source-power terms are excluded from this particular analytic contract: although smaller than all reciprocal-logarithmic orders, they need a separate differentiable remainder argument. Results of the supported operations retain the same holomorphic coefficient-function contract.

30.1 A common positive logarithmic coordinate

Write

$$t = \frac{\epsilon}{\log y} > 0, \quad \epsilon = \begin{cases} -1, & y \rightarrow 0^+, \\ 1, & y \rightarrow +\infty. \end{cases} \quad g(y) = Cy^\alpha A(t), \quad (30.1)$$

Here $C > 0$, α is a fixed exact real number, and A is holomorphic near $t = 0$. For an original inverse unit, $A(0) = 1$. Differentiation can produce a coefficient function whose constant

vanishes or has either sign; these outcomes remain valid carriers for further differentiation and as outer functions in composition.

Suppose the original inverse has positive leading amplitude a and internal leading source power $p \neq 0$. Its recorded small coordinate is $\tau = -1/\log((y/a)^{1/p})$, and therefore

$$\tau = \frac{|p|t}{1 - \epsilon(\log a)t}, \quad \epsilon = -\text{sign } p. \quad (30.2)$$

This is an analytic change of variable at zero with positive derivative. It transports an analytic $\mathcal{O}(\tau^\beta)$ tail to $\mathcal{O}(t^\beta)$. For an internal observable power r_* , the monomial carrier has $\alpha = r_*/p$ and $C = \sigma^r a^{-\alpha}$. The branch guards require this C to be positive. In addition to the original target domain, the canonical chart requires $y > 0$ and $\epsilon \log y > 0$. These are eventual branch conditions; they are not a computed radius of convergence.

All stored jets have nonnegative integer powers of t and constant real coefficients. A finite exclusive cutoff h retains powers $n < h$; the actual error exponent is limited by the inherited Taylor remainder and any explicitly discarded coefficient. Coefficient cancellation can make the recorded error order larger than the requested display cutoff.

30.2 Composition of two monomial carriers

Let the outer function and inner function be

$$g(v) = C_o v^\alpha A\left(\frac{\epsilon_o}{\log v}\right), \quad h(y) = C_i y^b V(t), \quad t = \frac{\epsilon_i}{\log y}, \quad V(0) = 1.$$

A positive nonunit constant in the inner coefficient function is first absorbed into C_i . The composition requires $b \neq 0$ and

$$\epsilon_o b \epsilon_i > 0. \quad (30.3)$$

This condition says that the positive inner function approaches the outer function's selected endpoint. It permits a negative power to interchange zero and infinity when the two specified endpoints agree with that change.

Taking the real logarithm of the inner carrier gives the exact identities

$$T(t) = \frac{\epsilon_o t}{b\epsilon_i + t(\log C_i + \log V(t))}, \quad (30.4)$$

$$g(h(y)) = C_o C_i^\alpha y^{\alpha b} V(t)^\alpha A(T(t)). \quad (30.5)$$

The denominator at zero is nonzero, $T(0) = 0$, and $T'(0) = \epsilon_o/(b\epsilon_i) > 0$. The functions $\log V$ and V^α use their analytic branches with $V(0) = 1$. Consequently the new coefficient function is holomorphic near zero and the original monomial carrier remains exact.

If the outer and inner Taylor remainders are $\mathcal{O}(T^{\beta_o})$ and $\mathcal{O}(t^{\beta_i})$, respectively, a conservative relative remainder for (30.5) is

$$\mathcal{O}\left(t^{\min(\beta_o, \beta_i)}\right). \quad (30.6)$$

To see this, an inner perturbation of size $\mathcal{O}(t^{\beta_i})$ changes $\log V$ and V^α by that order, while it changes T by $\mathcal{O}(t^{\beta_i+2})$. The outer function has bounded derivatives near zero, and its original remainder is evaluated at a coordinate comparable to t . The precision-aware polynomial operations preserve these bounds and can improve them when a constant factor or derivative vanishes. Truncating the displayed polynomial adds its own first discarded order. A larger requested cutoff never removes either operand's unknown tail.

For example, if both functions have the form $yW(-1/\log y)$ near zero, the coefficient function of their composition is

$$W(t) W\left(\frac{t}{1 - t \log W(t)}\right).$$

At positive infinity the denominator is instead $1 + t \log W(t)$. These formulas provide short independent ordinary Taylor-series oracles. For general powers and scales, (30.4) retains both the exponent factor b and the additive logarithm $\log C_i$; dropping either gives incorrect coefficients already at low reciprocal-logarithmic orders.

An inner coefficient function with a zero constant term is excluded from this composition rule. Its logarithm introduces an additional $\log t$ scale. A negative constant also fails the positive inner branch condition. Those failures describe a change in the required scale or branch, rather than a reason to take a logarithm of an unproved positive approximation.

30.3 Differentiating the analytic remainder

The coordinate derivative is

$$\frac{dt}{dy} = -\frac{\epsilon t^2}{y}.$$

Hence

$$\frac{d}{dy}(Cy^\alpha A(t)) = Cy^{\alpha-1} \left(\alpha A(t) - \epsilon t^2 A'(t) \right). \quad (30.7)$$

For the n th derivative, the coefficient function is obtained by applying the successive operators

$$\prod_{k=0}^{n-1} \left(\alpha - k - \epsilon t^2 \frac{d}{dt} \right) \quad \text{to } A(t), \quad (30.8)$$

and the carrier is $Cy^{\alpha-n}$. The factors differ only by constants and commute; the implementation nevertheless applies them in derivative order to track precision and cancellation explicitly.

If a holomorphic remainder has a zero of order at least β at zero, then its j th derivative is $\mathcal{O}(t^{\beta-j})$ for every fixed j . This follows directly from its convergent Taylor series, or from Cauchy estimates on a smaller disk. Thus the $t^2 A'$ contribution in (30.7) has remainder $\mathcal{O}(t^{\beta+1})$, while a nonzero αA contribution has remainder $\mathcal{O}(t^\beta)$. If $\alpha = 0$, the latter contribution disappears and the relative logarithmic error improves by one power. This is a proved derivative contract for the exact rational-logarithmic implicit model; it is not a rule for differentiating an arbitrary magnitude Big- $\mathcal{O}$.

For the inverse of $x + x/\log x$ near zero, the retained coefficient function begins

$$A(t) = 1 + t + t^2 + 2t^3 + \frac{7}{2}t^4 + \mathcal{O}(t^5).$$

The first derivative therefore has the expansion

$$1 + t + 2t^2 + 4t^3 + \frac{19}{2}t^4 + \mathcal{O}(t^5).$$

Its carrier is constant. Differentiating once more gives a carrier y^{-1} and a relative remainder $\mathcal{O}(t^6)$, provided the known terms generated through degree five are retained. This tests the precision improvement when a carrier derivative vanishes. At infinity the signs of the odd powers and the coordinate derivative change consistently.

30.4 Provenance, refinement, and evidence

```
a = AsymptoticLogarithmicInverse[
  x + x/Log[x], {x, 0}, {z, 4}];
b = AsymptoticLogarithmicInverse[
  x + 2 x/Log[x], {x, 0}, {y, 4}];
c = SeriesCompose[a, b];
d = SeriesDifferentiate[c];
SeriesRefine[d, 6]
```

The resulting scale is **ReciprocalLogCalculus**. Its exact carrier, analytic coefficient jet, endpoint sign, operand domains, and operation recipe are retained alongside the ordinary explicit **SeriesRepresentation**. Refinement replays the original reciprocal-log inverse constructors and then the stored composition or derivative recipe. It therefore gains new coefficient evidence before reporting a stronger remainder. A truncation of the existing expression alone cannot accomplish that refinement. Constants and convergence thresholds remain existential; the analytic contract sets **NumericCertificate** to **False**.

The regressions compare composition against ordinary native Taylor expansion of independently written substitution formulas, including fractional carrier powers and amplitude shifts. Derivative tests use both explicit low-order coefficient formulas and direct differentiation of the retained expression with independently tracked order. Numerical tests start from a known source point x_* , form the exact forward composition $y = f_i(f_o(x_*))$, and compare both the inverse composition and its derivative with x_* and $1/(f'_o(x_*)f'_i(f_o(x_*)))$. Further cases verify both operand-error caps, compatible and incompatible endpoint changes, signed inner values, excluded higher-power sectors, resource bounds, and public refinement recipes. These checks support the algebra and implementation; the holomorphic argument supplies the complete remainder theorem.

References

- [1] V. Reshetnikov, *How to get an asymptotic of the real-valued branch of the inverse function?*, Mathematica Stack Exchange, question 236367, December 2020. <https://mathematica.stackexchange.com/questions/236367/>
- [2] V. Reshetnikov, *Asymptotic expansion for a function containing irrational exponents*, Mathematica Stack Exchange. (Local copy in docs/.)
- [3] NIST Digital Library of Mathematical Functions, §1.10(vii) Lagrange inversion, §4.13 Lambert W -function. <https://dlmf.nist.gov/>
- [4] I. M. Gessel, *Lagrange inversion*, J. Combin. Theory Ser. A **144** (2016), 212–249. arXiv:1609.05988.
- [5] E. Surya, L. Warnke, *Lagrange inversion formula by induction*, arXiv:2305.17576.
- [6] G. A. Edgar, *Transseries for beginners*, Real Anal. Exchange **35** (2009/10), 253–309. arXiv:0801.4877.
- [7] G. A. Edgar, *Transseries: composition, recursion, and convergence*, arXiv:0909.1259.
- [8] J. van der Hoeven, *Transseries and Real Differential Algebra*, Lecture Notes in Mathematics 1888, Springer, 2006.

- [9] B. Salvy, J. Shackell, *Symbolic asymptotics: multiseries of inverse functions*, J. Symbolic Comput. **27** (1999), 543–563.
- [10] R. M. Corless, G. H. Gonnet, D. E. G. Hare, D. J. Jeffrey, D. E. Knuth, *On the Lambert W function*, Adv. Comput. Math. **5** (1996), 329–359.
- [11] V. Reshetnikov, *Transseries: the polynomial-logarithmic calculus, series reversal at infinity, and the inversion of rapidly growing functions* and *Combinatorial transseries and their inverses*, ProveIt repository, directory `Analysis/FabiusFunction/docs/semi-formalized-research-frontiers/drafts/series-and-transseries`, September 2026. <https://github.com/VladimirReshetnikov/ProveIt>
- [12] The nine reports in the directory `reports/` of this repository (`reports/COMPARISON.md`), 7 September 2026.
- [13] Wolfram Research, *Wolfram Language documentation*. Reference pages for series expansion and inversion, asymptotic computation, algebraic-number reduction, and the Lambert W function. <https://reference.wolfram.com/language/>